

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import plotly.offline as py
py.init_notebook_mode()
%matplotlib inline
import seaborn as sns
from pylab import rcParams
```

```
In [2]: df = pd.read_csv("Rose.csv", parse_dates=True, index_col=0)
```

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 187 entries, 1980-01-01 to 1995-07-01
Data columns (total 1 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0   Rose    185 non-null      float64
dtypes: float64(1)
memory usage: 2.9 KB
```

```
In [4]: df.head()
```

```
Out[4]:
```

	Rose
YearMonth	
1980-01-01	112.0
1980-02-01	118.0
1980-03-01	129.0
1980-04-01	99.0
1980-05-01	116.0

```
In [5]: df.tail()
```

```
Out[5]:
```

	Rose
YearMonth	
1995-03-01	45.0
1995-04-01	52.0
1995-05-01	28.0
1995-06-01	40.0
1995-07-01	62.0

```
In [6]: df.shape
```


```
Out[6]: (187, 1)
```

```
In [7]: df.isnull().sum()
```

```
Out[7]: Rose      2  
dtype: int64
```

```
In [8]: rcParams['figure.figsize'] = 12,8  
df.plot();  
plt.grid()  
plt.title('Rose Sales Distribution',color='blue',fontsize=18)  
plt.xlabel('Year/Month',color='blue',fontsize=16)  
plt.ylabel('Rose Sales',color='blue',fontsize=16)
```

```
Out[8]: Text(0, 0.5, 'Rose Sales')
```



```
In [9]: # Missing value imputation using interpolation method:
```

```
data=df.interpolate(method='polynomial',order=2)
```

```
In [10]: data.shape
```

```
Out[10]: (187, 1)
```


```
In [11]: data.isna().sum()
```

```
Out[11]: Rose      0  
dtype: int64
```

In [12]: *#Plotting the data with interpolated values:*

```
rcParams['figure.figsize'] = 12,8
data.plot();
plt.grid()
plt.title('Rose Sales Distribution',color='blue',fontsize=18)
plt.xlabel('Year/Month',color='blue',fontsize=16)
plt.ylabel('Rose Sales',color='blue',fontsize=16)
```

Out[12]: Text(0, 0.5, 'Rose Sales')




```
In [13]: data.describe()
```

Out[13]:

Rose	
count	187.000000
mean	89.907184
std	39.246679
min	28.000000
25%	62.500000
50%	85.000000
75%	111.000000
max	267.000000

```
In [14]: # Analyzing data:  
# Year-on-year plot of the sales:  
  
fig, ax = plt.subplots(figsize=(20,10))  
sns.boxplot(data.index.year, data.values[:,0], ax=ax,whis=1.5)  
plt.grid();  
plt.xlabel('Year/Month',color='blue',fontsize=18);  
plt.ylabel('Rose Sales',color='blue',fontsize=18);  
plt.title('Year-on-Year Rose Sales Distribution',color='blue',fontsize=20)
```

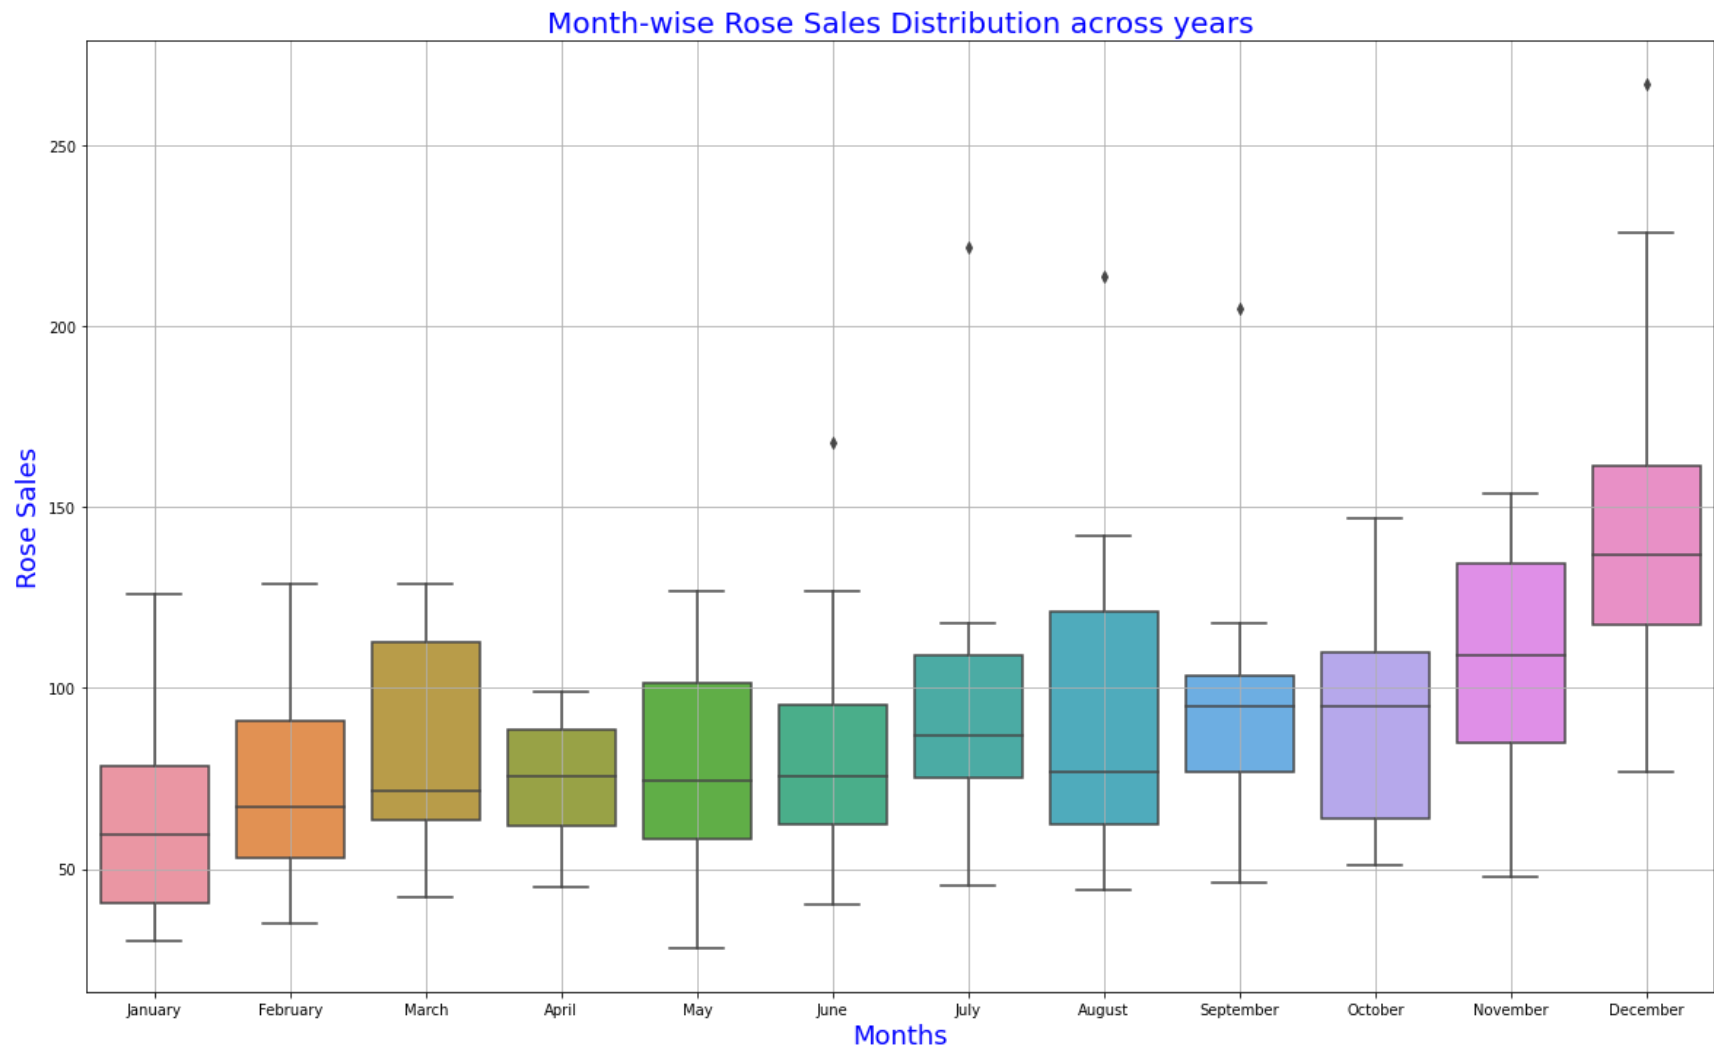
Out[14]: Text(0.5, 1.0, 'Year-on-Year Rose Sales Distribution')



In [15]: *#Plotting month-wise sales distribution across years:*

```
fig, ax = plt.subplots(figsize=(20,12))
sns.boxplot(data.index.month_name(), data.values[:,0], ax=ax,whis=1.5)
plt.grid();
plt.xlabel('Months',color='blue',fontsize=18);
plt.ylabel('Rose Sales',color='blue',fontsize=18);
plt.title('Month-wise Rose Sales Distribution across years',color='blue',fontsize=20)
```

Out[15]: Text(0.5, 1.0, 'Month-wise Rose Sales Distribution across years')



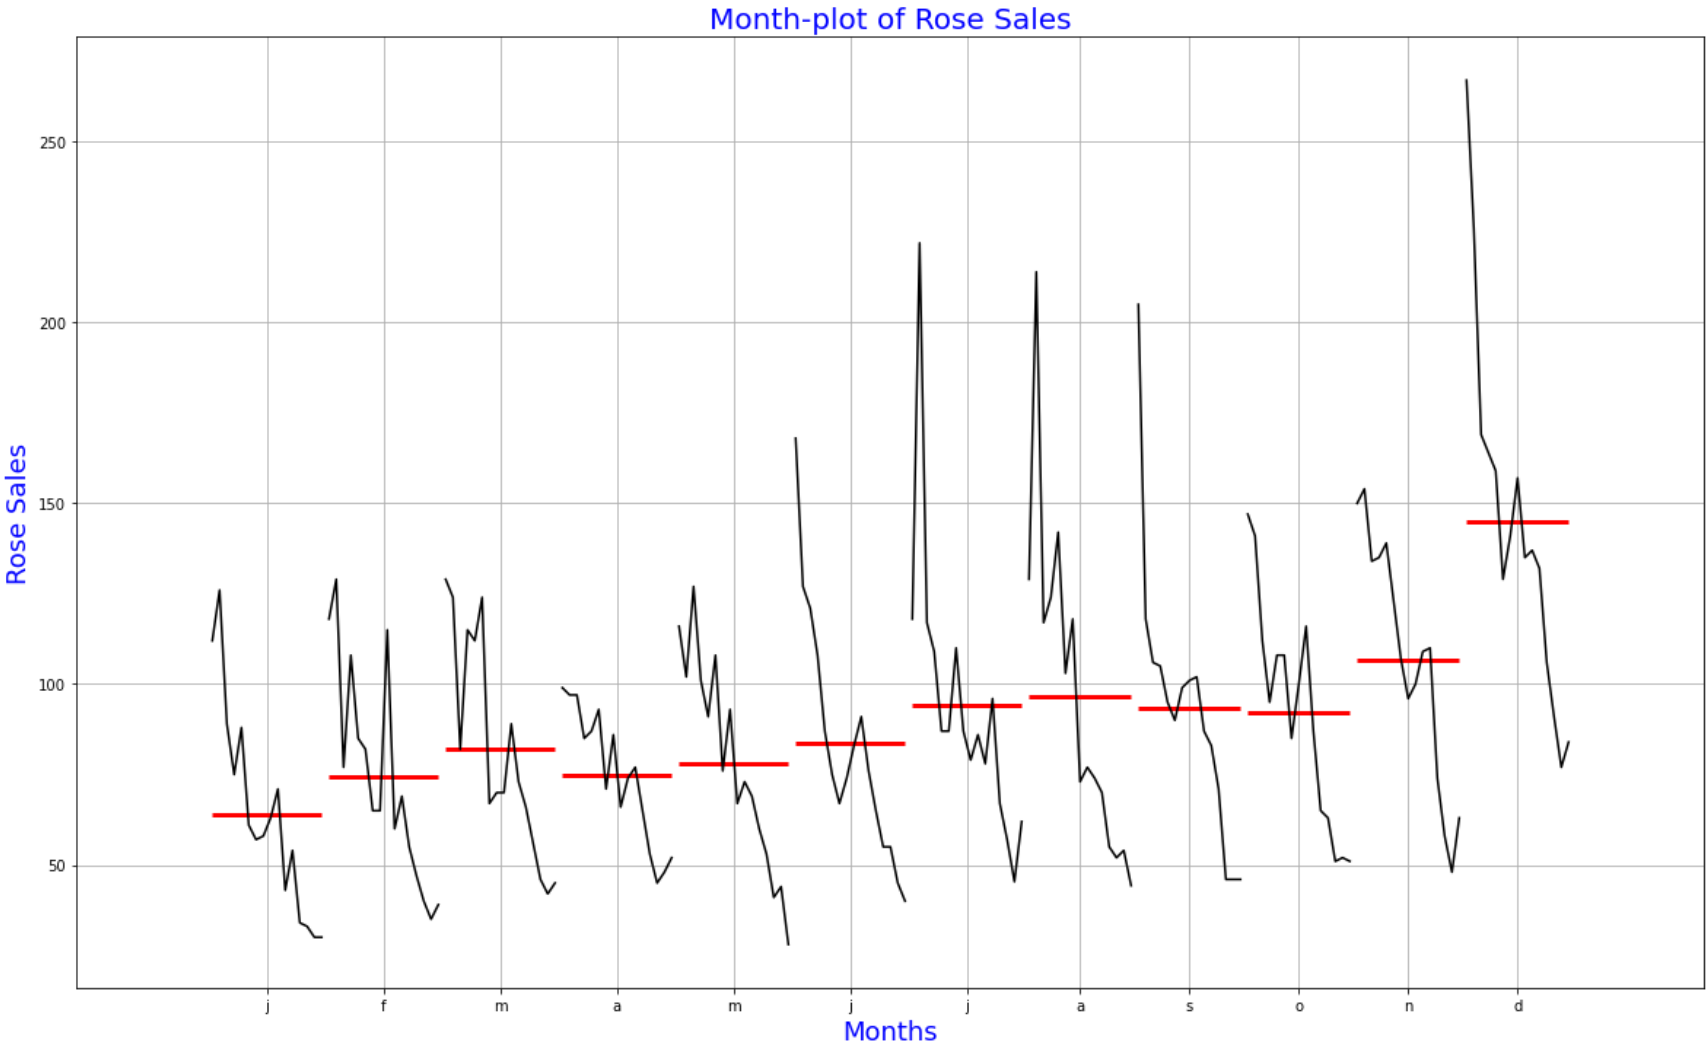
```
In [16]: # Plotting a month-plot:

from statsmodels.graphics.tsaplots import month_plot

fig, ax = plt.subplots(figsize=(20,12))

month_plot(data,ax=ax)
plt.grid();
plt.xlabel('Months',color='blue',fontsize=18);
plt.ylabel('Rose Sales',color='blue',fontsize=18);
plt.title('Month-plot of Rose Sales',color='blue',fontsize=20)
```

```
Out[16]: Text(0.5, 1.0, 'Month-plot of Rose Sales')
```



In [17]: *# Computing and plotting the per month sales for each year:*


```
monthly_sales_across_years = pd.pivot_table(data, values = 'Rose', columns = data.index.month_name(), in
monthly_sales_across_years
```

Out[17]:

YearMonth	April	August	December	February	January	July	June	March	May	November	October	September
YearMonth												
1980	99.0	129.000000	267.0	118.0	112.0	118.000000	168.0	129.0	116.0	150.0	147.0	205.0
1981	97.0	214.000000	226.0	129.0	126.0	222.000000	127.0	124.0	102.0	154.0	141.0	118.0
1982	97.0	117.000000	169.0	77.0	89.0	117.000000	121.0	82.0	127.0	134.0	112.0	106.0
1983	85.0	124.000000	164.0	108.0	75.0	109.000000	108.0	115.0	101.0	135.0	95.0	105.0
1984	87.0	142.000000	159.0	85.0	88.0	87.000000	87.0	112.0	91.0	139.0	108.0	95.0
1985	93.0	103.000000	129.0	82.0	61.0	87.000000	75.0	124.0	108.0	123.0	108.0	90.0
1986	71.0	118.000000	141.0	65.0	57.0	110.000000	67.0	67.0	76.0	107.0	85.0	99.0
1987	86.0	73.000000	157.0	65.0	58.0	87.000000	74.0	70.0	93.0	96.0	100.0	101.0
1988	66.0	77.000000	135.0	115.0	63.0	79.000000	83.0	70.0	67.0	100.0	116.0	102.0
1989	74.0	74.000000	137.0	60.0	71.0	86.000000	91.0	89.0	73.0	109.0	87.0	87.0
1990	77.0	70.000000	132.0	69.0	43.0	78.000000	76.0	73.0	69.0	110.0	65.0	83.0
1991	65.0	55.000000	106.0	55.0	54.0	96.000000	65.0	66.0	60.0	74.0	63.0	71.0
1992	53.0	52.000000	91.0	47.0	34.0	67.000000	55.0	56.0	53.0	58.0	51.0	46.0
1993	45.0	54.000000	77.0	40.0	33.0	57.000000	55.0	46.0	41.0	48.0	52.0	46.0
1994	48.0	44.279246	84.0	35.0	30.0	45.364189	45.0	42.0	44.0	63.0	51.0	46.0
1995	52.0	NaN	NaN	39.0	30.0	62.000000	40.0	45.0	28.0	NaN	NaN	NaN

```
In [18]: monthly_sales_across_years.plot(figsize=(20,10))
plt.grid()
plt.legend(loc='best');
plt.xlabel('Years',color='blue',fontsize=18);
plt.ylabel('Rose Sales',color='blue',fontsize=18);
plt.title('Monthly Rose Sales Distribution across Years',color='blue',fontsize=20)
```

Out[18]: Text(0.5, 1.0, 'Monthly Rose Sales Distribution across Years')



```
In [19]: # Computing and plotting mean sales for each year:
```


```
data_yearly_mean = data.resample('Y').mean()  
data_yearly_mean.head()
```

```
Out[19]:
```

Rose	
YearMonth	
1980-12-31	146.500000
1981-12-31	148.333333
1982-12-31	112.333333
1983-12-31	110.333333
1984-12-31	106.666667

```
In [20]: data_yearly_mean.plot();  
plt.grid()  
plt.xlabel('Mean Sales per year',color='blue',fontSize=18);  
plt.ylabel('Mean Sales',color='blue',fontSize=18);  
plt.title('Mean Rose Sales distribution across Years',color='blue',fontSize=20)
```

```
Out[20]: Text(0.5, 1.0, 'Mean Rose Sales distribution across Years')
```



In [21]: *# Computing and plotting mean sales for each quarter:*


```
data_quarterly_mean = data.resample('Q').mean()  
data_quarterly_mean.head()
```

Out[21]:

Rose	
YearMonth	
1980-03-31	119.666667
1980-06-30	127.666667
1980-09-30	150.666667
1980-12-31	188.000000
1981-03-31	126.333333


```
In [22]: data_quarterly_mean.plot(figsize=(20,10));  
plt.grid()  
plt.xlabel('Mean Sales per quarter',color='blue',fontsize=18);  
plt.ylabel('Mean Sales',color='blue',fontsize=18);  
plt.title('Mean Rose Sales distribution across Quarters',color='blue',fontsize=20)
```

```
Out[22]: Text(0.5, 1.0, 'Mean Rose Sales distribution across Quarters')
```




In [23]: *#Decomposing the Time Series:*

```
from statsmodels.tsa.seasonal import seasonal_decompose
```


```
In [24]: # Additive decomposition:
```

```
decomposition_add = seasonal_decompose(data,model='additive')  
decomposition_add.plot();
```



```
In [25]: # Multiplicative decomposition:
```

```
decomposition_multi = seasonal_decompose(data,model='multiplicative')  
decomposition_multi.plot();
```



```
In [26]: # we will choose multiplicative decomposition as residual is more random, centralised around 1.0
```

```
In [27]: # Computing the various components of the decomposed data:
```

```
trend = decomposition_multi.trend  
seasonality = decomposition_multi.seasonal  
residual = decomposition_multi.resid
```

```
In [28]: # Checking the components:
```

```
print('Trend in Sparkling Sales', '\n', trend.head(12), '\n')
print('Seasonality in Sparkling Sales', '\n', seasonality.head(12), '\n')
print('Residual', '\n', residual.head(12), '\n')
```

Trend in Sparkling Sales

YearMonth	
1980-01-01	NaN
1980-02-01	NaN
1980-03-01	NaN
1980-04-01	NaN
1980-05-01	NaN
1980-06-01	NaN
1980-07-01	147.083333
1980-08-01	148.125000
1980-09-01	148.375000
1980-10-01	148.083333
1980-11-01	147.416667
1980-12-01	145.125000

Name: trend, dtype: float64

Seasonality in Sparkling Sales

YearMonth	
1980-01-01	0.670208
1980-02-01	0.806225
1980-03-01	0.901320
1980-04-01	0.854202
1980-05-01	0.889574
1980-06-01	0.924142
1980-07-01	1.058226
1980-08-01	1.034110
1980-09-01	1.017792
1980-10-01	1.022731
1980-11-01	1.192548
1980-12-01	1.628922

Name: seasonal, dtype: float64

Residual

YearMonth	
1980-01-01	NaN
1980-02-01	NaN
1980-03-01	NaN

```
1980-04-01      NaN
1980-05-01      NaN
1980-06-01      NaN
1980-07-01    0.758124
1980-08-01    0.842160
1980-09-01    1.357482
1980-10-01    0.970621
1980-11-01    0.853235
1980-12-01    1.129454
Name: resid, dtype: float64
```

In [29]: *# Checking how the data looks without seasonality:*


```
deaseasonalized_ts = trend + residual
deaseasonalized_ts.head(12)
```

Out[29]:

```
YearMonth
1980-01-01      NaN
1980-02-01      NaN
1980-03-01      NaN
1980-04-01      NaN
1980-05-01      NaN
1980-06-01      NaN
1980-07-01    147.841457
1980-08-01    148.967160
1980-09-01    149.732482
1980-10-01    149.053954
1980-11-01    148.269902
1980-12-01    146.254454
dtype: float64
```

```
In [30]: data.plot()  
deaseasonalized_ts.plot()  
plt.legend(["Original Time Series", "Deseasonalised Time Series"]);  
plt.xlabel('Years',color='blue',fontsize=18);  
plt.ylabel('Rose Sales',color='blue',fontsize=18);  
plt.title('Original v/s Deseasonalised Rose Sales distribution',color='blue',fontsize=20)
```

Out[30]: Text(0.5, 1.0, 'Original v/s Deseasonalised Rose Sales distribution')



In [31]: *#Splitting data into train and test set:*

```
train = data[data.index<'1991']  
test  = data[data.index>='1991']
```

In [32]: `print(train.shape)`
`print(test.shape)`

```
(132, 1)  
(55, 1)
```

```
In [33]: print('First few rows of Training Data', '\n', train.head(), '\n')
print('Last few rows of Training Data', '\n', train.tail(), '\n')
print('First few rows of Test Data', '\n', test.head(), '\n')
print('Last few rows of Test Data', '\n', test.tail(), '\n')
```

First few rows of Training Data

	Rose
YearMonth	
1980-01-01	112.0
1980-02-01	118.0
1980-03-01	129.0
1980-04-01	99.0
1980-05-01	116.0

Last few rows of Training Data

	Rose
YearMonth	
1990-08-01	70.0
1990-09-01	83.0
1990-10-01	65.0
1990-11-01	110.0
1990-12-01	132.0

First few rows of Test Data


	Rose
YearMonth	
1991-01-01	54.0
1991-02-01	55.0
1991-03-01	66.0
1991-04-01	65.0
1991-05-01	60.0

Last few rows of Test Data

	Rose
YearMonth	
1995-03-01	45.0
1995-04-01	52.0
1995-05-01	28.0
1995-06-01	40.0
1995-07-01	62.0

```
In [34]: # Plotting the train and test data:
```

```
train['Rose'].plot(figsize=(15,8), fontsize=14)
test['Rose'].plot(figsize=(15,8), fontsize=14)
plt.grid()
plt.legend(['Training Data', 'Test Data'])
plt.xlabel('Years',color='blue',fontsize=18);
plt.ylabel('Rose Sales',color='blue',fontsize=18);
plt.title('Rose Sales - Train & Test Data',color='blue',fontsize=20)
plt.show()
```



```
In [35]: # Exponential smoothing models:
```

In [36]: *#Triple exponential smothing using the Holt-Winter's method: (since there is seasonality)*

```
import statsmodels.tools.eval_measures as em
from sklearn.metrics import mean_squared_error
from statsmodels.tsa.api import ExponentialSmoothing, SimpleExpSmoothing, Holt
from IPython.display import display
from pylab import rcParams
```

In [37]: *# Finding the best parameters:*

```
model_TES = ExponentialSmoothing(train,trend='multiplicative',seasonal='multiplicative')

# Fitting the model
model_TES = model_TES.fit()
```

```
print('')
print('~~~ Holt Winters model Exponential Smoothing Estimated Parameters ~~~')
print('')
print(model_TES.params)
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

~~~ Holt Winters model Exponential Smoothing Estimated Parameters ~~~

```
{'smoothing_level': 0.06995960293207605, 'smoothing_slope': 8.106800492700789e-20, 'smoothing_seasonal': 0.0, 'damping_slope': nan, 'initial_level': 76.65465630264949, 'initial_slope': 0.9939044541460769, 'initial_seasons': array([1.45183135, 1.64400934, 1.79905601, 1.5750985 , 1.76974808, 1.90937901, 2.10054341, 2.24522639, 2.11358019, 2.07293345, 2.41587165, 3.31139382]), 'use_boxcox': False, 'lamda': None, 'remove_bias': False}
```

```
In [38]: # Forecasting using this model for the duration of the test set
```

```
TES_predict = model_TES.forecast(len(test))  
TES_predict
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:342: FutureWarning:

The 'freq' argument in Timestamp is deprecated and will be removed in a future version.

```
Out[38]: 1991-01-01      53.999888  
         1991-02-01      60.775090  
         1991-03-01      66.101399  
         1991-04-01      57.519926  
         1991-05-01      64.234255  
         1991-06-01      68.879823  
         1991-07-01      75.314080  
         1991-08-01      80.010925  
         1991-09-01      74.860465  
         1991-10-01      72.973266  
         1991-11-01      84.527285  
         1991-12-01     115.153868  
         1992-01-01      50.179751  
         1992-02-01      56.475652  
         1992-03-01      61.425160  
         1992-04-01      53.450770  
         1992-05-01      59.690103  
         1992-06-01      64.007028  
         1992-07-01      69.986104  
         1992-08-01      74.350678  
         1992-09-01      69.564579  
         1992-10-01      67.810887  
         1992-11-01      78.547535  
         1992-12-01     107.007489  
         1993-01-01      46.629864  
         1993-02-01      52.480371  
         1993-03-01      57.079733  
         1993-04-01      49.669479  
         1993-05-01      55.467421  
         1993-06-01      59.478952  
         1993-07-01      65.035047  
         1993-08-01      69.090856
```

|            |           |
|------------|-----------|
| 1993-09-01 | 64.643342 |
| 1993-10-01 | 63.013713 |
| 1993-11-01 | 72.990813 |
| 1993-12-01 | 99.437413 |
| 1994-01-01 | 43.331107 |
| 1994-02-01 | 48.767730 |
| 1994-03-01 | 53.041718 |
| 1994-04-01 | 46.155690 |
| 1994-05-01 | 51.543466 |
| 1994-06-01 | 55.271207 |
| 1994-07-01 | 60.434245 |
| 1994-08-01 | 64.203132 |
| 1994-09-01 | 60.070251 |
| 1994-10-01 | 58.555907 |
| 1994-11-01 | 67.827193 |
| 1994-12-01 | 92.402870 |
| 1995-01-01 | 40.265717 |
| 1995-02-01 | 45.317733 |
| 1995-03-01 | 49.289365 |
| 1995-04-01 | 42.890478 |
| 1995-05-01 | 47.897104 |
| 1995-06-01 | 51.361132 |
| 1995-07-01 | 56.158920 |


Freq: MS, dtype: float64

```
In [39]: ## Plotting the Training data, Test data and the forecasted values:
```

```
plt.plot(train, label='Train')
plt.plot(test, label='Test')

plt.plot(TES_predict, label='Alpha=0.07, Beta=8.11e-20, Gamma=0.0: TES predictions on Test Set')

plt.legend(loc='best')
plt.grid()
plt.title('Triple Exponential Smoothing Predictions',color='blue',fontsize=20);
plt.xlabel('Years',color='blue',fontsize=18);
plt.ylabel('Rose Sales',color='blue',fontsize=18);
```



In [40]: *#Evaluating the TES method using RSME:*

```
print('TES RMSE:',mean_squared_error(test.values,TES_predict.values,squared=False))
```

TES RMSE: 12.8311055961368

In [41]: *# Storing results to a dataframe:*

```
results = pd.DataFrame({'Test RMSE': [mean_squared_error(test.values,TES_predict.values,squared=False)]},  
                        index=['TES: Alpha=0.07, Beta=8.11e-20, Gamma=0.0'])  
results
```

Out[41]:

|                                           | Test RMSE |
|-------------------------------------------|-----------|
| TES: Alpha=0.07, Beta=8.11e-20, Gamma=0.0 | 12.831106 |



```
In [42]: # Double exponential smoothing using the Holt's method:
```

```
model_DES = Holt(train)
```

```
model_DES = model_DES.fit()
```

```
print('')
```

```
print('~~~ Holt DES model Estimated Parameters ~~~')
```

```
print('')
```

```
print(model_DES.params)
```

```
~~~ Holt DES model Estimated Parameters ~~~
```

```
{'smoothing_level': 0.15789473684210525, 'smoothing_slope': 0.15789473684210525, 'smoothing_seasonal': nan, 'damping_slope': nan, 'initial_level': 112.0, 'initial_slope': 6.0, 'initial_seasons': array([], dtype=float64), 'use_boxcox': False, 'lamda': None, 'remove_bias': False}
```

```
/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:
```

```
No frequency information was provided, so inferred frequency MS will be used.
```

```
/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/holtwinters.py:743: ConvergenceWarning:
```

```
Optimization failed to converge. Check mle_retvals.
```

```
In [43]: # Forecasting using this model for the duration of the test set
```

```
DES_predict = model_DES.forecast(len(test))
DES_predict
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:342: FutureWarning:

The 'freq' argument in Timestamp is deprecated and will be removed in a future version.

```
Out[43]: 1991-01-01 86.863579
 1991-02-01 88.028056
 1991-03-01 89.192534
 1991-04-01 90.357011
 1991-05-01 91.521488
 1991-06-01 92.685966
 1991-07-01 93.850443
 1991-08-01 95.014921
 1991-09-01 96.179398
 1991-10-01 97.343876
 1991-11-01 98.508353
 1991-12-01 99.672831
 1992-01-01 100.837308
 1992-02-01 102.001785
 1992-03-01 103.166263
 1992-04-01 104.330740
 1992-05-01 105.495218
 1992-06-01 106.659695
 1992-07-01 107.824173
 1992-08-01 108.988650
 1992-09-01 110.153127
 1992-10-01 111.317605
 1992-11-01 112.482082
 1992-12-01 113.646560
 1993-01-01 114.811037
 1993-02-01 115.975515
 1993-03-01 117.139992
 1993-04-01 118.304469
 1993-05-01 119.468947
 1993-06-01 120.633424
 1993-07-01 121.797902
 1993-08-01 122.962379
```

|            |            |
|------------|------------|
| 1993-09-01 | 124.126857 |
| 1993-10-01 | 125.291334 |
| 1993-11-01 | 126.455811 |
| 1993-12-01 | 127.620289 |
| 1994-01-01 | 128.784766 |
| 1994-02-01 | 129.949244 |
| 1994-03-01 | 131.113721 |
| 1994-04-01 | 132.278199 |
| 1994-05-01 | 133.442676 |
| 1994-06-01 | 134.607153 |
| 1994-07-01 | 135.771631 |
| 1994-08-01 | 136.936108 |
| 1994-09-01 | 138.100586 |
| 1994-10-01 | 139.265063 |
| 1994-11-01 | 140.429541 |
| 1994-12-01 | 141.594018 |
| 1995-01-01 | 142.758495 |
| 1995-02-01 | 143.922973 |
| 1995-03-01 | 145.087450 |
| 1995-04-01 | 146.251928 |
| 1995-05-01 | 147.416405 |
| 1995-06-01 | 148.580883 |
| 1995-07-01 | 149.745360 |

Freq: MS, dtype: float64

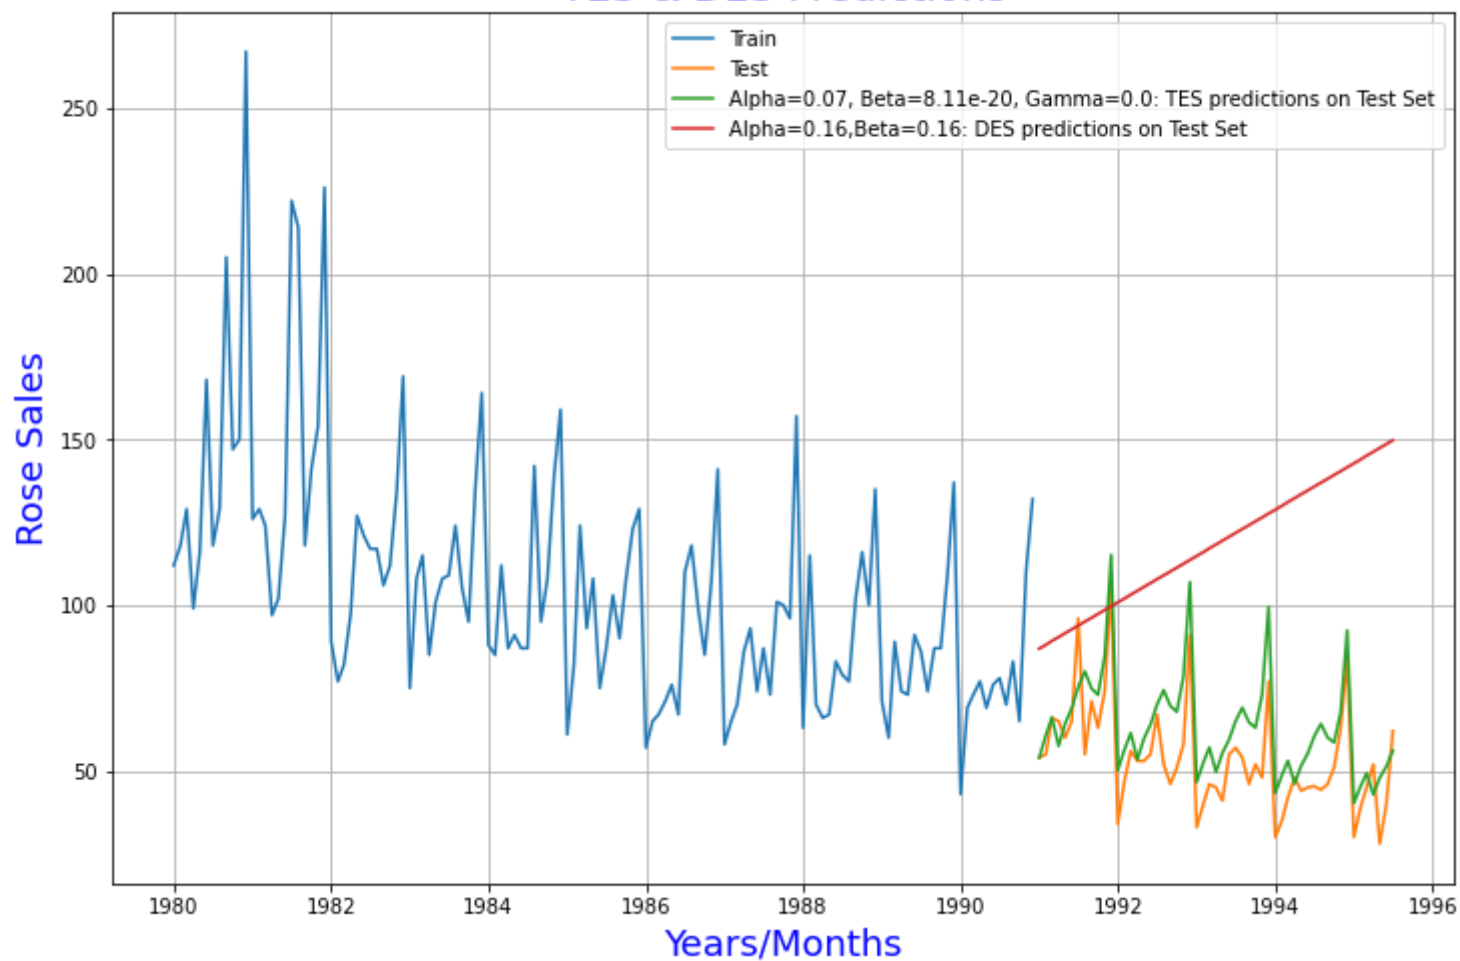
In [44]: *# Plotting the Training data, Test data and the forecasted values:*

```
plt.plot(train, label='Train')
plt.plot(test, label='Test')

plt.plot(TES_predict, label='Alpha=0.07, Beta=8.11e-20, Gamma=0.0: TES predictions on Test Set')
plt.plot(DES_predict, label='Alpha=0.16,Beta=0.16: DES predictions on Test Set')

plt.legend(loc='best')
plt.grid()
plt.title('TES & DES Predictions',color='blue',fontsize=20);
plt.xlabel('Years/Months',color='blue',fontsize=18);
plt.ylabel('Rose Sales',color='blue',fontsize=18);
```

## TES & DES Predictions



```
In [45]: #Evaluating the DES model:
```

```
print('DES RMSE:', mean_squared_error(test.values, DES_predict.values, squared=False))
```

```
DES RMSE: 70.60459768443859
```

In [46]: *# Storing results to a dataframe:*

```
results_smoothing_2 = pd.DataFrame({'Test RMSE': [mean_squared_error(test.values, DES_predict.values, squared_loss=loss)],
 ,index=['DES: Alpha=0.16,Beta=0.16']})
results = pd.concat([results, results_smoothing_2])
results
```

Out[46]:

|                                           | Test RMSE |
|-------------------------------------------|-----------|
| TES: Alpha=0.07, Beta=8.11e-20, Gamma=0.0 | 12.831106 |
| DES: Alpha=0.16,Beta=0.16                 | 70.604598 |

In [47]: *# Using the Linear Regression model for forecasting:*

*# Modifying the data to incorporate order against the sales values:*

```
train_time = [i+1 for i in range(len(train))]
test_time = [i+133 for i in range(len(test))]
print('Training Time instance', '\n', train_time)
print('Test Time instance', '\n', test_time)
```

Training Time instance

[1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75, 76, 77, 78, 79, 80, 81, 82, 83, 84, 85, 86, 87, 88, 89, 90, 91, 92, 93, 94, 95, 96, 97, 98, 99, 100, 101, 102, 103, 104, 105, 106, 107, 108, 109, 110, 111, 112, 113, 114, 115, 116, 117, 118, 119, 120, 121, 122, 123, 124, 125, 126, 127, 128, 129, 130, 131, 132]

Test Time instance

[133, 134, 135, 136, 137, 138, 139, 140, 141, 142, 143, 144, 145, 146, 147, 148, 149, 150, 151, 152, 153, 154, 155, 156, 157, 158, 159, 160, 161, 162, 163, 164, 165, 166, 167, 168, 169, 170, 171, 172, 173, 174, 175, 176, 177, 178, 179, 180, 181, 182, 183, 184, 185, 186, 187]

In [48]: *# Working on copies of Train & test data:*

```
LinearRegression_train = train.copy()
LinearRegression_test = test.copy()
```

In [49]: *#Cross-checking the data:*

```
LinearRegression_train['time'] = train_time
LinearRegression_test['time'] = test_time

print('First few rows of Training Data')
display(LinearRegression_train.head())
print('Last few rows of Training Data')
display(LinearRegression_train.tail())
print('First few rows of Test Data')
display(LinearRegression_test.head())
print('Last few rows of Test Data')
display(LinearRegression_test.tail())
```

First few rows of Training Data

|            | Rose  | time |
|------------|-------|------|
| YearMonth  |       |      |
| 1980-01-01 | 112.0 | 1    |
| 1980-02-01 | 118.0 | 2    |
| 1980-03-01 | 129.0 | 3    |
| 1980-04-01 | 99.0  | 4    |
| 1980-05-01 | 116.0 | 5    |

Last few rows of Training Data

|            | Rose | time |
|------------|------|------|
| YearMonth  |      |      |
| 1990-08-01 | 70.0 | 128  |
| 1990-09-01 | 83.0 | 129  |
| 1990-10-01 | 65.0 | 130  |

| Rose time  |       |     |
|------------|-------|-----|
| YearMonth  |       |     |
| <hr/>      |       |     |
| 1990-11-01 | 110.0 | 131 |
| 1990-12-01 | 132.0 | 132 |

First few rows of Test Data

| Rose time  |      |     |
|------------|------|-----|
| YearMonth  |      |     |
| <hr/>      |      |     |
| 1991-01-01 | 54.0 | 133 |
| 1991-02-01 | 55.0 | 134 |
| 1991-03-01 | 66.0 | 135 |
| 1991-04-01 | 65.0 | 136 |
| 1991-05-01 | 60.0 | 137 |

Last few rows of Test Data

| Rose time  |      |     |
|------------|------|-----|
| YearMonth  |      |     |
| <hr/>      |      |     |
| 1995-03-01 | 45.0 | 183 |
| 1995-04-01 | 52.0 | 184 |
| 1995-05-01 | 28.0 | 185 |
| 1995-06-01 | 40.0 | 186 |
| 1995-07-01 | 62.0 | 187 |



In [50]: *#Building the LR model:*

```
from sklearn.linear_model import LinearRegression

lr = LinearRegression()
lr.fit(LinearRegression_train[['time']],LinearRegression_train['Rose'])
```

Out[50]: LinearRegression()

In [51]: *#Predicting values:*

```
train_predictions_lr = lr.predict(LinearRegression_train[['time']])
LinearRegression_train['LR_on_time'] = train_predictions_lr


test_predictions_lr = lr.predict(LinearRegression_test[['time']])
LinearRegression_test['LR_on_time'] = test_predictions_lr

plt.plot(train['Rose'], label='Train')
plt.plot(test['Rose'], label='Test')
plt.plot(TES_predict, label='Alpha=0.07, Beta=8.11e-20, Gamma=0.0: TES predictions on Test Set')
plt.plot(DES_predict, label='Alpha=0.16, Beta=0.16: DES predictions on Test Set')

plt.plot(LinearRegression_test['LR_on_time'], label='Regression On Time_Test Data')

plt.legend(loc='best')
plt.grid();
plt.title('TES, DES & LR Predictions', color='blue', fontsize=20);
plt.xlabel('Years/Months', color='blue', fontsize=18);
plt.ylabel('Rose Sales', color='blue', fontsize=18);
```

## TES, DES & LR Predictions



```
In [52]: # Evaluating the model:
```

```
print('LR RMSE:', mean_squared_error(test['Rose'], test_predictions_lr, squared=False))
```

```
LR RMSE: 15.278368802792677
```

```
In [53]: results_smoothing_3 = pd.DataFrame({'Test RMSE': [mean_squared_error(test['Rose'], test_predictions_lr, scores=test_predictions_lr, index=['LR RMSE'])],
 index=['LR RMSE'])

results = pd.concat([results, results_smoothing_3])
results
```

Out[53]:

|                                           | Test RMSE |
|-------------------------------------------|-----------|
| TES: Alpha=0.07, Beta=8.11e-20, Gamma=0.0 | 12.831106 |
| DES: Alpha=0.16, Beta=0.16                | 70.604598 |
| LR RMSE                                   | 15.278369 |

```
In [54]: # Using the Naive Approach for forecasting:

Working on copies of Train & test data:

Naive_train = train.copy()
Naive_test = test.copy()
```

```
In [55]: train.head()
```

Out[55]:

|            | Rose  |
|------------|-------|
| YearMonth  |       |
| 1980-01-01 | 112.0 |
| 1980-02-01 | 118.0 |
| 1980-03-01 | 129.0 |
| 1980-04-01 | 99.0  |
| 1980-05-01 | 116.0 |

```
In [56]: test.head()
```

```
Out[56]:
```

|            | Rose |
|------------|------|
| YearMonth  |      |
| 1991-01-01 | 54.0 |
| 1991-02-01 | 55.0 |
| 1991-03-01 | 66.0 |
| 1991-04-01 | 65.0 |
| 1991-05-01 | 60.0 |

```
In [57]: Naive_test['naive'] = np.asarray(train['Rose'])[len(np.asarray(train['Rose']))-1]
Naive_test['naive'].head()
```


```
Out[57]: YearMonth
1991-01-01 132.0
1991-02-01 132.0
1991-03-01 132.0
1991-04-01 132.0
1991-05-01 132.0
Name: naive, dtype: float64
```

```
In [58]: plt.plot(Naive_train['Rose'], label='Train')
plt.plot(test['Rose'], label='Test')

plt.plot(TES_predict, label='Alpha=0.07, Beta=8.11e-20, Gamma=0.0: TES predictions on Test Set')
plt.plot(DES_predict, label='Alpha=0.16, Beta=0.16: DES predictions on Test Set')
plt.plot(LinearRegression_test['LR_on_time'], label='Regression On Time_Test Data')

plt.plot(Naive_test['naive'], label='Naive Forecast on Test Data')

plt.legend(loc='best')
plt.title("Various Model Forecasts")
plt.grid();
```



In [59]: *# Evaluating the model:*

```
print('Naive RMSE:', mean_squared_error(test['Rose'], Naive_test['naive'], squared=False))
```

Naive RMSE: 79.74569745398024

```
In [60]: results_smoothing_4 = pd.DataFrame({'Test RMSE': [mean_squared_error(test['Rose'], Naive_test['naive'], squared=False),
], index=['Naive RMSE']})

results = pd.concat([results, results_smoothing_4])
results
```

Out[60]:

|                                                  | Test RMSE |
|--------------------------------------------------|-----------|
| <b>TES: Alpha=0.07, Beta=8.11e-20, Gamma=0.0</b> | 12.831106 |
| <b>DES: Alpha=0.16, Beta=0.16</b>                | 70.604598 |
| <b>LR RMSE</b>                                   | 15.278369 |
| <b>Naive RMSE</b>                                | 79.745697 |

In [61]: *# Using the Simple Average method:*

```
Working on copies of Train & test data:
```

```
SA_train = train.copy()
SA_test = test.copy()
```

```
In [62]: SA_test['mean_forecast'] = train['Rose'].mean()
SA_test.head()
```

Out[62]:

|            | Rose | mean_forecast |
|------------|------|---------------|
| YearMonth  |      |               |
| 1991-01-01 | 54.0 | 104.939394    |
| 1991-02-01 | 55.0 | 104.939394    |
| 1991-03-01 | 66.0 | 104.939394    |
| 1991-04-01 | 65.0 | 104.939394    |
| 1991-05-01 | 60.0 | 104.939394    |



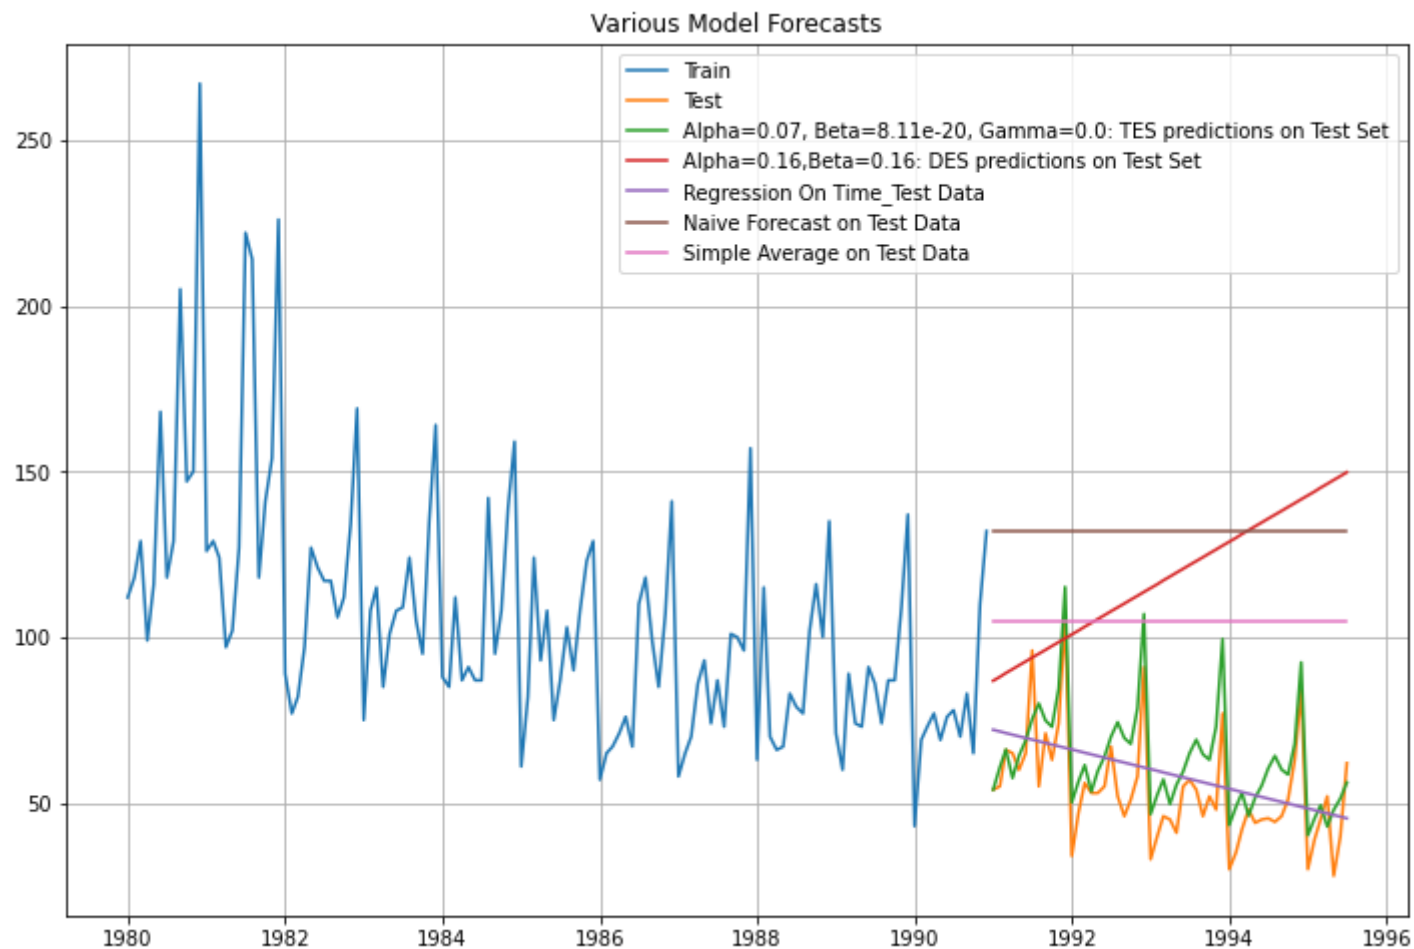
In [63]: *# Plotting predictions of various models:*

```
plt.plot(SA_train['Rose'], label='Train')
plt.plot(SA_test['Rose'], label='Test')

plt.plot(TES_predict, label='Alpha=0.07, Beta=8.11e-20, Gamma=0.0: TES predictions on Test Set')
plt.plot(DES_predict, label='Alpha=0.16,Beta=0.16: DES predictions on Test Set')
plt.plot(LinearRegression_test['LR_on_time'], label='Regression On Time_Test Data')
plt.plot(Naive_test['naive'], label='Naive Forecast on Test Data')

plt.plot(SA_test['mean_forecast'], label='Simple Average on Test Data')

plt.legend(loc='best')
plt.title("Various Model Forecasts")
plt.grid();
```



```
In [64]: # Evaluating the model:
```

```
print('SA RMSE:', mean_squared_error(test['Rose'], SA_test['mean_forecast'], squared=False))
```

```
SA RMSE: 53.48823282885678
```

```
In [65]: results_smoothing_5 = pd.DataFrame({'Test RMSE': [mean_squared_error(test['Rose'],SA_test['mean_forecast'],index=['SA RMSE'])

results = pd.concat([results, results_smoothing_5])
results
```

Out[65]:

|                                                  | Test RMSE |
|--------------------------------------------------|-----------|
| <b>TES: Alpha=0.07, Beta=8.11e-20, Gamma=0.0</b> | 12.831106 |
| <b>DES: Alpha=0.16,Beta=0.16</b>                 | 70.604598 |
| <b>LR RMSE</b>                                   | 15.278369 |
| <b>Naive RMSE</b>                                | 79.745697 |
| <b>SA RMSE</b>                                   | 53.488233 |

```
In [66]: #Using the Moving Average method on a copy of the original data:
```

```
MA = data.copy()
MA.head()
```

Out[66]:

|                   | Rose  |
|-------------------|-------|
| YearMonth         |       |
| <b>1980-01-01</b> | 112.0 |
| <b>1980-02-01</b> | 118.0 |
| <b>1980-03-01</b> | 129.0 |
| <b>1980-04-01</b> | 99.0  |
| <b>1980-05-01</b> | 116.0 |

In [67]: *# Using various parameters:*

```
MA['Trailing_2'] = MA['Rose'].rolling(2).mean()
MA['Trailing_4'] = MA['Rose'].rolling(4).mean()
MA['Trailing_6'] = MA['Rose'].rolling(6).mean()
MA['Trailing_9'] = MA['Rose'].rolling(9).mean()

MA.head(10)
```

Out[67]:


|            | Rose  | Trailing_2 | Trailing_4 | Trailing_6 | Trailing_9 |
|------------|-------|------------|------------|------------|------------|
| YearMonth  |       |            |            |            |            |
| 1980-01-01 | 112.0 | NaN        | NaN        | NaN        | NaN        |
| 1980-02-01 | 118.0 | 115.0      | NaN        | NaN        | NaN        |
| 1980-03-01 | 129.0 | 123.5      | NaN        | NaN        | NaN        |
| 1980-04-01 | 99.0  | 114.0      | 114.50     | NaN        | NaN        |
| 1980-05-01 | 116.0 | 107.5      | 115.50     | NaN        | NaN        |
| 1980-06-01 | 168.0 | 142.0      | 128.00     | 123.666667 | NaN        |
| 1980-07-01 | 118.0 | 143.0      | 125.25     | 124.666667 | NaN        |
| 1980-08-01 | 129.0 | 123.5      | 132.75     | 126.500000 | NaN        |
| 1980-09-01 | 205.0 | 167.0      | 155.00     | 139.166667 | 132.666667 |
| 1980-10-01 | 147.0 | 176.0      | 149.75     | 147.166667 | 136.555556 |

In [68]: *# Plotting on the entire data:*

```
plt.plot(MA['Rose'], label='Train')
plt.plot(MA['Trailing_2'], label='2 Point Moving Average')
plt.plot(MA['Trailing_4'], label='4 Point Moving Average')
plt.plot(MA['Trailing_6'], label='6 Point Moving Average')
plt.plot(MA['Trailing_9'], label='9 Point Moving Average')

plt.legend(loc = 'best')
plt.grid();
plt.title('Various MA Forecasts on entire data',color='blue',fontsize=20);
plt.xlabel('Years / Months',color='blue',fontsize=18);
plt.ylabel('Rose Sales',color='blue',fontsize=18);
```

## Various MA Forecasts on entire data



```
In [69]: #Creating train and test set for MA method:
```

```
trailing_MA_train = MA[MA.index<'1991']
trailing_MA_test= MA[MA.index>='1991']
```

```
In [70]: trailing_MA_train.tail()
```

```
Out[70]:
```

|            | Rose  | Trailing_2 | Trailing_4 | Trailing_6 | Trailing_9 |
|------------|-------|------------|------------|------------|------------|
| YearMonth  |       |            |            |            |            |
| 1990-08-01 | 70.0  | 74.0       | 73.25      | 73.833333  | 76.888889  |
| 1990-09-01 | 83.0  | 76.5       | 76.75      | 75.500000  | 70.888889  |
| 1990-10-01 | 65.0  | 74.0       | 74.00      | 73.500000  | 73.333333  |
| 1990-11-01 | 110.0 | 87.5       | 82.00      | 80.333333  | 77.888889  |
| 1990-12-01 | 132.0 | 121.0      | 97.50      | 89.666667  | 84.444444  |

```
In [71]: trailing_MA_test.head()
```

```
Out[71]:
```

|            | Rose | Trailing_2 | Trailing_4 | Trailing_6 | Trailing_9 |
|------------|------|------------|------------|------------|------------|
| YearMonth  |      |            |            |            |            |
| 1991-01-01 | 54.0 | 93.0       | 90.25      | 85.666667  | 81.888889  |
| 1991-02-01 | 55.0 | 54.5       | 87.75      | 83.166667  | 80.333333  |
| 1991-03-01 | 66.0 | 60.5       | 76.75      | 80.333333  | 79.222222  |
| 1991-04-01 | 65.0 | 65.5       | 60.00      | 80.333333  | 77.777778  |
| 1991-05-01 | 60.0 | 62.5       | 61.50      | 72.000000  | 76.666667  |

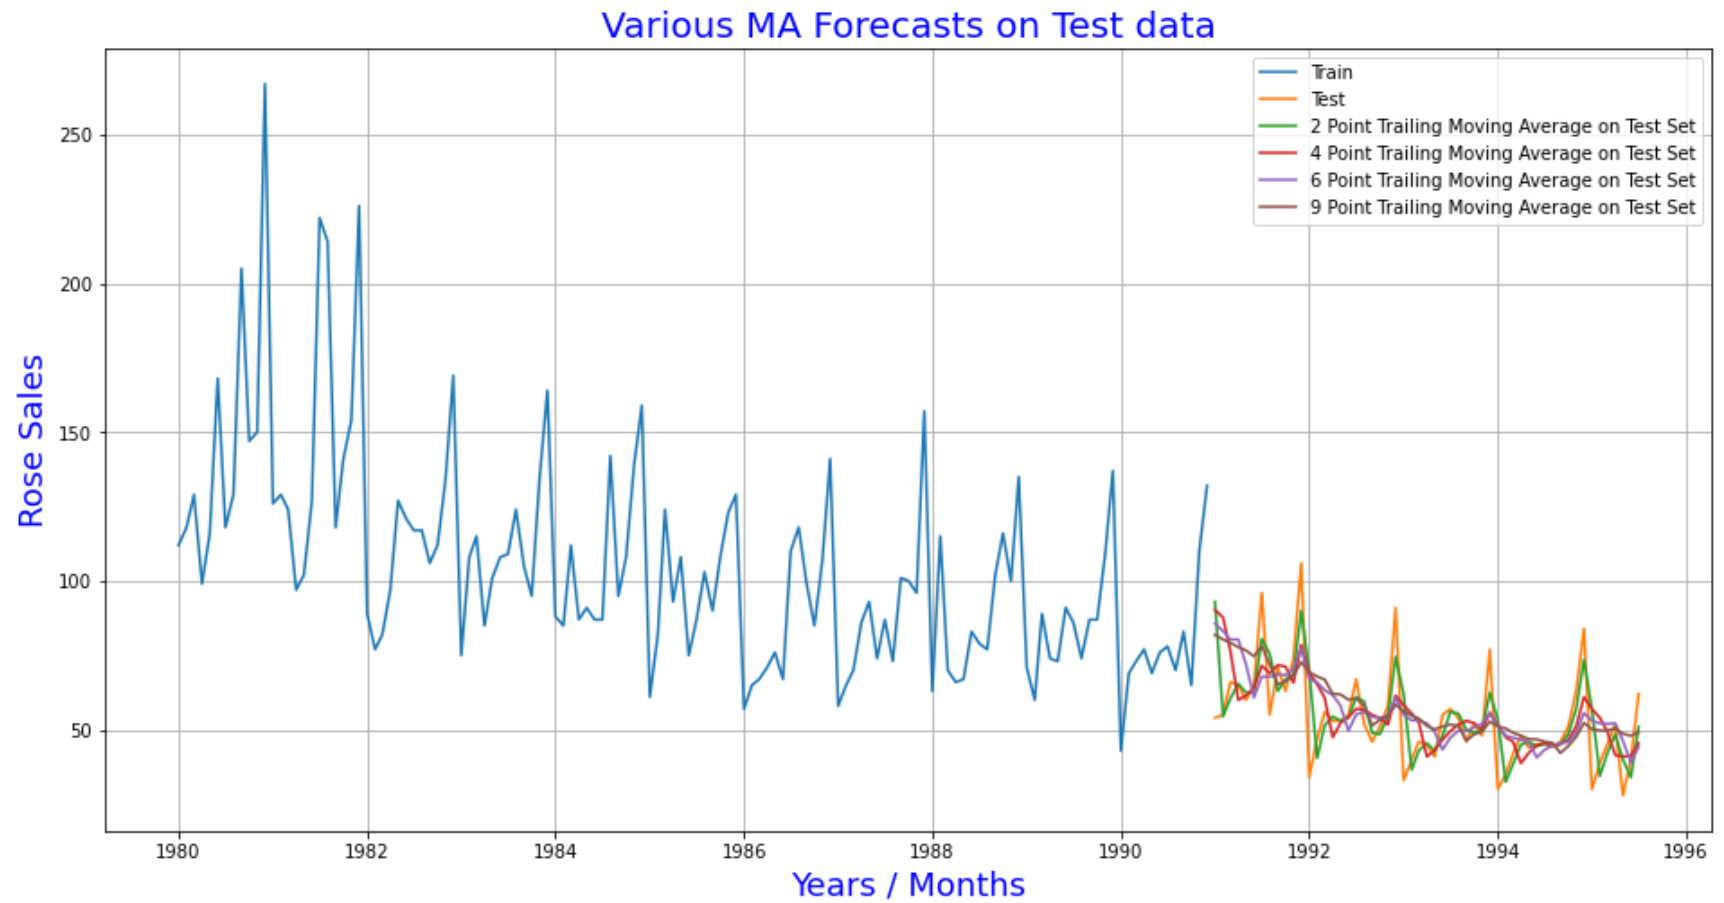
In [72]: *# Plotting on Test data:*

```
plt.figure(figsize=(16,8))
plt.plot(trailing_MA_train['Rose'], label='Train')
plt.plot(trailing_MA_test['Rose'], label='Test')

plt.plot(trailing_MA_test['Trailing_2'],label = '2 Point Trailing Moving Average on Test Set')
plt.plot(trailing_MA_test['Trailing_4'],label = '4 Point Trailing Moving Average on Test Set')
plt.plot(trailing_MA_test['Trailing_6'],label = '6 Point Trailing Moving Average on Test Set')
plt.plot(trailing_MA_test['Trailing_9'],label = '9 Point Trailing Moving Average on Test Set')

plt.legend(loc = 'best')
plt.grid();
plt.title('Various MA Forecasts on Test data',color='blue',fontsize=20);
plt.xlabel('Years / Months',color='blue',fontsize=18);
plt.ylabel('Rose Sales',color='blue',fontsize=18);
```





In [73]: *# Evaluating using RSME:*

```
from sklearn import metrics
```

```
2 point Trailing RSME
```

```
rmse_test_2 = metrics.mean_squared_error(test['Rose'],trailing_MA_test['Trailing_2'],squared=False)
print("For 2 point MA Model, RMSE is %3.3f" %(rmse_test_2))
```

```
4 point Trailing RSME
```

```
rmse_test_4 = metrics.mean_squared_error(test['Rose'],trailing_MA_test['Trailing_4'],squared=False)
print("For 4 point MA Model, RMSE is %3.3f" %(rmse_test_4))
```

```
6 point Trailing RSME
```

```
rmse_test_6 = metrics.mean_squared_error(test['Rose'],trailing_MA_test['Trailing_6'],squared=False)
print("For 6 point MA Model, RMSE is %3.3f" %(rmse_test_6))
```

```
9 point Trailing RSME
```

```
rmse_test_9 = metrics.mean_squared_error(test['Rose'],trailing_MA_test['Trailing_9'],squared=False)
print("For 9 point MA Model, RMSE is %3.3f" %(rmse_test_9))
```

For 2 point MA Model, RMSE is 11.530

For 4 point MA Model, RMSE is 14.458

For 6 point MA Model, RMSE is 14.573

For 9 point MA Model, RMSE is 14.733

In [74]: *# Storing the results:*

```
results_smoothing_6 = pd.DataFrame({'Test RMSE': [rmse_test_2,rmse_test_4
 ,rmse_test_6,rmse_test_9]})
 ,index=['2-point MA','4-point MA','6-point MA','9-point MA'])

results = pd.concat([results, results_smoothing_6])
results
```

Out[74]:

|                                                  | Test RMSE |
|--------------------------------------------------|-----------|
| <b>TES: Alpha=0.07, Beta=8.11e-20, Gamma=0.0</b> | 12.831106 |
| <b>DES: Alpha=0.16,Beta=0.16</b>                 | 70.604598 |
| <b>LR RMSE</b>                                   | 15.278369 |
| <b>Naive RMSE</b>                                | 79.745697 |
| <b>SA RMSE</b>                                   | 53.488233 |
| <b>2-point MA</b>                                | 11.530054 |
| <b>4-point MA</b>                                | 14.458402 |
| <b>6-point MA</b>                                | 14.572976 |
| <b>9-point MA</b>                                | 14.732918 |

In [75]: *# Plotting the comparison of all model predictions:*

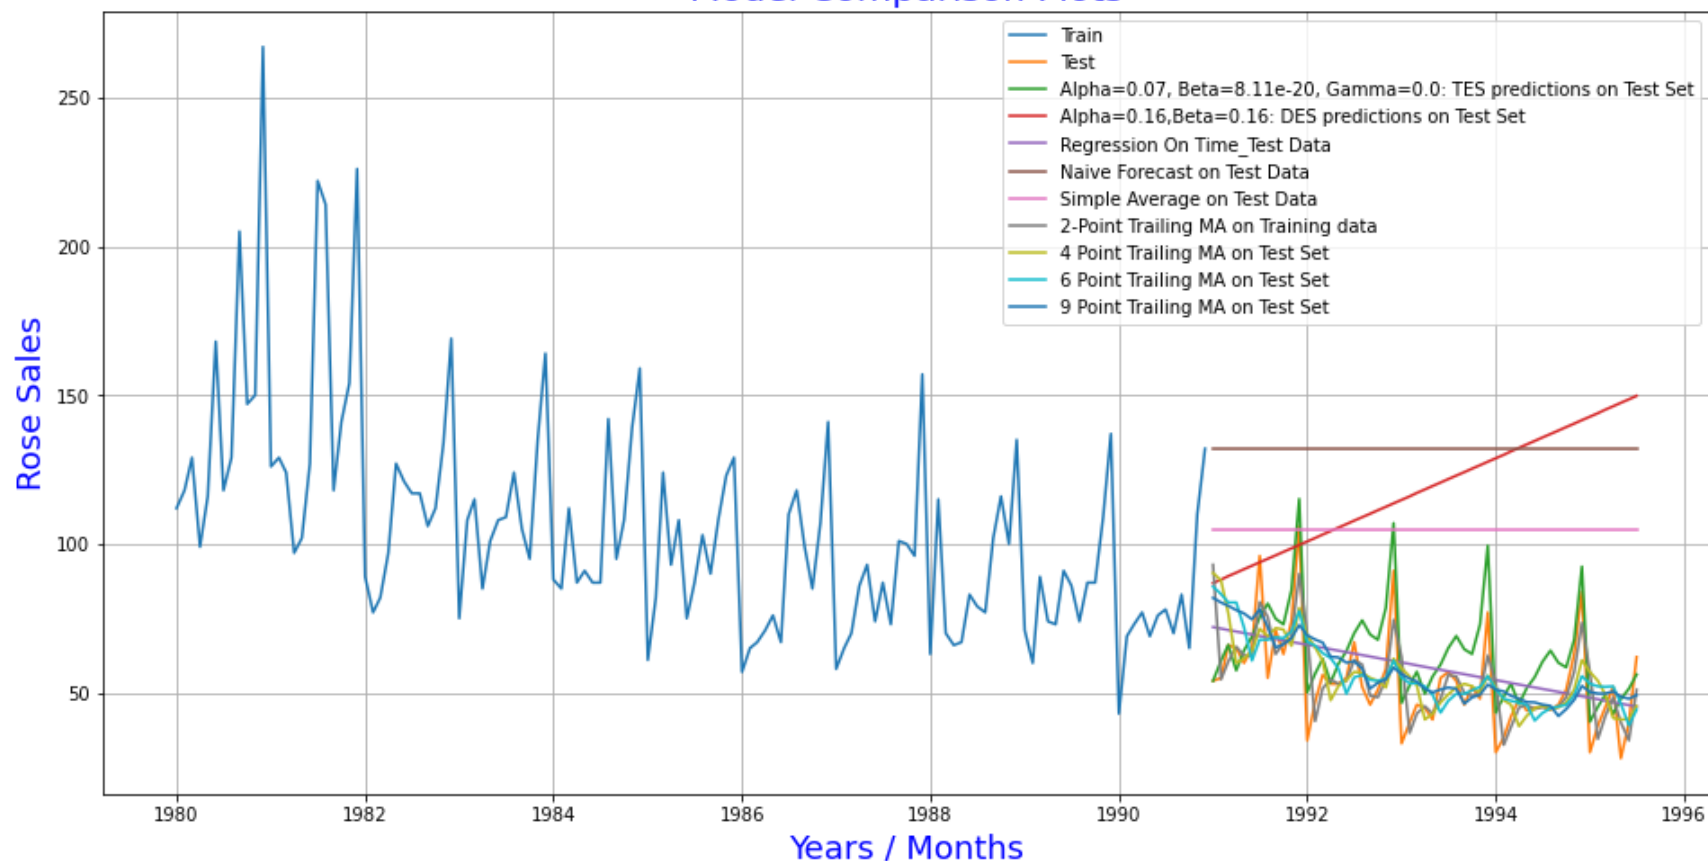
```
plt.figure(figsize=(16,8))
plt.plot(train['Rose'], label='Train')
plt.plot(test['Rose'], label='Test')

plt.plot(TES_predict, label='Alpha=0.07, Beta=8.11e-20, Gamma=0.0: TES predictions on Test Set')
plt.plot(DEs_predict, label='Alpha=0.16,Beta=0.16: DES predictions on Test Set')
plt.plot(LinearRegression_test['LR_on_time'], label='Regression On Time_Test Data')
plt.plot(Naive_test['naive'], label='Naive Forecast on Test Data')
plt.plot(SA_test['mean_forecast'], label='Simple Average on Test Data')

plt.plot(trailing_MA_test['Trailing_2'],label = '2-Point Trailing MA on Training data')
plt.plot(trailing_MA_test['Trailing_4'],label = '4 Point Trailing MA on Test Set')
plt.plot(trailing_MA_test['Trailing_6'],label = '6 Point Trailing MA on Test Set')
plt.plot(trailing_MA_test['Trailing_9'],label = '9 Point Trailing MA on Test Set')

plt.legend(loc='best')
plt.grid();
plt.title('Model Comparison Plots',color='blue',fontsize=20);
plt.xlabel('Years / Months',color='blue',fontsize=18);
plt.ylabel('Rose Sales',color='blue',fontsize=18);
```

## Model Comparison Plots



In [76]: *# Checking for stationarity of data using ADF test:*

```
from statsmodels.tsa.stattools import adfuller
```

```
In [77]: stat_test = adfuller(data, regression='ct')
print('Test statistic is %3.3f' % stat_test[0])
print('Test p-value is' , stat_test[1])
print('Number of lags used' , stat_test[2])
```

```
Test statistic is -2.242
Test p-value is 0.46628917681591
Number of lags used 13
```

In [78]: *# P-value > alpha, so data is non-stationary. Using level-1 differencing to make data stationary:*

```
stat_test = adfuller(data.diff().dropna(), regression='ct')
print('Test statistic is %3.3f' % stat_test[0])
print('Test p-value is' , stat_test[1])
print('Number of lags used' , stat_test[2])
```

Test statistic is -8.161


Test p-value is 3.037926501732391e-11

Number of lags used 12

In [79]: *# Now data is stationary. Plotting the differenced data:*

```
data.diff().dropna().plot(grid=True);
plt.title('Rose Sales Distribution (Post Differencing)',color='blue',fontsize=16)
plt.xlabel('Years/Months',color='blue',fontsize=14)
plt.ylabel('Rose Sales',color='blue',fontsize=14)
```

Out[79]: Text(0, 0.5, 'Rose Sales')








```
In [80]: # Plotting the autocorrelation and partial autocorrelation plots on data:


from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

plot_acf(data,alpha=0.05);
```





```
In [81]: plot_pacf(data, zero=False, alpha=0.05);
```



```
In [82]: #Splitting data to build the models:
train = data[data.index<'1991']
test = data[data.index>='1991']
```


```
In [83]: test.shape
```

```
Out[83]: (55, 1)
```

In [84]: *#Plotting the train data:*

```
train.plot(grid=True);
plt.title('Train data',color='blue',fontsize=16)
plt.xlabel('Years/Months',color='blue',fontsize=14)
plt.ylabel('Rose Sales',color='blue',fontsize=14)
```

Out[84]: Text(0, 0.5, 'Rose Sales')



In [85]: *#Checking for stationarity of the train data:*

```
stat_test = adfuller(train,regression='ct')
print('Test statistic is %3.3f' %stat_test[0])
print('Test p-value is' ,stat_test[1])
print('Number of lags used' ,stat_test[2])
```

```
Test statistic is -1.686
Test p-value is 0.7569093051047098
Number of lags used 13
```

In [86]: *# P-value > alpha, so data is non-stationary. Using level-1 differencing to make data stationary:*


```
stat_test = adfuller(train.diff().dropna(),regression='ct')
print('Test statistic is %3.3f' %stat_test[0])
print('Test p-value is' ,stat_test[1])
print('Number of lags used' ,stat_test[2])
```

```
Test statistic is -6.804
Test p-value is 3.8948313567828775e-08
Number of lags used 12
```

In [87]: *#Plotting the differenced train data:*

```
train.diff().dropna().plot(grid=True);
plt.title('Train data (Post Differencing)',color='blue',fontsize=16)
plt.xlabel('Years/Months',color='blue',fontsize=14)
plt.ylabel('Rose Sales',color='blue',fontsize=14)
```

Out[87]: Text(0, 0.5, 'Rose Sales')



```
In [88]: # Since there is seasonality in the data set, we will build SARIMA model:
 ## Building automated version of SARIMA:
```

```
In [89]: import itertools
p = q = range(0, 4)
d= range(1,2)
D = range(0,1)
pdq = list(itertools.product(p, d, q))
PDQ = [(x[0], x[1], x[2], 12) for x in list(itertools.product(p, D, q))]
print('Examples of the parameter combinations for the Model are')
for i in range(1,len(pdq)):
 print('Model: {}'.format(pdq[i], PDQ[i]))
```

Examples of the parameter combinations for the Model are

```
Model: (0, 1, 1)(0, 0, 1, 12)
Model: (0, 1, 2)(0, 0, 2, 12)
Model: (0, 1, 3)(0, 0, 3, 12)
Model: (1, 1, 0)(1, 0, 0, 12)
Model: (1, 1, 1)(1, 0, 1, 12)
Model: (1, 1, 2)(1, 0, 2, 12)
Model: (1, 1, 3)(1, 0, 3, 12)
Model: (2, 1, 0)(2, 0, 0, 12)
Model: (2, 1, 1)(2, 0, 1, 12)
Model: (2, 1, 2)(2, 0, 2, 12)
Model: (2, 1, 3)(2, 0, 3, 12)
Model: (3, 1, 0)(3, 0, 0, 12)
Model: (3, 1, 1)(3, 0, 1, 12)
Model: (3, 1, 2)(3, 0, 2, 12)
Model: (3, 1, 3)(3, 0, 3, 12)
```

```
In [90]: #Creating dataframe for storing AIC values:
```

```
SARIMA_AIC = pd.DataFrame(columns=['param', 'seasonal', 'AIC'])
SARIMA_AIC
```

```
Out[90]:
```

| param | seasonal | AIC |
|-------|----------|-----|
|-------|----------|-----|



```
In [91]: statsmodels.api as sm

sm in pdq:
param_seasonal in PDQ:
SARIMA_model = sm.tsa.statespace.SARIMAX(train['Rose'].values,
 order=param,
 seasonal_order=param_seasonal,
 enforce_stationarity=False,
 enforce_invertibility=False)

results_SARIMA = SARIMA_model.fit(maxiter=1000)
print('SARIMA{x} - AIC:{}'.format(param, param_seasonal, results_SARIMA.aic))
SARIMA_AIC = SARIMA_AIC.append({'param':param, 'seasonal':param_seasonal, 'AIC': results_SARIMA.aic}, ignore_index=True)
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/base/model.py:567: ConvergenceWarning:

Maximum Likelihood optimization failed to converge. Check mle\_retvals

```
SARIMA(2, 1, 1)x(3, 0, 3, 12) - AIC:2493.0270444578505
SARIMA(2, 1, 2)x(0, 0, 0, 12) - AIC:1253.91021165659
SARIMA(2, 1, 2)x(0, 0, 1, 12) - AIC:1085.9643681360224
SARIMA(2, 1, 2)x(0, 0, 2, 12) - AIC:916.3259134369662
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/base/model.py:567: ConvergenceWarning:

Maximum Likelihood optimization failed to converge. Check mle\_retvals

```
SARIMA(2, 1, 2)x(0, 0, 3, 12) - AIC:2034.5662408936164
SARIMA(2, 1, 2)x(1, 0, 0, 12) - AIC:1073.2912711166598
SARIMA(2, 1, 2)x(1, 0, 1, 12) - AIC:1044.1909349817224
```

```
In [92]: # Arranging models in ascending order to select least-AIC model:
```

```
SARIMA_AIC.sort_values(by=['AIC']).head()
```

```
Out[92]:
```

|            | param     | seasonal      | AIC        |
|------------|-----------|---------------|------------|
| <b>222</b> | (3, 1, 1) | (3, 0, 2, 12) | 774.400286 |
| <b>238</b> | (3, 1, 2) | (3, 0, 2, 12) | 774.880941 |
| <b>220</b> | (3, 1, 1) | (3, 0, 0, 12) | 775.426699 |
| <b>221</b> | (3, 1, 1) | (3, 0, 1, 12) | 775.495330 |
| <b>252</b> | (3, 1, 3) | (3, 0, 0, 12) | 775.561019 |

In [93]: *#Using the least-AIC model for SARIMAX computation:*

```
import statsmodels.api as sm

auto_SARIMA = sm.tsa.statespace.SARIMAX(train['Rose'],
 order=(3, 1, 1),
 seasonal_order=(3, 0, 2, 12),
 enforce_stationarity=False,
 enforce_invertibility=False)
results_auto_SARIMA = auto_SARIMA.fit(maxiter=1000)
print(results_auto_SARIMA.summary())
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

#### SARIMAX Results

```
=====
Dep. Variable: Rose No. Observations: 132
Model: SARIMAX(3, 1, 1)x(3, 0, [1, 2], 12) Log Likelihood -377.200
Date: Sat, 09 Oct 2021 AIC 774.400
Time: 09:28:34 BIC 799.618
Sample: 01-01-1980 HQIC 784.578
 - 12-01-1990
Covariance Type: opg
=====
```

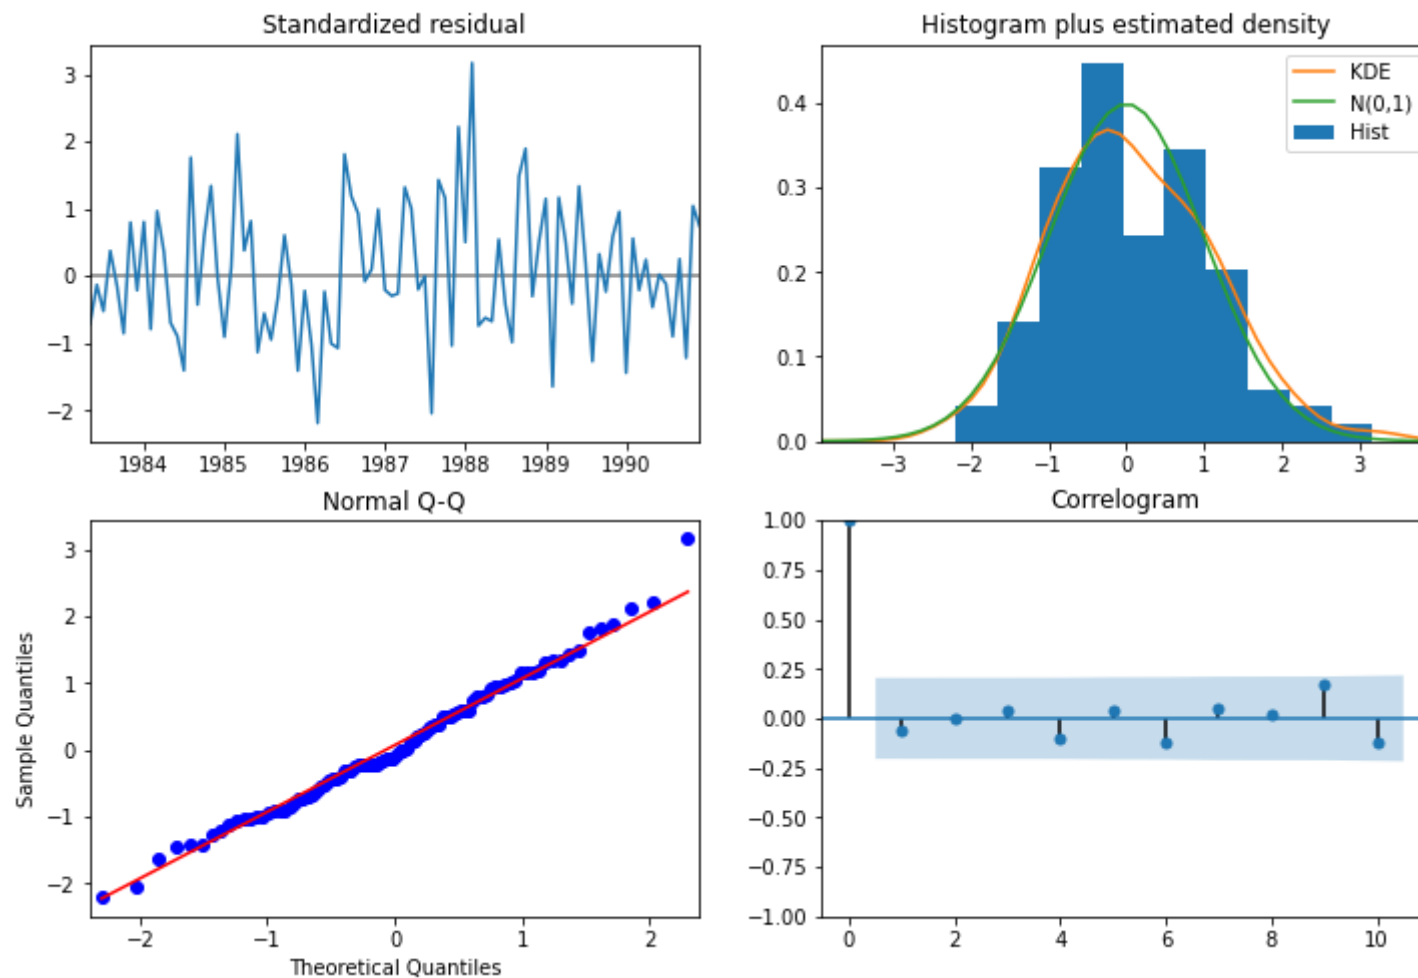
|          | coef    | std err | z       | P> z  | [0.025 | 0.975] |
|----------|---------|---------|---------|-------|--------|--------|
| ar.L1    | 0.0464  | 0.127   | 0.366   | 0.714 | -0.202 | 0.294  |
| ar.L2    | -0.0060 | 0.120   | -0.050  | 0.960 | -0.241 | 0.229  |
| ar.L3    | -0.1808 | 0.098   | -1.837  | 0.066 | -0.374 | 0.012  |
| ma.L1    | -0.9370 | 0.067   | -13.902 | 0.000 | -1.069 | -0.805 |
| ar.S.L12 | 0.7639  | 0.165   | 4.639   | 0.000 | 0.441  | 1.087  |
| ar.S.L24 | 0.0840  | 0.159   | 0.527   | 0.598 | -0.229 | 0.397  |
| ar.S.L36 | 0.0727  | 0.095   | 0.764   | 0.445 | -0.114 | 0.259  |
| ma.S.L12 | -0.4967 | 0.250   | -1.988  | 0.047 | -0.987 | -0.007 |

|                         |          |        |                   |       |         |         |
|-------------------------|----------|--------|-------------------|-------|---------|---------|
| ma.S.L24                | -0.2190  | 0.210  | -1.044            | 0.297 | -0.630  | 0.192   |
| sigma2                  | 192.1979 | 39.638 | 4.849             | 0.000 | 114.508 | 269.888 |
| =====                   |          |        |                   |       |         |         |
| Ljung-Box (Q):          |          | 34.22  | Jarque-Bera (JB): |       | 1.64    |         |
| Prob(Q):                |          | 0.73   | Prob(JB):         |       | 0.44    |         |
| Heteroskedasticity (H): |          | 1.11   | Skew:             |       | 0.33    |         |
| Prob(H) (two-sided):    |          | 0.78   | Kurtosis:         |       | 3.03    |         |
| =====                   |          |        |                   |       |         |         |

**Warnings:**

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [94]: results_auto_SARIMA.plot_diagnostics();
```



```
In [95]: # Using SARIMA model to predict test set:
```

```
predicted_auto_SARIMA = results_auto_SARIMA.get_forecast(steps=len(test))
```

```
In [96]: # Defining Mean Absolute Percentage Error (MAPE):

def mean_absolute_percentage_error(y_true, y_pred):
 return np.mean((np.abs(y_true-y_pred))/(y_true))*100

from sklearn.metrics import mean_squared_error
```

```
In [97]: # Evaluating the predictions:

rmse = mean_squared_error(test['Rose'],predicted_auto_SARIMA.predicted_mean,squared=False)
mape = mean_absolute_percentage_error(test['Rose'],predicted_auto_SARIMA.predicted_mean)
print('RMSE:',rmse,'\nMAPE:',mape)

RMSE: 18.91207676153986
MAPE: 36.45868367214725
```

```
In [98]: # Storing results for comparison:
results_models = pd.DataFrame({'Test RMSE': rmse,'MAPE':mape}
 ,index=['SARIMA(3,1,1)(3,0,2,12)'])


results_models
```

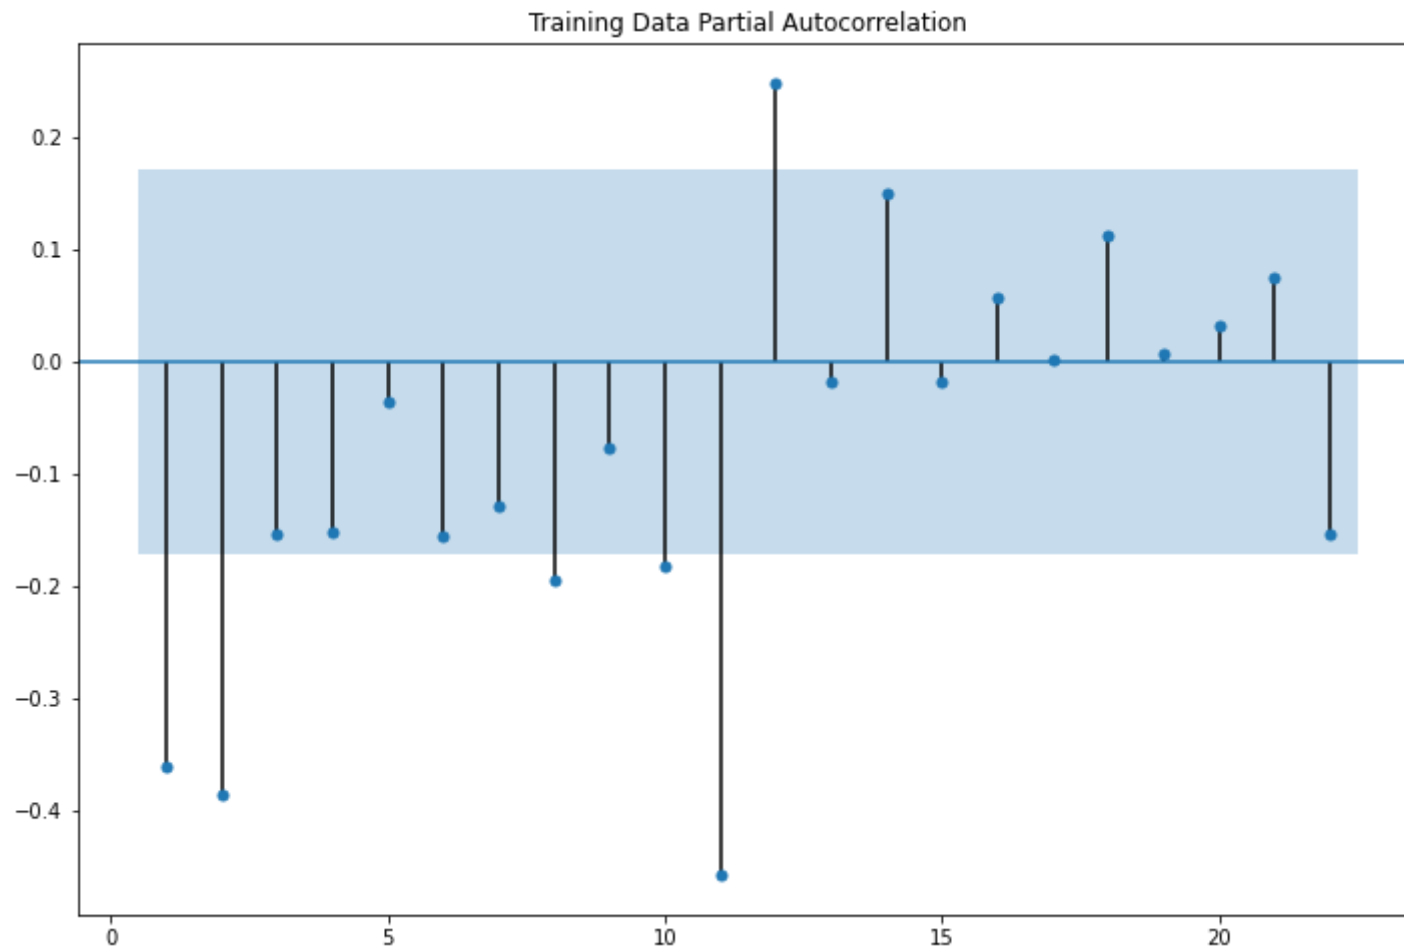
Out[98]:

|                                | Test RMSE | MAPE      |
|--------------------------------|-----------|-----------|
| <b>SARIMA(3,1,1)(3,0,2,12)</b> | 18.912077 | 36.458684 |

```
In [99]: # Building a manual SARIMA model by selecting values of p, q from correlation plots:

plot_acf(train.diff(),title='Training Data Autocorrelation',missing='drop')
plot_pacf(train.diff().dropna(),title='Training Data Partial Autocorrelation',zero=False)
plt.show()
```







```
In [100]: # As per the ACF and PACF plots, we will take the values as p=2, q=2, d=1 and seasonal componenets P=1,
```

```
In [101]: import statsmodels.api as sm
```

```
manual_SARIMA = sm.tsa.statespace.SARIMAX(train['Rose'],
 order=(2,1,2),
 seasonal_order=(1, 1, 1, 12),
 enforce_stationarity=False,
 enforce_invertibility=False)
results_manual_SARIMA = manual_SARIMA.fit(maxiter=1000)
print(results_manual_SARIMA.summary())
```

```
/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:
```

No frequency information was provided, so inferred frequency MS will be used.

```
/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:
```

No frequency information was provided, so inferred frequency MS will be used.

#### SARIMAX Results

```
=====
Dep. Variable: Rose No. Observations: 132
Model: SARIMAX(2, 1, 2)x(1, 1, [1], 12) Log Likelihood -450.847
Date: Sat, 09 Oct 2021 AIC 915.693
Time: 09:28:39 BIC 934.204
Sample: 01-01-1980 HQIC 923.193
 - 12-01-1990
```

```
Covariance Type: opg
```

```
=====
 coef std err z P>|z| [0.025 0.975]

ar.L1 1.1021 0.134 8.228 0.000 0.840 1.365
ar.L2 -0.3435 0.109 -3.141 0.002 -0.558 -0.129
ma.L1 -1.8137 0.106 -17.035 0.000 -2.022 -1.605
ma.L2 0.8654 0.095 9.133 0.000 0.680 1.051
ar.S.L12 -0.3879 0.069 -5.615 0.000 -0.523 -0.252
ma.S.L12 -0.0779 0.130 -0.598 0.550 -0.333 0.177
sigma2 338.2778 53.785 6.289 0.000 232.860 443.695
=====
```

```
Ljung-Box (Q): 25.41 Jarque-Bera (JB): 0.03
Prob(Q): 0.96 Prob(JB): 0.98
```

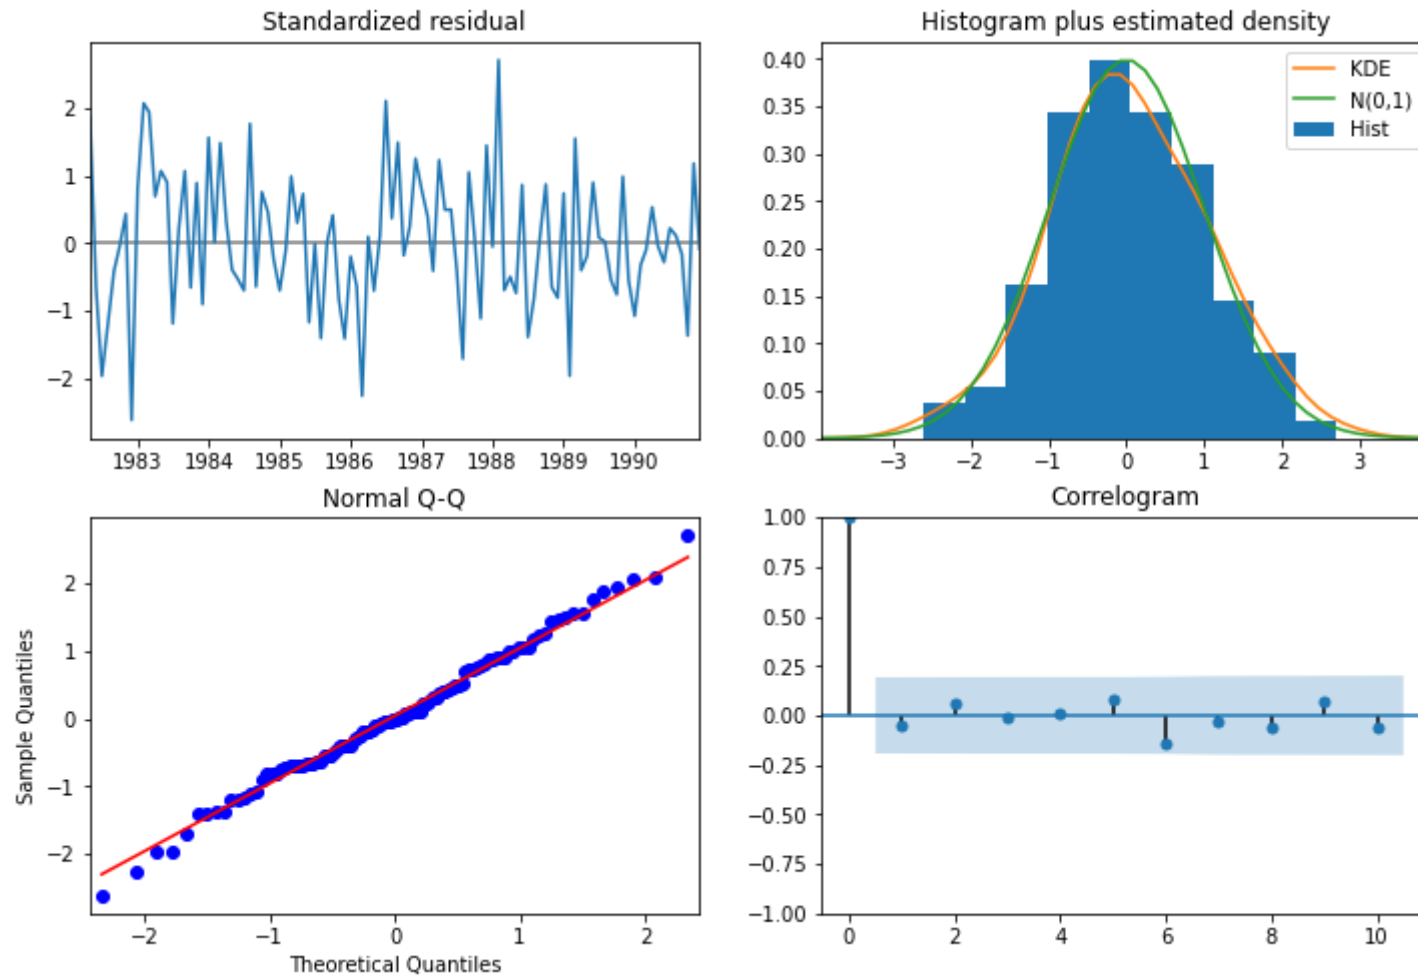
|                         |      |           |      |
|-------------------------|------|-----------|------|
| Heteroskedasticity (H): | 0.66 | Skew:     | 0.04 |
| Prob(H) (two-sided):    | 0.22 | Kurtosis: | 2.97 |

=====

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [102]: results_manual_SARIMA.plot_diagnostics()
plt.show()
```



In [103]: *# Forecasting using the new model:*

```
predicted_manual_SARIMA = results_manual_SARIMA.get_forecast(steps=len(test))
```

In [104]: *# Evaluating the model:*

```
rmse = mean_squared_error(test['Rose'], predicted_manual_SARIMA.predicted_mean, squared=False)
mape = mean_absolute_percentage_error(test['Rose'], predicted_manual_SARIMA.predicted_mean)
print('RMSE:', rmse, '\nMAPE:', mape)
```

RMSE: 13.278680854490426

MAPE: 17.133916945618907

In [105]: *# Storing the results:*

```
results_1 = pd.DataFrame({'Test RMSE': [rmse], 'MAPE': mape}
 , index=['SARIMA(2,1,2)(1,1,1,12)'])
results_models = pd.concat([results_models, results_1])
results_models
```

Out[105]:

|                                | Test RMSE | MAPE      |
|--------------------------------|-----------|-----------|
| <b>SARIMA(3,1,1)(3,0,2,12)</b> | 18.912077 | 36.458684 |
| <b>SARIMA(2,1,2)(1,1,1,12)</b> | 13.278681 | 17.133917 |

In [106]: *# Trying various parameters:*

```
manual_SARIMA = sm.tsa.statespace.SARIMAX(train['Rose'],
 order=(1,1,1),
 seasonal_order=(0, 0, 0, 12),
 enforce_stationarity=False,
 enforce_invertibility=False)
results_manual_SARIMA = manual_SARIMA.fit(maxiter=1000)
print(results_manual_SARIMA.summary())
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

#### SARIMAX Results

```
=====
Dep. Variable: Rose No. Observations: 132
Model: SARIMAX(1, 1, 1) Log Likelihood -628.092
Date: Sat, 09 Oct 2021 AIC 1262.184
Time: 09:28:40 BIC 1270.763
Sample: 01-01-1980 HQIC 1265.670
 - 12-01-1990
Covariance Type: opg
=====
```

|        | coef     | std err | z       | P> z  | [0.025  | 0.975]   |
|--------|----------|---------|---------|-------|---------|----------|
| ar.L1  | 0.1713   | 0.074   | 2.311   | 0.021 | 0.026   | 0.317    |
| ma.L1  | -0.9172  | 0.056   | -16.407 | 0.000 | -1.027  | -0.808   |
| sigma2 | 982.8469 | 95.174  | 10.327  | 0.000 | 796.310 | 1169.384 |

```
=====
Ljung-Box (Q): 116.72 Jarque-Bera (JB): 35.90
Prob(Q): 0.00 Prob(JB): 0.00
Heteroskedasticity (H): 0.35 Skew: 0.85
Prob(H) (two-sided): 0.00 Kurtosis: 4.95
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [107]: predicted_manual_SARIMA = results_manual_SARIMA.get_forecast(steps=len(test))

rmse = mean_squared_error(test['Rose'],predicted_manual_SARIMA.predicted_mean,squared=False)
mape = mean_absolute_percentage_error(test['Rose'],predicted_manual_SARIMA.predicted_mean)
print('RMSE:',rmse,'\nMAPE:',mape)
```

RMSE: 37.446643609350396

MAPE: 77.28515074326211

```
In [108]: results_2 = pd.DataFrame({'Test RMSE': [rmse], 'MAPE':mape}
 ,index=['SARIMA(1,1,1)(0,0,0,12)'])
results_models = pd.concat([results_models,results_2])
results_models
```

Out[108]:

|                                | Test RMSE | MAPE      |
|--------------------------------|-----------|-----------|
| <b>SARIMA(3,1,1)(3,0,2,12)</b> | 18.912077 | 36.458684 |
| <b>SARIMA(2,1,2)(1,1,1,12)</b> | 13.278681 | 17.133917 |
| <b>SARIMA(1,1,1)(0,0,0,12)</b> | 37.446644 | 77.285151 |

```
In [109]: manual_SARIMA = sm.tsa.statespace.SARIMAX(train['Rose'],
 order=(3,1,0),
 seasonal_order=(1, 0, 1, 12),
 enforce_stationarity=False,
 enforce_invertibility=False)
results_manual_SARIMA = manual_SARIMA.fit(maxiter=1000)
print(results_manual_SARIMA.summary())
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

#### SARIMAX Results

```
=====
Dep. Variable: Rose No. Observations: 132
Model: SARIMAX(3, 1, 0)x(1, 0, [1], 12) Log Likelihood -526.778
Date: Sat, 09 Oct 2021 AIC 1065.557
Time: 09:28:41 BIC 1082.079
Sample: 01-01-1980 HQIC 1072.264
 - 12-01-1990
Covariance Type: opg
=====
```

|          | coef     | std err | z      | P> z  | [0.025  | 0.975]  |
|----------|----------|---------|--------|-------|---------|---------|
| ar.L1    | -0.5474  | 0.090   | -6.087 | 0.000 | -0.724  | -0.371  |
| ar.L2    | -0.5233  | 0.066   | -7.875 | 0.000 | -0.654  | -0.393  |
| ar.L3    | -0.2008  | 0.117   | -1.712 | 0.087 | -0.431  | 0.029   |
| ar.S.L12 | 0.8809   | 0.031   | 27.976 | 0.000 | 0.819   | 0.943   |
| ma.S.L12 | -0.8080  | 0.156   | -5.187 | 0.000 | -1.113  | -0.503  |
| sigma2   | 463.0906 | 77.996  | 5.937  | 0.000 | 310.222 | 615.959 |

```
=====
Ljung-Box (Q): 27.62 Jarque-Bera (JB): 53.77
Prob(Q): 0.93 Prob(JB): 0.00
Heteroskedasticity (H): 0.43 Skew: 0.40
Prob(H) (two-sided): 0.01 Kurtosis: 6.24
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [110]: predicted_manual_SARIMA = results_manual_SARIMA.get_forecast(steps=len(test))

rmse = mean_squared_error(test['Rose'],predicted_manual_SARIMA.predicted_mean,squared=False)
mape = mean_absolute_percentage_error(test['Rose'],predicted_manual_SARIMA.predicted_mean)
print('RMSE:',rmse,'\nMAPE:',mape)
```

RMSE: 32.65456537387812

MAPE: 67.15727281651775

```
In [111]: results_3 = pd.DataFrame({'Test RMSE': [rmse], 'MAPE':mape}
 ,index=['SARIMA(3,1,0)(1,0,1,12)'])
results_models = pd.concat([results_models,results_3])
results_models
```

Out[111]:

|                                | Test RMSE | MAPE      |
|--------------------------------|-----------|-----------|
| <b>SARIMA(3,1,1)(3,0,2,12)</b> | 18.912077 | 36.458684 |
| <b>SARIMA(2,1,2)(1,1,1,12)</b> | 13.278681 | 17.133917 |
| <b>SARIMA(1,1,1)(0,0,0,12)</b> | 37.446644 | 77.285151 |
| <b>SARIMA(3,1,0)(1,0,1,12)</b> | 32.654565 | 67.157273 |



```
In [112]: manual_SARIMA = sm.tsa.statespace.SARIMAX(train['Rose'],
 order=(2,1,1),
 seasonal_order=(1, 1, 0, 12),
 enforce_stationarity=False,
 enforce_invertibility=False)
results_manual_SARIMA = manual_SARIMA.fit(maxiter=1000)
print(results_manual_SARIMA.summary())
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

#### SARIMAX Results

```
=====
Dep. Variable: Rose No. Observations: 132
Model: SARIMAX(2, 1, 1)x(1, 1, [], 12) Log Likelihood -458.160
Date: Sat, 09 Oct 2021 AIC 926.320
Time: 09:28:42 BIC 939.590
Sample: 01-01-1980 HQIC 931.697
 - 12-01-1990
Covariance Type: opg
=====
```

|          | coef     | std err | z       | P> z  | [0.025  | 0.975]  |
|----------|----------|---------|---------|-------|---------|---------|
| ar.L1    | 0.2202   | 0.113   | 1.954   | 0.051 | -0.001  | 0.441   |
| ar.L2    | -0.1033  | 0.105   | -0.980  | 0.327 | -0.310  | 0.103   |
| ma.L1    | -0.9140  | 0.061   | -15.070 | 0.000 | -1.033  | -0.795  |
| ar.S.L12 | -0.4028  | 0.059   | -6.801  | 0.000 | -0.519  | -0.287  |
| sigma2   | 358.1205 | 49.037  | 7.303   | 0.000 | 262.010 | 454.231 |

```
=====
Ljung-Box (Q): 24.63 Jarque-Bera (JB): 0.38
Prob(Q): 0.97 Prob(JB): 0.83
Heteroskedasticity (H): 0.62 Skew: 0.14
Prob(H) (two-sided): 0.16 Kurtosis: 3.09
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [113]: predicted_manual_SARIMA = results_manual_SARIMA.get_forecast(steps=len(test))

rmse = mean_squared_error(test['Rose'], predicted_manual_SARIMA.predicted_mean, squared=False)
mape = mean_absolute_percentage_error(test['Rose'], predicted_manual_SARIMA.predicted_mean)
print('RMSE:', rmse, '\nMAPE:', mape)
```

RMSE: 16.4632845999064

MAPE: 23.552683718744845

```
In [114]: results_4 = pd.DataFrame({'Test RMSE': [rmse], 'MAPE': mape}
 , index=['SARIMA(2,1,1)(1,1,0,12)'])
results_models = pd.concat([results_models, results_4])
results_models
```

Out[114]:

|                                | Test RMSE | MAPE      |
|--------------------------------|-----------|-----------|
| <b>SARIMA(3,1,1)(3,0,2,12)</b> | 18.912077 | 36.458684 |
| <b>SARIMA(2,1,2)(1,1,1,12)</b> | 13.278681 | 17.133917 |
| <b>SARIMA(1,1,1)(0,0,0,12)</b> | 37.446644 | 77.285151 |
| <b>SARIMA(3,1,0)(1,0,1,12)</b> | 32.654565 | 67.157273 |
| <b>SARIMA(2,1,1)(1,1,0,12)</b> | 16.463285 | 23.552684 |

```
In [115]: manual_SARIMA = sm.tsa.statespace.SARIMAX(train['Rose'],
 order=(3,1,1),
 seasonal_order=(1, 0, 1, 6),
 enforce_stationarity=False,
 enforce_invertibility=False)
results_manual_SARIMA = manual_SARIMA.fit(maxiter=1000)
print(results_manual_SARIMA.summary())
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

#### SARIMAX Results

```
=====
Dep. Variable: Rose No. Observations: 132
Model: SARIMAX(3, 1, 1)x(1, 0, 1, 6) Log Likelihood -568.546
Date: Sat, 09 Oct 2021 AIC 1151.092
Time: 09:28:43 BIC 1170.720
Sample: 01-01-1980 HQIC 1159.065
 - 12-01-1990
Covariance Type: opg
=====
```

|         | coef     | std err  | z       | P> z  | [0.025    | 0.975]   |
|---------|----------|----------|---------|-------|-----------|----------|
| ar.L1   | 0.0102   | 0.083    | 0.123   | 0.902 | -0.152    | 0.172    |
| ar.L2   | -0.2493  | 0.106    | -2.345  | 0.019 | -0.458    | -0.041   |
| ar.L3   | 0.0025   | 0.085    | 0.029   | 0.977 | -0.165    | 0.170    |
| ma.L1   | -0.8397  | 0.075    | -11.222 | 0.000 | -0.986    | -0.693   |
| ar.S.L6 | -0.9530  | 0.020    | -46.944 | 0.000 | -0.993    | -0.913   |
| ma.S.L6 | 1.0000   | 429.160  | 0.002   | 0.998 | -840.138  | 842.138  |
| sigma2  | 564.8418 | 2.42e+05 | 0.002   | 0.998 | -4.75e+05 | 4.76e+05 |

```
=====
Ljung-Box (Q): 73.65 Jarque-Bera (JB): 27.50
Prob(Q): 0.00 Prob(JB): 0.00
Heteroskedasticity (H): 0.39 Skew: 0.73
Prob(H) (two-sided): 0.00 Kurtosis: 4.82
=====
```

=====

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [116]: predicted_manual_SARIMA = results_manual_SARIMA.get_forecast(steps=len(test))

rmse = mean_squared_error(test['Rose'], predicted_manual_SARIMA.predicted_mean, squared=False)
mape = mean_absolute_percentage_error(test['Rose'], predicted_manual_SARIMA.predicted_mean)
print('RMSE:', rmse, '\nMAPE:', mape)
```

RMSE: 31.19001082060666

MAPE: 62.97904655504848

```
In [117]: results_5 = pd.DataFrame({'Test RMSE': [rmse], 'MAPE': mape}
 , index=['SARIMA(3,1,1)(1,0,1,6)'])
results_models = pd.concat([results_models, results_5])
results_models
```

Out[117]:

|                                | Test RMSE | MAPE      |
|--------------------------------|-----------|-----------|
| <b>SARIMA(3,1,1)(3,0,2,12)</b> | 18.912077 | 36.458684 |
| <b>SARIMA(2,1,2)(1,1,1,12)</b> | 13.278681 | 17.133917 |
| <b>SARIMA(1,1,1)(0,0,0,12)</b> | 37.446644 | 77.285151 |
| <b>SARIMA(3,1,0)(1,0,1,12)</b> | 32.654565 | 67.157273 |
| <b>SARIMA(2,1,1)(1,1,0,12)</b> | 16.463285 | 23.552684 |
| <b>SARIMA(3,1,1)(1,0,1,6)</b>  | 31.190011 | 62.979047 |

In [118]: *#Building a comparison table for comparing all the different models built:*

```
frames=[results,results_models]
result=pd.concat(frames)
result
```

Out[118]:

|                                                  | Test RMSE | MAPE      |
|--------------------------------------------------|-----------|-----------|
| <b>TES: Alpha=0.07, Beta=8.11e-20, Gamma=0.0</b> | 12.831106 | NaN       |
| <b>DES: Alpha=0.16,Beta=0.16</b>                 | 70.604598 | NaN       |
| <b>LR RMSE</b>                                   | 15.278369 | NaN       |
| <b>Naive RMSE</b>                                | 79.745697 | NaN       |
| <b>SA RMSE</b>                                   | 53.488233 | NaN       |
| <b>2-point MA</b>                                | 11.530054 | NaN       |
| <b>4-point MA</b>                                | 14.458402 | NaN       |
| <b>6-point MA</b>                                | 14.572976 | NaN       |
| <b>9-point MA</b>                                | 14.732918 | NaN       |
| <b>SARIMA(3,1,1)(3,0,2,12)</b>                   | 18.912077 | 36.458684 |
| <b>SARIMA(2,1,2)(1,1,1,12)</b>                   | 13.278681 | 17.133917 |
| <b>SARIMA(1,1,1)(0,0,0,12)</b>                   | 37.446644 | 77.285151 |
| <b>SARIMA(3,1,0)(1,0,1,12)</b>                   | 32.654565 | 67.157273 |
| <b>SARIMA(2,1,1)(1,1,0,12)</b>                   | 16.463285 | 23.552684 |
| <b>SARIMA(3,1,1)(1,0,1,6)</b>                    | 31.190011 | 62.979047 |

In [119]: *#Comparing RMSE of all models, we can go with 2-point Moving Average, TES & SARIMA(2,1,2)(1,1,1,12) as t*

In [120]: *#Final forecasting using TES model:*

```
model_final = ExponentialSmoothing(data,trend='multiplicative',seasonal='multiplicative')
```

```
Fitting the model
```

```
model_final = model_final.fit()
```

```
print('')
```

```
print('~~~ Holt Winters model Exponential Smoothing Estimated Parameters ~~~')
```

```
print('')
```

```
print(model_final.params)
```

```
/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:
ng:
```

No frequency information was provided, so inferred frequency MS will be used.

~~~ Holt Winters model Exponential Smoothing Estimated Parameters ~~~


```
{'smoothing_level': 0.08409605754446657, 'smoothing_slope': 1.9059183137266802e-28, 'smoothing_seasona
l': 0.0, 'damping_slope': nan, 'initial_level': 64.03828414961501, 'initial_slope': 0.993307504511918
6, 'initial_seasons': array([1.72703925, 1.96783625, 2.17516895, 1.9448051 , 2.12245456,
 2.30708257, 2.58780767, 2.66156174, 2.52889052, 2.49412681,
 2.90036268, 4.00865224]), 'use_boxcox': False, 'lamda': None, 'remove_bias': False}
```

In [130]: *# Plotting the predictions:*

```
data.plot()
model_final.forecast(steps=12).plot()
plt.legend(['Observed series', 'Forecasted series'])
plt.title('Forecasted Rose Sales (TES Model)',color='blue',fontsize=20)
plt.xlabel('Years/Months',color='blue',fontsize=18)
plt.ylabel('Rose Sales',color='blue',fontsize=18)
plt.grid();
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:342: FutureWarning:

The 'freq' argument in Timestamp is deprecated and will be removed in a future version.





In [122]: *#Building an additional forecast using SARIMA(2,1,2)(1,1,1,12):*

```
full_data_model = sm.tsa.statespace.SARIMAX(data['Rose'],
 order=(2,1,2),
 seasonal_order=(1, 1, 1, 12),
 enforce_stationarity=False,
 enforce_invertibility=False)
results_full_data_model = full_data_model.fit(maxiter=1000)
print(results_full_data_model.summary())
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

#### SARIMAX Results

```
=====
Dep. Variable: Rose No. Observations: 187
Model: SARIMAX(2, 1, 2)x(1, 1, [1], 12) Log Likelihood -665.358
Date: Sat, 09 Oct 2021 AIC 1344.716
Time: 09:28:48 BIC 1366.199
Sample: 01-01-1980 HQIC 1353.440
 - 07-01-1995
```

Covariance Type: opg

```
=====
 coef std err z P>|z| [0.025 0.975]

ar.L1 1.1071 0.093 11.850 0.000 0.924 1.290
ar.L2 -0.3258 0.079 -4.098 0.000 -0.482 -0.170
ma.L1 -1.8269 0.068 -26.777 0.000 -1.961 -1.693
ma.L2 0.8776 0.061 14.465 0.000 0.759 0.996
ar.S.L12 -0.3824 0.049 -7.752 0.000 -0.479 -0.286
ma.S.L12 -0.0853 0.093 -0.916 0.360 -0.268 0.097
sigma2 251.0546 26.973 9.308 0.000 198.189 303.921
=====
```

```
Ljung-Box (Q): 35.74 Jarque-Bera (JB): 2.97
```

```

Prob(Q): 0.66 Prob(JB): 0.23
Heteroskedasticity (H): 0.22 Skew: 0.04
Prob(H) (two-sided): 0.00 Kurtosis: 3.66
=====

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [123]: predicted_full_data = results_full_data_model.get_forecast(steps=12)
```

```
In [124]: pred_SARIMA = predicted_full_data.summary_frame(alpha=0.05)
pred_SARIMA.head()
```

Out[124]:

| Rose       | mean      | mean_se   | mean_ci_lower | mean_ci_upper |
|------------|-----------|-----------|---------------|---------------|
| 1995-08-01 | 52.762088 | 15.844704 | 21.707038     | 83.817137     |
| 1995-09-01 | 47.558565 | 16.454734 | 15.307878     | 79.809252     |
| 1995-10-01 | 52.319731 | 16.464088 | 20.050711     | 84.588752     |
| 1995-11-01 | 57.435555 | 16.464115 | 25.166483     | 89.704627     |
| 1995-12-01 | 82.686197 | 16.474636 | 50.396504     | 114.975890    |

```
In [125]: pred_SARIMA.tail()
```

Out[125]:

| Rose       | mean      | mean_se   | mean_ci_lower | mean_ci_upper |
|------------|-----------|-----------|---------------|---------------|
| 1996-03-01 | 45.939048 | 16.921644 | 12.773236     | 79.104860     |
| 1996-04-01 | 52.302785 | 17.215946 | 18.560152     | 86.045418     |
| 1996-05-01 | 37.739694 | 17.548626 | 3.345019      | 72.134369     |
| 1996-06-01 | 45.111871 | 17.901053 | 10.026451     | 80.197291     |
| 1996-07-01 | 57.115584 | 18.260924 | 21.324831     | 92.906336     |


In [129]: *# Plotting the forecast with confidence interval:*

```
axis = data['Rose'].plot(label='Observed series')
pred_SARIMA['mean'].plot(ax=axis, label='Forecasted series', alpha=0.7)

axis.fill_between(pred_SARIMA['mean'].index, pred_SARIMA['mean_ci_lower'], pred_SARIMA['mean_ci_upper'], c

plt.title('Forecasted Rose Sales (SARIMA model)', color='blue', fontsize=20)
plt.xlabel('Years/Months', color='blue', fontsize=18)
plt.ylabel('Rose Sales', color='blue', fontsize=18)

plt.legend(loc='best')
plt.grid();
```



```
In []: ## THE END!
```