

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import plotly.offline as py
py.init_notebook_mode()
%matplotlib inline
import seaborn as sns
from pylab import rcParams
```

```
In [2]: df = pd.read_csv("Sparkling.csv", parse_dates=True, index_col=0)
```

```
In [3]: df.head(10)
```

Out[3]:

Sparkling	
YearMonth	
1980-01-01	1686
1980-02-01	1591
1980-03-01	2304
1980-04-01	1712
1980-05-01	1471
1980-06-01	1377
1980-07-01	1966
1980-08-01	2453
1980-09-01	1984
1980-10-01	2596

```
In [4]: df.tail(10)
```


```
Out[4]:
```

Sparkling	
YearMonth	
1994-10-01	3385
1994-11-01	3729
1994-12-01	5999
1995-01-01	1070
1995-02-01	1402
1995-03-01	1897
1995-04-01	1862
1995-05-01	1670
1995-06-01	1688
1995-07-01	2031

##Plotting the distribution of Sparkling sales:

```
In [5]: rcParams['figure.figsize'] = 12,8
df.plot();
plt.grid()
plt.title('Sparkling Sales Distribution',color='blue',fontsize=16)
plt.xlabel('Year/Month',color='blue',fontsize=14)
plt.ylabel('Sparkling Sales',color='blue',fontsize=14)
```

```
Out[5]: Text(0, 0.5, 'Sparkling Sales')
```



```
In [6]: df.shape
```

```
Out[6]: (187, 1)
```

```
In [7]: df.describe()
```

```
Out[7]:
```

	Sparkling
count	187.000000
mean	2402.417112
std	1295.111540
min	1070.000000
25%	1605.000000
50%	1874.000000
75%	2549.000000
max	7242.000000


```
In [8]: df.isnull().sum()
```

```
Out[8]: Sparkling      0  
dtype: int64
```

#Plotting the year-on-year sales of 'Sparkling':

```
In [9]: fig, ax = plt.subplots(figsize=(20,10))
sns.boxplot(df.index.year, df.values[:,0], ax=ax,whis=1.5)
plt.grid();
plt.xlabel('Years',color='blue',fontsize=18);
plt.ylabel('Sparkling Sales',color='blue',fontsize=18);
plt.title('Year-on-Year Sparkling Sales Distribution',color='blue',fontsize=20)
```


```
Out[9]: Text(0.5, 1.0, 'Year-on-Year Sparkling Sales Distribution')
```



```
In [10]: #Plotting month-wise sales distribution across years:
```

```
fig, ax = plt.subplots(figsize=(20,12))
sns.boxplot(df.index.month_name(), df.values[:,0], ax=ax,whis=1.5)
plt.grid();
plt.xlabel('Months',color='blue',fontsize=18);
plt.ylabel('Sparkling Sales',color='blue',fontsize=18);
plt.title('Month-wise Sparkling Sales Distribution across years',color='blue',fontsize=20)
```

```
Out[10]: Text(0.5, 1.0, 'Month-wise Sparkling Sales Distribution across years')
```




```
In [11]: from statsmodels.graphics.tsaplots import month_plot
```

```
fig, ax = plt.subplots(figsize=(20,12))
```

```
month_plot(df,ax=ax)
```


```
plt.grid();
```

```
plt.xlabel('Months',color='blue',fontsize=18);
```

```
plt.ylabel('Sparkling Sales',color='blue',fontsize=18);
```

```
plt.title('Year-wise distribution across months of Sparkling Sales',color='blue',fontsize=20)
```

```
Out[11]: Text(0.5, 1.0, 'Year-wise distribution across months of Sparkling Sales')
```



In [12]: *# Computing and plotting the per month sales for each year:*


```
monthly_sales_across_years = pd.pivot_table(df, values = 'Sparkling', columns = df.index.month_name(), i
monthly_sales_across_years
```

Out[12]:

YearMonth	April	August	December	February	January	July	June	March	May	November	October	September
YearMonth												
1980	1712.0	2453.0	5179.0	1591.0	1686.0	1966.0	1377.0	2304.0	1471.0	4087.0	2596.0	1984.0
1981	1976.0	2472.0	4551.0	1523.0	1530.0	1781.0	1480.0	1633.0	1170.0	3857.0	2273.0	1981.0
1982	1790.0	1897.0	4524.0	1329.0	1510.0	1954.0	1449.0	1518.0	1537.0	3593.0	2514.0	1706.0
1983	1375.0	2298.0	4923.0	1638.0	1609.0	1600.0	1245.0	2030.0	1320.0	3440.0	2511.0	2191.0
1984	1789.0	3159.0	5274.0	1435.0	1609.0	1597.0	1404.0	2061.0	1567.0	4273.0	2504.0	1759.0
1985	1589.0	2512.0	5434.0	1682.0	1771.0	1645.0	1379.0	1846.0	1896.0	4388.0	3727.0	1771.0
1986	1605.0	3318.0	5891.0	1523.0	1606.0	2584.0	1403.0	1577.0	1765.0	3987.0	2349.0	1562.0
1987	1935.0	1930.0	7242.0	1442.0	1389.0	1847.0	1250.0	1548.0	1518.0	4405.0	3114.0	2638.0
1988	2336.0	1645.0	6757.0	1779.0	1853.0	2230.0	1661.0	2108.0	1728.0	4988.0	3740.0	2421.0
1989	1650.0	1968.0	6694.0	1394.0	1757.0	1971.0	1406.0	1982.0	1654.0	4514.0	3845.0	2608.0
1990	1628.0	1605.0	6047.0	1321.0	1720.0	1899.0	1457.0	1859.0	1615.0	4286.0	3116.0	2424.0
1991	1279.0	1857.0	6153.0	2049.0	1902.0	2214.0	1540.0	1874.0	1432.0	3627.0	3252.0	2408.0
1992	1997.0	1773.0	6119.0	1667.0	1577.0	2076.0	1625.0	1993.0	1783.0	4096.0	3088.0	2377.0
1993	2121.0	2795.0	6410.0	1564.0	1494.0	2048.0	1515.0	1898.0	1831.0	4227.0	3339.0	1749.0
1994	1725.0	1495.0	5999.0	1968.0	1197.0	2031.0	1693.0	1720.0	1674.0	3729.0	3385.0	2968.0
1995	1862.0	NaN	NaN	1402.0	1070.0	2031.0	1688.0	1897.0	1670.0	NaN	NaN	NaN

```
In [13]: monthly_sales_across_years.plot(figsize=(20,10))
plt.grid()
plt.legend(loc='best');
plt.xlabel('Years',color='blue',fontsize=18);
plt.ylabel('Sparkling Sales',color='blue',fontsize=18);
plt.title('Monthly Sparkling Sales Distribution across Years',color='blue',fontsize=20)
```

Out[13]: Text(0.5, 1.0, 'Monthly Sparkling Sales Distribution across Years')



In [14]: *# Computing and plotting mean sales for each year:*


```
df_yearly_mean = df.resample('Y').mean()  
df_yearly_mean.head()
```

Out[14]:

Sparkling	
YearMonth	
1980-12-31	2367.166667
1981-12-31	2185.583333
1982-12-31	2110.083333
1983-12-31	2181.666667
1984-12-31	2369.250000

```
In [15]: df_yearly_mean.plot();  
plt.grid()  
plt.xlabel('Mean Sales per year',color='blue',fontsize=18);  
plt.ylabel('Mean Sales',color='blue',fontsize=18);  
plt.title('Mean Sparkling Sales distribution across Years',color='blue',fontsize=20)
```

```
Out[15]: Text(0.5, 1.0, 'Mean Sparkling Sales distribution across Years')
```



In [16]: *# Computing and plotting mean sales for each quarter:*

```
df_quarterly_mean = df.resample('Q').mean()  
df_quarterly_mean.head()
```

Out[16]:

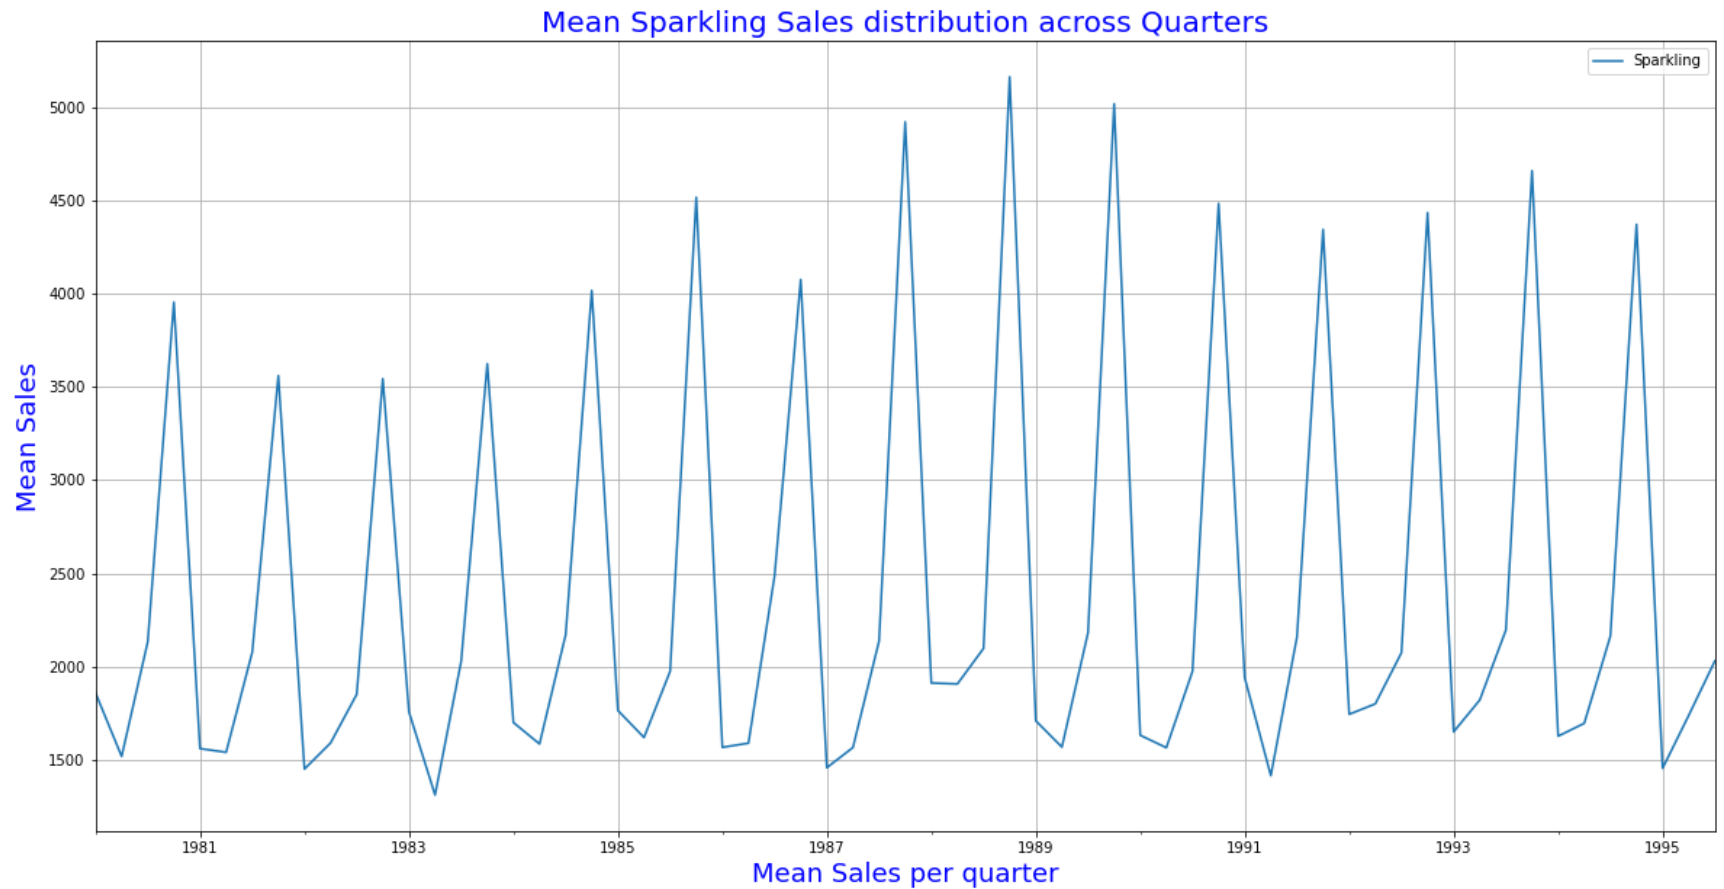
Sparkling

YearMonth

1980-03-31	1860.333333
1980-06-30	1520.000000
1980-09-30	2134.333333
1980-12-31	3954.000000
1981-03-31	1562.000000

```
In [17]: df_quarterly_mean.plot(figsize=(20,10));  
plt.grid()  
plt.xlabel('Mean Sales per quarter',color='blue',fontsize=18);  
plt.ylabel('Mean Sales',color='blue',fontsize=18);  
plt.title('Mean Sparkling Sales distribution across Quarters',color='blue',fontsize=20)
```


Out[17]: Text(0.5, 1.0, 'Mean Sparkling Sales distribution across Quarters')




#Decomposing the Time Series:

```
In [18]: from statsmodels.tsa.seasonal import seasonal_decompose
```

```
In [19]: decomposition_add = seasonal_decompose(df,model='additive')  
decomposition_add.plot();
```



```
In [20]: decomposition_multi = seasonal_decompose(df,model='multiplicative')  
decomposition_multi.plot();
```




```
In [21]: # Computing the various components of the decomposed data:
```

```
trend = decomposition_multi.trend  
seasonality = decomposition_multi.seasonal  
residual = decomposition_multi.resid
```

```
In [22]: # Checking the components:
```

```
print('Trend in Sparkling Sales', '\n', trend.head(12), '\n')
print('Seasonality in Sparkling Sales', '\n', seasonality.head(12), '\n')
print('Residual', '\n', residual.head(12), '\n')
```

Trend in Sparkling Sales

YearMonth	
1980-01-01	NaN
1980-02-01	NaN
1980-03-01	NaN
1980-04-01	NaN
1980-05-01	NaN
1980-06-01	NaN
1980-07-01	2360.666667
1980-08-01	2351.333333
1980-09-01	2320.541667
1980-10-01	2303.583333
1980-11-01	2302.041667
1980-12-01	2293.791667

Name: trend, dtype: float64

Seasonality in Sparkling Sales

YearMonth	
1980-01-01	0.649843
1980-02-01	0.659214
1980-03-01	0.757440
1980-04-01	0.730351
1980-05-01	0.660609
1980-06-01	0.603468
1980-07-01	0.809164
1980-08-01	0.918822
1980-09-01	0.894367
1980-10-01	1.241789
1980-11-01	1.690158
1980-12-01	2.384776

Name: seasonal, dtype: float64

Residual

YearMonth	
1980-01-01	NaN
1980-02-01	NaN
1980-03-01	NaN

```
1980-04-01      NaN
1980-05-01      NaN
1980-06-01      NaN
1980-07-01    1.029230
1980-08-01    1.135407
1980-09-01    0.955954
1980-10-01    0.907513
1980-11-01    1.050423
1980-12-01    0.946770
Name: resid, dtype: float64
```


```
In [23]: # Checking how the data looks without seasonality:
```

```
deaseasonalized_ts = trend + residual
deaseasonalized_ts.head(12)
```

```
Out[23]: YearMonth
1980-01-01      NaN
1980-02-01      NaN
1980-03-01      NaN
1980-04-01      NaN
1980-05-01      NaN
1980-06-01      NaN
1980-07-01    2361.695896
1980-08-01    2352.468741
1980-09-01    2321.497620
1980-10-01    2304.490847
1980-11-01    2303.092089
1980-12-01    2294.738436
dtype: float64
```

```
In [24]: df.plot()  
deaseasonalized_ts.plot()  
plt.legend(["Original Time Series", "Deseasonalised Time Series"]);  
plt.xlabel('Years',color='blue',fontsize=18);  
plt.ylabel('Sparkling Sales',color='blue',fontsize=18);  
plt.title('Original v/s Deseasonalised Sparkling Sales distribution',color='blue',fontsize=20)
```

Out[24]: Text(0.5, 1.0, 'Original v/s Deseasonalised Sparkling Sales distribution')



#Splitting data into train and test set:

```
In [25]: train = df[df.index<'1991']  
test    = df[df.index>='1991']
```

```
In [26]: print(train.shape)  
         print(test.shape)
```

```
(132, 1)
```

```
(55, 1)
```

```
In [27]: print('First few rows of Training Data', '\n', train.head(), '\n')
print('Last few rows of Training Data', '\n', train.tail(), '\n')
print('First few rows of Test Data', '\n', test.head(), '\n')
print('Last few rows of Test Data', '\n', test.tail(), '\n')
```

First few rows of Training Data

Sparkling

YearMonth

1980-01-01	1686
1980-02-01	1591
1980-03-01	2304
1980-04-01	1712
1980-05-01	1471

Last few rows of Training Data

Sparkling

YearMonth

1990-08-01	1605
1990-09-01	2424
1990-10-01	3116
1990-11-01	4286
1990-12-01	6047

First few rows of Test Data

Sparkling

YearMonth

1991-01-01	1902
1991-02-01	2049
1991-03-01	1874
1991-04-01	1279
1991-05-01	1432

Last few rows of Test Data


Sparkling

YearMonth

1995-03-01	1897
1995-04-01	1862
1995-05-01	1670
1995-06-01	1688
1995-07-01	2031

```
In [28]: # Plotting the train and test data:
```

```
train['Sparkling'].plot(figsize=(15,8), fontsize=14)
test['Sparkling'].plot(figsize=(15,8), fontsize=14)
plt.grid()
plt.legend(['Training Data', 'Test Data'])
plt.xlabel('Years',color='blue',fontsize=18);
plt.ylabel('Sparkling Sales',color='blue',fontsize=18);
plt.title('Sparkling Sales - Train & Test Data',color='blue',fontsize=20)
plt.show()
```



```
In [29]: #Triple exponential smothing using the Holt-Winter's method:
```

```
In [30]: import statsmodels.tools.eval_measures as em
from sklearn.metrics import mean_squared_error
from statsmodels.tsa.api import ExponentialSmoothing, SimpleExpSmoothing, Holt
from IPython.display import display
from pylab import rcParams
```

```
In [31]: model_TES = ExponentialSmoothing(train,trend='multiplicative',seasonal='multiplicative')
```

```
# Fitting the model
```

```
model_TES = model_TES.fit()
```

```
print('')
```

```
print('~~~ Holt Winters model Exponential Smoothing Estimated Parameters ~~~')
```

```
print('')
```

```
print(model_TES.params)
```

```
/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:
No frequency information was provided, so inferred frequency MS will be used.
```

```
~~~ Holt Winters model Exponential Smoothing Estimated Parameters ~~~
```

```
{'smoothing_level': 0.1533370898171079, 'smoothing_slope': 1.3387629728833717e-20, 'smoothing_seasonal': 0.369040099605268, 'damping_slope': nan, 'initial_level': 1640.0000699266104, 'initial_slope': 1.002822904757003, 'initial_seasons': array([1.00842317, 0.96873745, 1.24208978, 1.13203929, 0.93995306, 0.93800969, 1.22519687, 1.5458432 , 1.27400584, 1.63515799, 2.48733686, 3.12532974]), 'use_boxcox': False, 'lamda': None, 'remove_bias': False}
```



```
In [32]: # Forecasting using this model for the duration of the test set
TES_predict = model_TES.forecast(len(test))
TES_predict
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:342: FutureWarning:

The 'freq' argument in Timestamp is deprecated and will be removed in a future version.

```
Out[32]: 1991-01-01    1603.408052
          1991-02-01    1375.868826
          1991-03-01    1808.629451
          1991-04-01    1706.429435
          1991-05-01    1603.480417
          1991-06-01    1416.360471
          1991-07-01    1946.556945
          1991-08-01    1914.505494
          1991-09-01    2435.737064
          1991-10-01    3335.385360
          1991-11-01    4411.830520
          1991-12-01    6335.097945
          1992-01-01    1658.574554
          1992-02-01    1423.206663
          1992-03-01    1870.856754
          1992-04-01    1765.140467
          1992-05-01    1658.649408
          1992-06-01    1465.091455
          1992-07-01    2013.529751
          1992-08-01    1980.375544
          1992-09-01    2519.540492
          1992-10-01    3450.141888
          1992-11-01    4563.622980
          1992-12-01    6553.061918
          1993-01-01    1715.639101
          1993-02-01    1472.173196
          1993-03-01    1935.225036
          1993-04-01    1825.871498
          1993-05-01    1715.716531
          1993-06-01    1515.499066
          1993-07-01    2082.806808
          1993-08-01    2048.511905
          1993-09-01    2606.227243
```

1993-10-01	3568.846703
1993-11-01	4720.637979
1993-12-01	6778.525110
1994-01-01	1774.666999
1994-02-01	1522.824460
1994-03-01	2001.807959
1994-04-01	1888.692027
1994-05-01	1774.747092
1994-06-01	1567.640990
1994-07-01	2154.467396
1994-08-01	2118.992551
1994-09-01	2695.896518
1994-10-01	3691.635650
1994-11-01	4883.055200
1994-12-01	7011.745539
1995-01-01	1835.725797
1995-02-01	1575.218420
1995-03-01	2070.681718
1995-04-01	1953.673945
1995-05-01	1835.808647
1995-06-01	1621.576897
1995-07-01	2228.593523

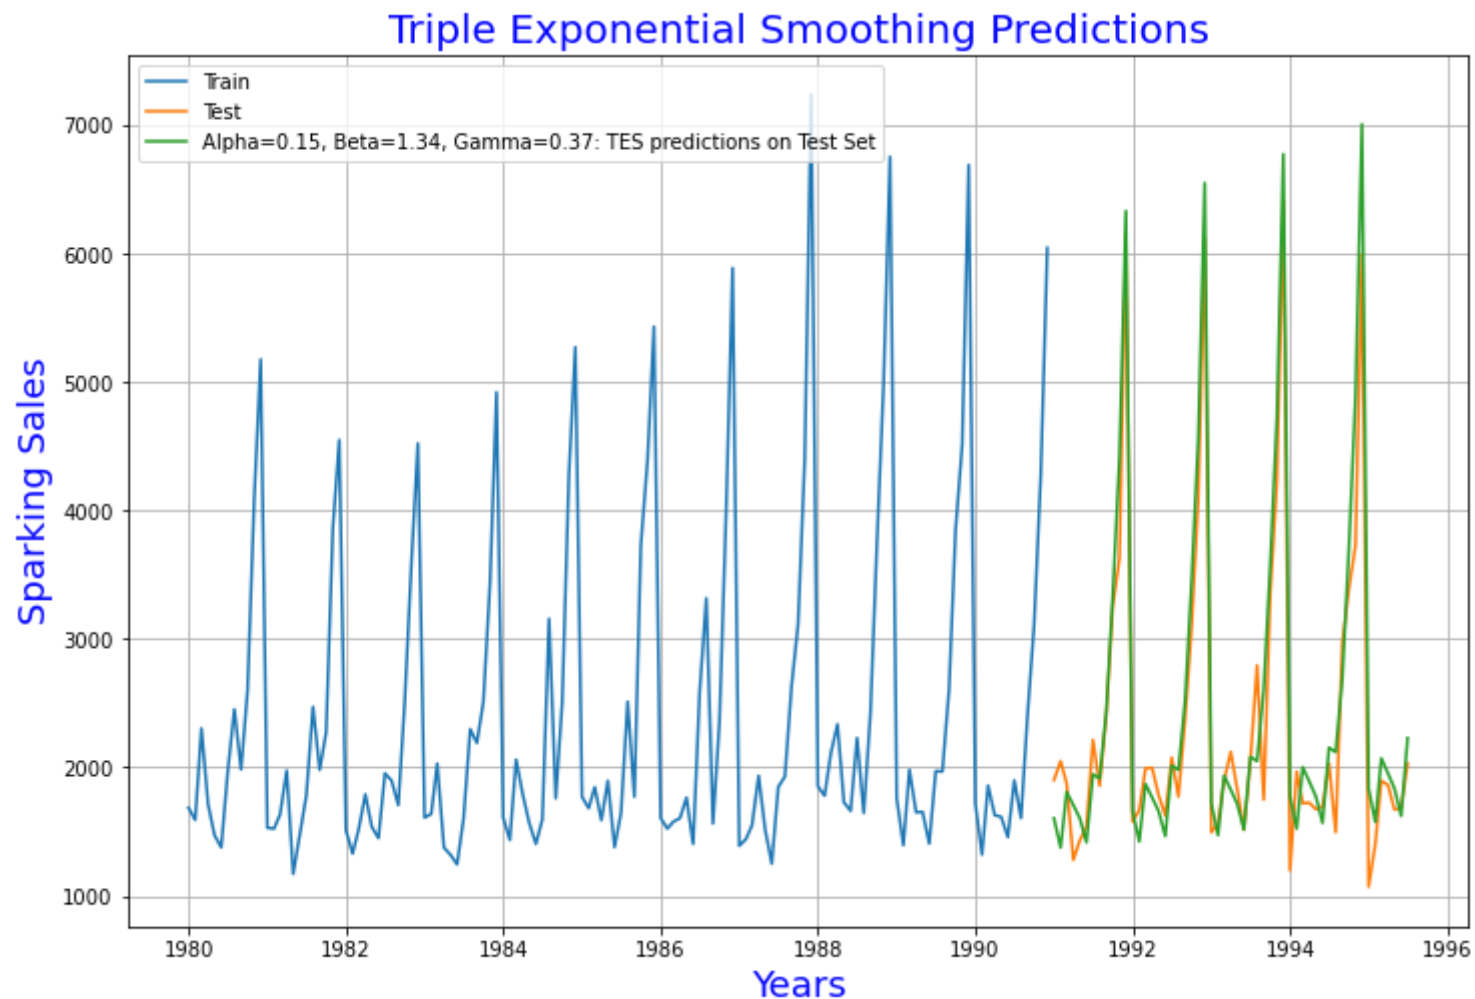
Freq: MS, dtype: float64

```
In [33]: ## Plotting the Training data, Test data and the forecasted values
```

```
plt.plot(train, label='Train')
plt.plot(test, label='Test')

plt.plot(TES_predict, label='Alpha=0.15, Beta=1.34, Gamma=0.37: TES predictions on Test Set')

plt.legend(loc='best')
plt.grid()
plt.title('Triple Exponential Smoothing Predictions',color='blue',fontsize=20);
plt.xlabel('Years',color='blue',fontsize=18);
plt.ylabel('Sparkling Sales',color='blue',fontsize=18);
```



In [34]: *#Evaluating the TES method using RSME:*

In [35]: `print('TES RMSE:',mean_squared_error(test.values, TES_predict.values,squared=False))`

TES RMSE: 392.932696056841

In [36]: `results = pd.DataFrame({'Test RMSE': [mean_squared_error(test.values, TES_predict.values,squared=False)]},
index=['TES: Alpha=0.15, Beta=1.34, Gamma=0.37'])`
results

Out[36]:

	Test RMSE
TES: Alpha=0.15, Beta=1.34, Gamma=0.37	392.932696

In [37]: *#Triple exponential smothing using Holt-Winter's method with tweaked parameters:*

```
In [38]: model_TES_tweaked = ExponentialSmoothing(train,trend='additive',seasonal='multiplicative')
# Fitting the model
model_TES_tweaked = model_TES_tweaked.fit()

print('')
print('~~~ Holt Winters Exponential Smoothing Tweaked Parameters ~~~')
print('')
print(model_TES_tweaked.params)
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

~~~ Holt Winters Exponential Smoothing Tweaked Parameters ~~~

```
{'smoothing_level': 0.1542165016472889, 'smoothing_slope': 1.2783091295878023e-21, 'smoothing_seasonal': 0.37133343678021546, 'damping_slope': nan, 'initial_level': 1639.999344596867, 'initial_slope': 4.847082379796527, 'initial_seasons': array([1.00842292, 0.96898445, 1.24179517, 1.13206135, 0.93982422, 0.93811215, 1.2245931, 1.54431416, 1.27337131, 1.63199235, 2.48297943, 3.11867021]), 'use_boxcox': False, 'lamda': None, 'remove_bias': False}
```

In [39]: *#Forecasting data using tweaked TES model:*

```
TES_predict_tweaked = model_TES_tweaked.forecast(len(test))  
TES_predict_tweaked
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:342: FutureWarning:

The 'freq' argument in Timestamp is deprecated and will be removed in a future version.

```
Out[39]: 1991-01-01    1602.184272  
1991-02-01    1373.875569  
1991-03-01    1807.428191  
1991-04-01    1704.559896  
1991-05-01    1602.365129  
1991-06-01    1415.471251  
1991-07-01    1944.839173  
1991-08-01    1910.035018  
1991-09-01    2435.185926  
1991-10-01    3333.435905  
1991-11-01    4407.750452  
1991-12-01    6328.489537  
1992-01-01    1656.041881  
1992-02-01    1419.929545  
1992-03-01    1867.846602  
1992-04-01    1761.381358  
1992-05-01    1655.631960  
1992-06-01    1462.395248  
1992-07-01    2009.134517  
1992-08-01    1973.006273  
1992-09-01    2515.250723  
1992-10-01    3442.734180  
1992-11-01    4551.880046  
1992-12-01    6534.863270  
1993-01-01    1709.899489  
1993-02-01    1465.983521  
1993-03-01    1928.265013  
1993-04-01    1818.202820  
1993-05-01    1708.898791  
1993-06-01    1509.319245  
1993-07-01    2073.429862  
1993-08-01    2035.977528
```

|            |             |
|------------|-------------|
| 1993-09-01 | 2595.315519 |
| 1993-10-01 | 3552.032456 |
| 1993-11-01 | 4696.009639 |
| 1993-12-01 | 6741.237002 |
| 1994-01-01 | 1763.757098 |
| 1994-02-01 | 1512.037496 |
| 1994-03-01 | 1988.683424 |
| 1994-04-01 | 1875.024282 |
| 1994-05-01 | 1762.165622 |
| 1994-06-01 | 1556.243242 |
| 1994-07-01 | 2137.725206 |
| 1994-08-01 | 2098.948783 |
| 1994-09-01 | 2675.380316 |
| 1994-10-01 | 3661.330731 |
| 1994-11-01 | 4840.139233 |
| 1994-12-01 | 6947.610735 |
| 1995-01-01 | 1817.614706 |
| 1995-02-01 | 1558.091472 |
| 1995-03-01 | 2049.101835 |
| 1995-04-01 | 1931.845744 |
| 1995-05-01 | 1815.432453 |
| 1995-06-01 | 1603.167240 |
| 1995-07-01 | 2202.020551 |

Freq: MS, dtype: float64

In [40]: *# Plotting the Training data, Test data and the forecasted values:*


```
plt.plot(train, label='Train')
plt.plot(test, label='Test')

plt.plot(TES_predict, label='Alpha=0.15, Beta=1.34, Gamma=0.37:TES predictions on Test Set')
plt.plot(TES_predict_tweaked, label='Alpha=0.15,Beta=1.28e-21,Gamma=0.37:Tweaked TES predictions on Test')

plt.legend(loc='best')
plt.grid()
plt.title('TES Predictions - Normal & Tweaked',color='blue',fontsize=20);
plt.xlabel('Years',color='blue',fontsize=18);
plt.ylabel('Sparking Sales',color='blue',fontsize=18);
```



## TES Predictions - Normal &amp; Tweaked



```
In [41]: print('TES_tweaked RMSE:', mean_squared_error(test.values, TES_predict_tweaked.values, squared=False))
```

TES\_tweaked RMSE: 383.1384642466706

```
In [42]: results_smoothing_1 = pd.DataFrame({'Test RMSE': [mean_squared_error(test.values, TES_predict_tweaked.values, index=['TES_tweaked: Alpha=0.15, Beta=1.28e-21, Gamma=0.37'])], index=['TES_tweaked: Alpha=0.15, Beta=1.28e-21, Gamma=0.37'])
results = pd.concat([results, results_smoothing_1])
results
```

Out[42]:

|                                                    | Test RMSE  |
|----------------------------------------------------|------------|
| TES: Alpha=0.15, Beta=1.34, Gamma=0.37             | 392.932696 |
| TES_tweaked: Alpha=0.15, Beta=1.28e-21, Gamma=0.37 | 383.138464 |

```
In [43]: # Double exponential smothing using the Holt's method:
```

```
In [44]: model_DES = Holt(train)
# Fitting the model
model_DES = model_DES.fit()

print('')
print('~~~ Holt DES model Estimated Parameters ~~~')
print('')
print(model_DES.params)
```

~~~ Holt DES model Estimated Parameters ~~~

```
{'smoothing_level': 0.6478091609267566, 'smoothing_slope': 0.0, 'smoothing_seasonal': nan, 'damping_slope': nan, 'initial_level': 1686.0838036944974, 'initial_slope': 27.068228572915256, 'initial_seasons': array([], dtype=float64), 'use_boxcox': False, 'lamda': None, 'remove_bias': False}
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

```
In [45]: # Forecasting using this model for the duration of the test set
DES_predict = model_DES.forecast(len(test))
DES_predict
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:342: FutureWarning:

The 'freq' argument in Timestamp is deprecated and will be removed in a future version.

```
Out[45]: 1991-01-01    5281.501604
          1991-02-01    5308.569832
          1991-03-01    5335.638061
          1991-04-01    5362.706289
          1991-05-01    5389.774518
          1991-06-01    5416.842746
          1991-07-01    5443.910975
          1991-08-01    5470.979204
          1991-09-01    5498.047432
          1991-10-01    5525.115661
          1991-11-01    5552.183889
          1991-12-01    5579.252118
          1992-01-01    5606.320347
          1992-02-01    5633.388575
          1992-03-01    5660.456804
          1992-04-01    5687.525032
          1992-05-01    5714.593261
          1992-06-01    5741.661489
          1992-07-01    5768.729718
          1992-08-01    5795.797947
          1992-09-01    5822.866175
          1992-10-01    5849.934404
          1992-11-01    5877.002632
          1992-12-01    5904.070861
          1993-01-01    5931.139089
          1993-02-01    5958.207318
          1993-03-01    5985.275547
          1993-04-01    6012.343775
          1993-05-01    6039.412004
          1993-06-01    6066.480232
          1993-07-01    6093.548461
          1993-08-01    6120.616689
          1993-09-01    6147.684918
```

| | |
|------------|-------------|
| 1993-10-01 | 6174.753147 |
| 1993-11-01 | 6201.821375 |
| 1993-12-01 | 6228.889604 |
| 1994-01-01 | 6255.957832 |
| 1994-02-01 | 6283.026061 |
| 1994-03-01 | 6310.094289 |
| 1994-04-01 | 6337.162518 |
| 1994-05-01 | 6364.230747 |
| 1994-06-01 | 6391.298975 |
| 1994-07-01 | 6418.367204 |
| 1994-08-01 | 6445.435432 |
| 1994-09-01 | 6472.503661 |
| 1994-10-01 | 6499.571889 |
| 1994-11-01 | 6526.640118 |
| 1994-12-01 | 6553.708347 |
| 1995-01-01 | 6580.776575 |
| 1995-02-01 | 6607.844804 |
| 1995-03-01 | 6634.913032 |
| 1995-04-01 | 6661.981261 |
| 1995-05-01 | 6689.049489 |
| 1995-06-01 | 6716.117718 |
| 1995-07-01 | 6743.185947 |

Freq: MS, dtype: float64

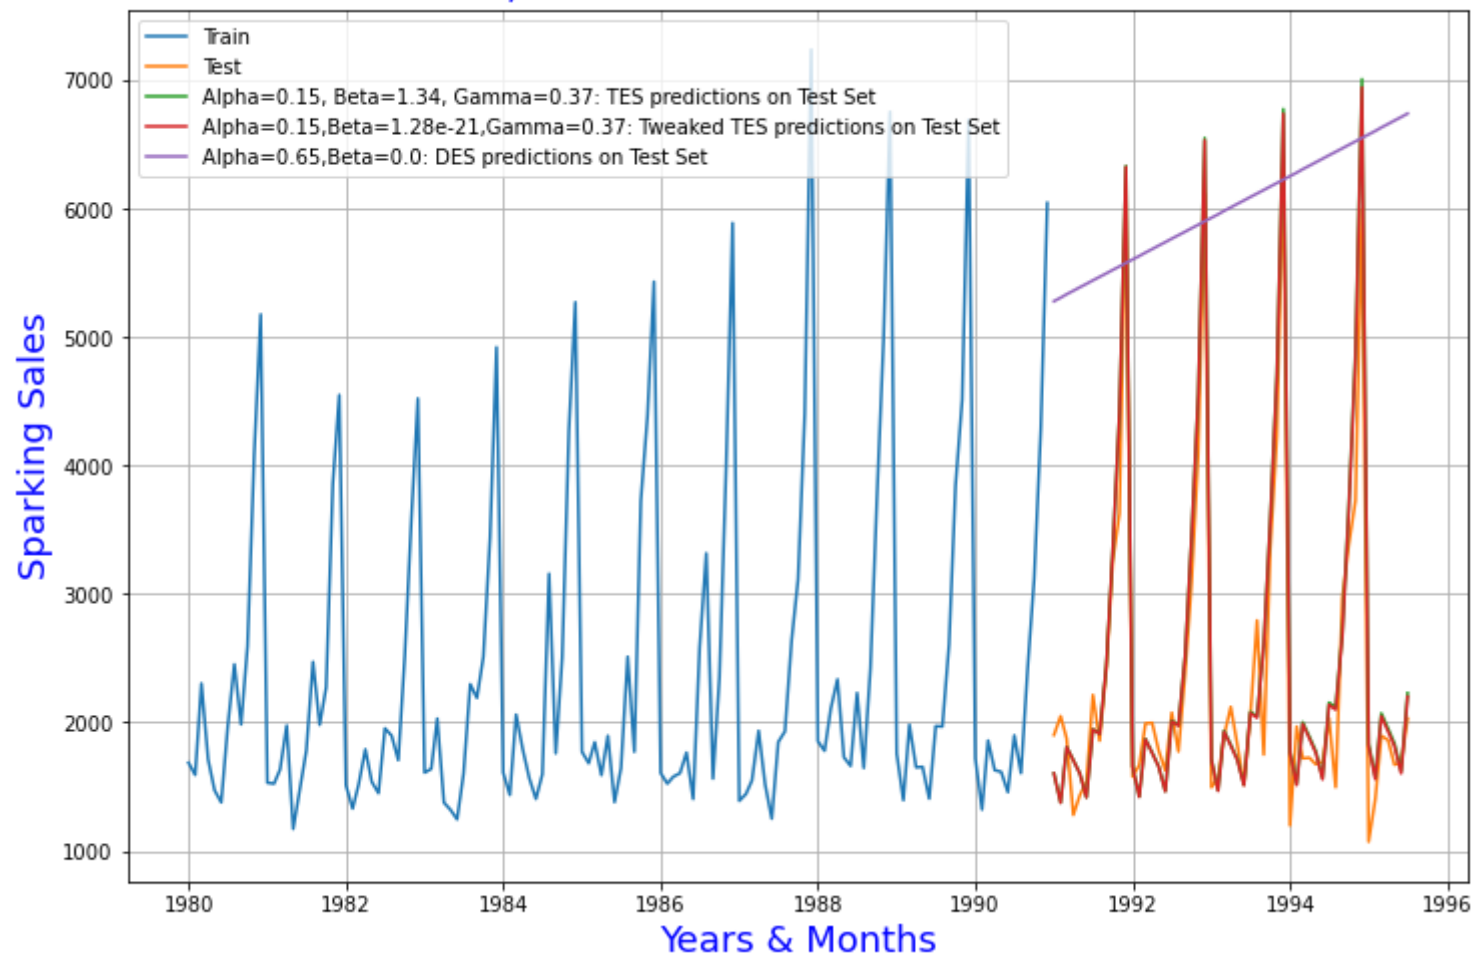
In [46]: *# Plotting the Training data, Test data and the forecasted values:*

```
plt.plot(train, label='Train')
plt.plot(test, label='Test')

plt.plot(TES_predict, label='Alpha=0.15, Beta=1.34, Gamma=0.37: TES predictions on Test Set')
plt.plot(TES_predict_tweaked, label='Alpha=0.15,Beta=1.28e-21,Gamma=0.37: Tweaked TES predictions on Test Set')
plt.plot(DES_predict, label='Alpha=0.65,Beta=0.0: DES predictions on Test Set')

plt.legend(loc='best')
plt.grid()
plt.title('TES, Tweaked TES & DES Predictions',color='blue',fontsize=20);
plt.xlabel('Years & Months',color='blue',fontsize=18);
plt.ylabel('Sparkling Sales',color='blue',fontsize=18);
```

TES, Tweaked TES & DES Predictions



```
In [47]: print('DES RMSE:', mean_squared_error(test.values, DES_predict.values, squared=False))
```

DES RMSE: 3851.2790161127123

```
In [48]: results_smoothing_2 = pd.DataFrame({'Test RMSE': [mean_squared_error(test.values, DES_predict.values, squared_difference=test.values - DES_predict.values),
                                                         ,index=['DES: Alpha=0.65,Beta=0.0']])
results = pd.concat([results, results_smoothing_2])
results
```

Out[48]:

| | Test RMSE |
|--|-------------|
| TES: Alpha=0.15, Beta=1.34, Gamma=0.37 | 392.932696 |
| TES_tweaked: Alpha=0.15, Beta=1.28e-21, Gamma=0.37 | 383.138464 |
| DES: Alpha=0.65,Beta=0.0 | 3851.279016 |

```
In [49]: # Using the Linear Regression model for forecasting:
```

```
In [50]: # Modifying the data to incorporate order against the sales values:
```

```
train_time = [i+1 for i in range(len(train))]
test_time = [i+133 for i in range(len(test))]
print('Training Time instance', '\n', train_time)
print('Test Time instance', '\n', test_time)
```

Training Time instance

[1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75, 76, 77, 78, 79, 80, 81, 82, 83, 84, 85, 86, 87, 88, 89, 90, 91, 92, 93, 94, 95, 96, 97, 98, 99, 100, 101, 102, 103, 104, 105, 106, 107, 108, 109, 110, 111, 112, 113, 114, 115, 116, 117, 118, 119, 120, 121, 122, 123, 124, 125, 126, 127, 128, 129, 130, 131, 132]

Test Time instance

[133, 134, 135, 136, 137, 138, 139, 140, 141, 142, 143, 144, 145, 146, 147, 148, 149, 150, 151, 152, 153, 154, 155, 156, 157, 158, 159, 160, 161, 162, 163, 164, 165, 166, 167, 168, 169, 170, 171, 172, 173, 174, 175, 176, 177, 178, 179, 180, 181, 182, 183, 184, 185, 186, 187]

```
In [51]: # Working on copies of Train & test data:
```

```
LinearRegression_train = train.copy()
LinearRegression_test = test.copy()
```

In [52]: *#Cross-checking the data:*

```
LinearRegression_train['time'] = train_time
LinearRegression_test['time'] = test_time

print('First few rows of Training Data')
display(LinearRegression_train.head())
print('Last few rows of Training Data')
display(LinearRegression_train.tail())
print('First few rows of Test Data')
display(LinearRegression_test.head())
print('Last few rows of Test Data')
display(LinearRegression_test.tail())
```

First few rows of Training Data

| Sparkling time | | |
|----------------|------|---|
| YearMonth | | |
| 1980-01-01 | 1686 | 1 |
| 1980-02-01 | 1591 | 2 |
| 1980-03-01 | 2304 | 3 |
| 1980-04-01 | 1712 | 4 |
| 1980-05-01 | 1471 | 5 |

Last few rows of Training Data

| Sparkling time | | |
|----------------|------|-----|
| YearMonth | | |
| 1990-08-01 | 1605 | 128 |
| 1990-09-01 | 2424 | 129 |
| 1990-10-01 | 3116 | 130 |

| Sparkling time | | |
|----------------|------|-----|
| YearMonth | | |
| 1990-11-01 | 4286 | 131 |
| 1990-12-01 | 6047 | 132 |

First few rows of Test Data

| Sparkling time | | |
|----------------|------|-----|
| YearMonth | | |
| 1991-01-01 | 1902 | 133 |
| 1991-02-01 | 2049 | 134 |
| 1991-03-01 | 1874 | 135 |
| 1991-04-01 | 1279 | 136 |
| 1991-05-01 | 1432 | 137 |

Last few rows of Test Data

| Sparkling time | | |
|----------------|------|-----|
| YearMonth | | |
| 1995-03-01 | 1897 | 183 |
| 1995-04-01 | 1862 | 184 |
| 1995-05-01 | 1670 | 185 |
| 1995-06-01 | 1688 | 186 |
| 1995-07-01 | 2031 | 187 |

In [53]: *#Building the LR model:*

```
from sklearn.linear_model import LinearRegression

lr = LinearRegression()
lr.fit(LinearRegression_train[['time']],LinearRegression_train['Sparkling'])
```

Out[53]: LinearRegression()

In [54]: *#Predicting values:*

```
train_predictions_lr = lr.predict(LinearRegression_train[['time']])
LinearRegression_train['LR_on_time'] = train_predictions_lr


test_predictions_lr = lr.predict(LinearRegression_test[['time']])
LinearRegression_test['LR_on_time'] = test_predictions_lr

plt.plot(train['Sparkling'], label='Train')
plt.plot(test['Sparkling'], label='Test')
plt.plot(TES_predict, label='Alpha=0.15, Beta=1.34, Gamma=0.37: TES predictions on Test Set')
plt.plot(TES_predict_tweaked, label='Alpha=0.15,Beta=1.28e-21,Gamma=0.37: Tweaked TES predictions on Test Set')
plt.plot(DES_predict, label='Alpha=0.65,Beta=0.0: DES predictions on Test Set')

plt.plot(LinearRegression_test['LR_on_time'], label='Regression On Time_Test Data')

plt.legend(loc='best')
plt.grid();
plt.title('TES, Tweaked TES, DES & LR Predictions',color='blue',fontsize=20);
plt.xlabel('Years & Months',color='blue',fontsize=18);
plt.ylabel('Sparkling Sales',color='blue',fontsize=18);
```

TES, Tweaked TES, DES & LR Predictions



```
In [55]: # Evaluating the model:  
print('LR RMSE:', mean_squared_error(test['Sparkling'], test_predictions_lr, squared=False))
```

LR RMSE: 1389.135174897992

```
In [56]: results_smoothing_3 = pd.DataFrame({'Test RMSE': [mean_squared_error(test['Sparkling'], test_predictions_
                                             ,index=[ 'LR RSME' ] )

results = pd.concat([results, results_smoothing_3])
results
```

Out[56]:

| | Test RMSE |
|---|-------------|
| TES: Alpha=0.15, Beta=1.34, Gamma=0.37 | 392.932696 |
| TES_tweaked: Alpha=0.15, Beta=1.28e-21, Gamma=0.37 | 383.138464 |
| DES: Alpha=0.65, Beta=0.0 | 3851.279016 |
| LR RSME | 1389.135175 |

```
In [57]: # Using the Naive Approach for forecasting:
```

```
In [58]: # Working on copies of Train & test data:
```

```
Naive_train = train.copy()
Naive_test = test.copy()
```

```
In [59]: train.head()
```

Out[59]:

| | Sparkling |
|-------------------|-----------|
| YearMonth | |
| 1980-01-01 | 1686 |
| 1980-02-01 | 1591 |
| 1980-03-01 | 2304 |
| 1980-04-01 | 1712 |
| 1980-05-01 | 1471 |

```
In [60]: train.tail()
```

Out[60]:

| Sparkling | |
|------------|------|
| YearMonth | |
| 1990-08-01 | 1605 |
| 1990-09-01 | 2424 |
| 1990-10-01 | 3116 |
| 1990-11-01 | 4286 |
| 1990-12-01 | 6047 |

```
In [61]: Naive_test['naive'] = np.asarray(train['Sparkling'])[len(np.asarray(train['Sparkling']))-1]
Naive_test['naive'].head()
```

Out[61]:

| YearMonth | |
|------------|------|
| 1991-01-01 | 6047 |
| 1991-02-01 | 6047 |
| 1991-03-01 | 6047 |
| 1991-04-01 | 6047 |
| 1991-05-01 | 6047 |

Name: naive, dtype: int64

```

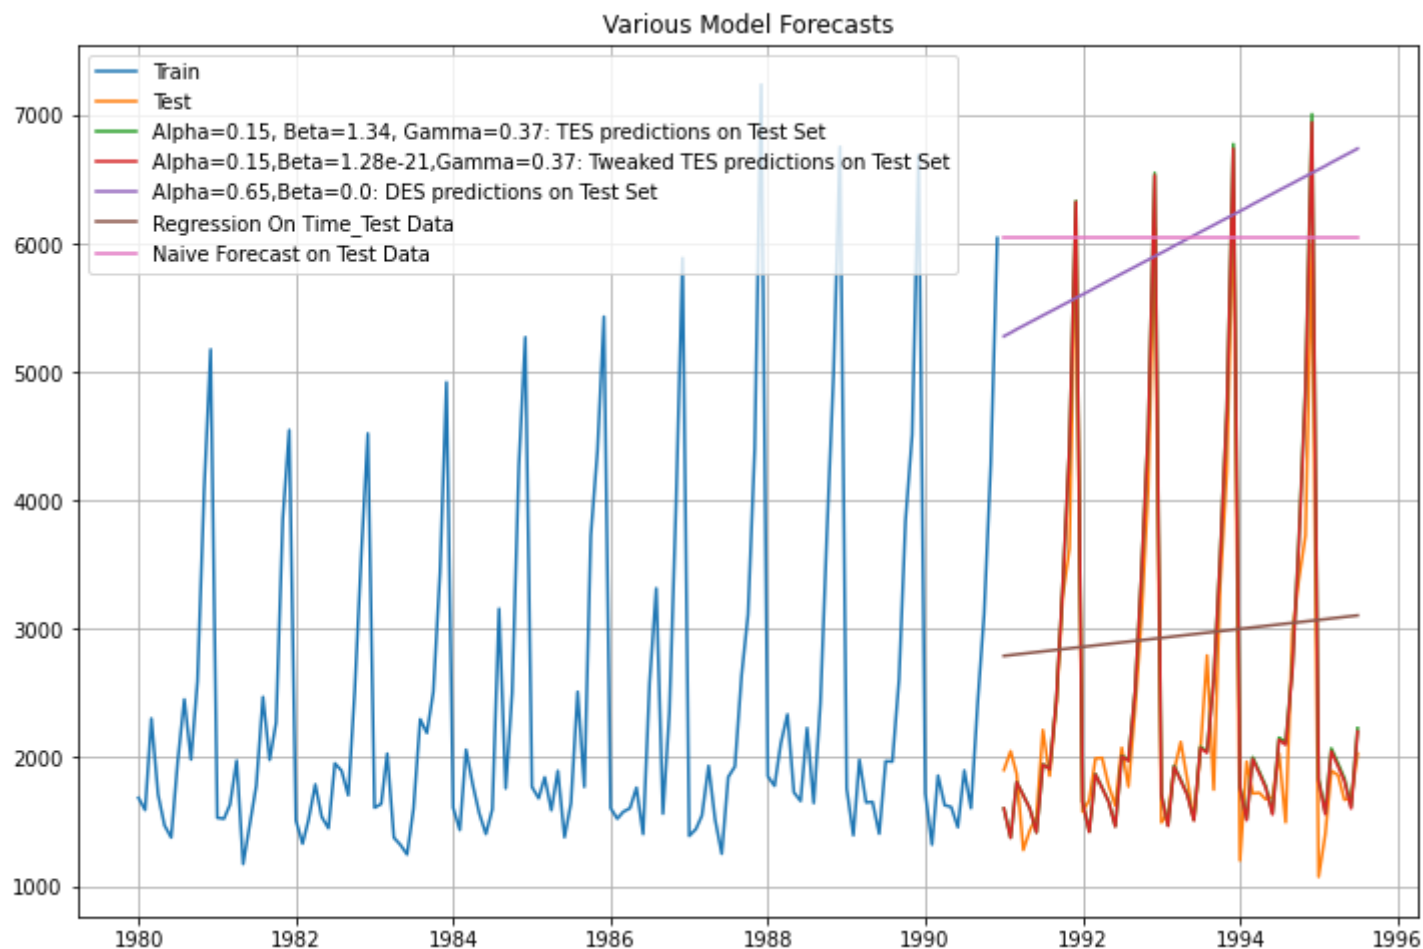
In [62]: plt.plot(Naive_train['Sparkling'], label='Train')
plt.plot(test['Sparkling'], label='Test')

plt.plot(TES_predict, label='Alpha=0.15, Beta=1.34, Gamma=0.37: TES predictions on Test Set')
plt.plot(TES_predict_tweaked, label='Alpha=0.15, Beta=1.28e-21, Gamma=0.37: Tweaked TES predictions on Test Set')
plt.plot(DEs_predict, label='Alpha=0.65, Beta=0.0: DES predictions on Test Set')
plt.plot(LinearRegression_test['LR_on_time'], label='Regression On Time_Test Data')

plt.plot(Naive_test['naive'], label='Naive Forecast on Test Data')

plt.legend(loc='best')
plt.title("Various Model Forecasts")
plt.grid();

```



```
In [63]: print('Naive RMSE:', mean_squared_error(test['Sparkling'], Naive_test['naive'], squared=False))
```

Naive RMSE: 3864.2793518443914

```
In [64]: results_smoothing_4 = pd.DataFrame({'Test RMSE': [mean_squared_error(test['Sparkling'], Naive_test['naive'], squared=False)],
                                             , index=['Naive RSME']})

results = pd.concat([results, results_smoothing_4])
results
```

Out[64]:

| | Test RMSE |
|---|-------------|
| TES: Alpha=0.15, Beta=1.34, Gamma=0.37 | 392.932696 |
| TES_tweaked: Alpha=0.15, Beta=1.28e-21, Gamma=0.37 | 383.138464 |
| DES: Alpha=0.65, Beta=0.0 | 3851.279016 |
| LR RSME | 1389.135175 |
| Naive RSME | 3864.279352 |

```
In [65]: # Using the Simple Average method:
```



```
In [66]: # Working on copies of Train & test data:
```

```
SA_train = train.copy()  
SA_test = test.copy()
```

```
In [67]: SA_test['mean_forecast'] = train['Sparkling'].mean()  
SA_test.head()
```

Out[67]:


| | Sparkling | mean_forecast |
|------------|-----------|---------------|
| YearMonth | | |
| 1991-01-01 | 1902 | 2403.780303 |
| 1991-02-01 | 2049 | 2403.780303 |
| 1991-03-01 | 1874 | 2403.780303 |
| 1991-04-01 | 1279 | 2403.780303 |
| 1991-05-01 | 1432 | 2403.780303 |

```
In [68]: plt.plot(SA_train['Sparkling'], label='Train')
plt.plot(SA_test['Sparkling'], label='Test')

plt.plot(TES_predict, label='Alpha=0.15, Beta=1.34, Gamma=0.37: TES predictions on Test Set')
plt.plot(TES_predict_tweaked, label='Alpha=0.15,Beta=1.28e-21,Gamma=0.37: Tweaked TES predictions on Test Set')
plt.plot(DESPredict, label='Alpha=0.65,Beta=0.0: DES predictions on Test Set')
plt.plot(LinearRegression_test['LR_on_time'], label='Regression On Time_Test Data')
plt.plot(Naive_test['naive'], label='Naive Forecast on Test Data')

plt.plot(SA_test['mean_forecast'], label='Simple Average on Test Data')

plt.legend(loc='best')
plt.title("Various Model Forecasts")
plt.grid();
```



```
In [69]: print('SA RMSE:', mean_squared_error(test['Sparkling'], SA_test['mean_forecast'], squared=False))
```

SA RMSE: 1275.0818036965309

```
In [70]: results_smoothing_5 = pd.DataFrame({'Test RMSE': [mean_squared_error(test['Sparkling'],SA_test['mean_for
                                     ,index=['SA RSME'])

results = pd.concat([results, results_smoothing_5])
results
```

Out[70]:

| | Test RMSE |
|---|-------------|
| TES: Alpha=0.15, Beta=1.34, Gamma=0.37 | 392.932696 |
| TES_tweaked: Alpha=0.15, Beta=1.28e-21, Gamma=0.37 | 383.138464 |
| DES: Alpha=0.65,Beta=0.0 | 3851.279016 |
| LR RSME | 1389.135175 |
| Naive RSME | 3864.279352 |
| SA RSME | 1275.081804 |

```
In [71]: #Using the Moving Average method on a copy of the original data:
```

```
MA = df.copy()
MA.head()
```

Out[71]:

| | Sparkling |
|-------------------|-----------|
| YearMonth | |
| 1980-01-01 | 1686 |
| 1980-02-01 | 1591 |
| 1980-03-01 | 2304 |
| 1980-04-01 | 1712 |
| 1980-05-01 | 1471 |

```
In [72]: MA['Trailing_2'] = MA['Sparkling'].rolling(2).mean()  
MA['Trailing_4'] = MA['Sparkling'].rolling(4).mean()  
MA['Trailing_6'] = MA['Sparkling'].rolling(6).mean()  
MA['Trailing_9'] = MA['Sparkling'].rolling(9).mean()  
  
MA.head(10)
```

Out[72]:

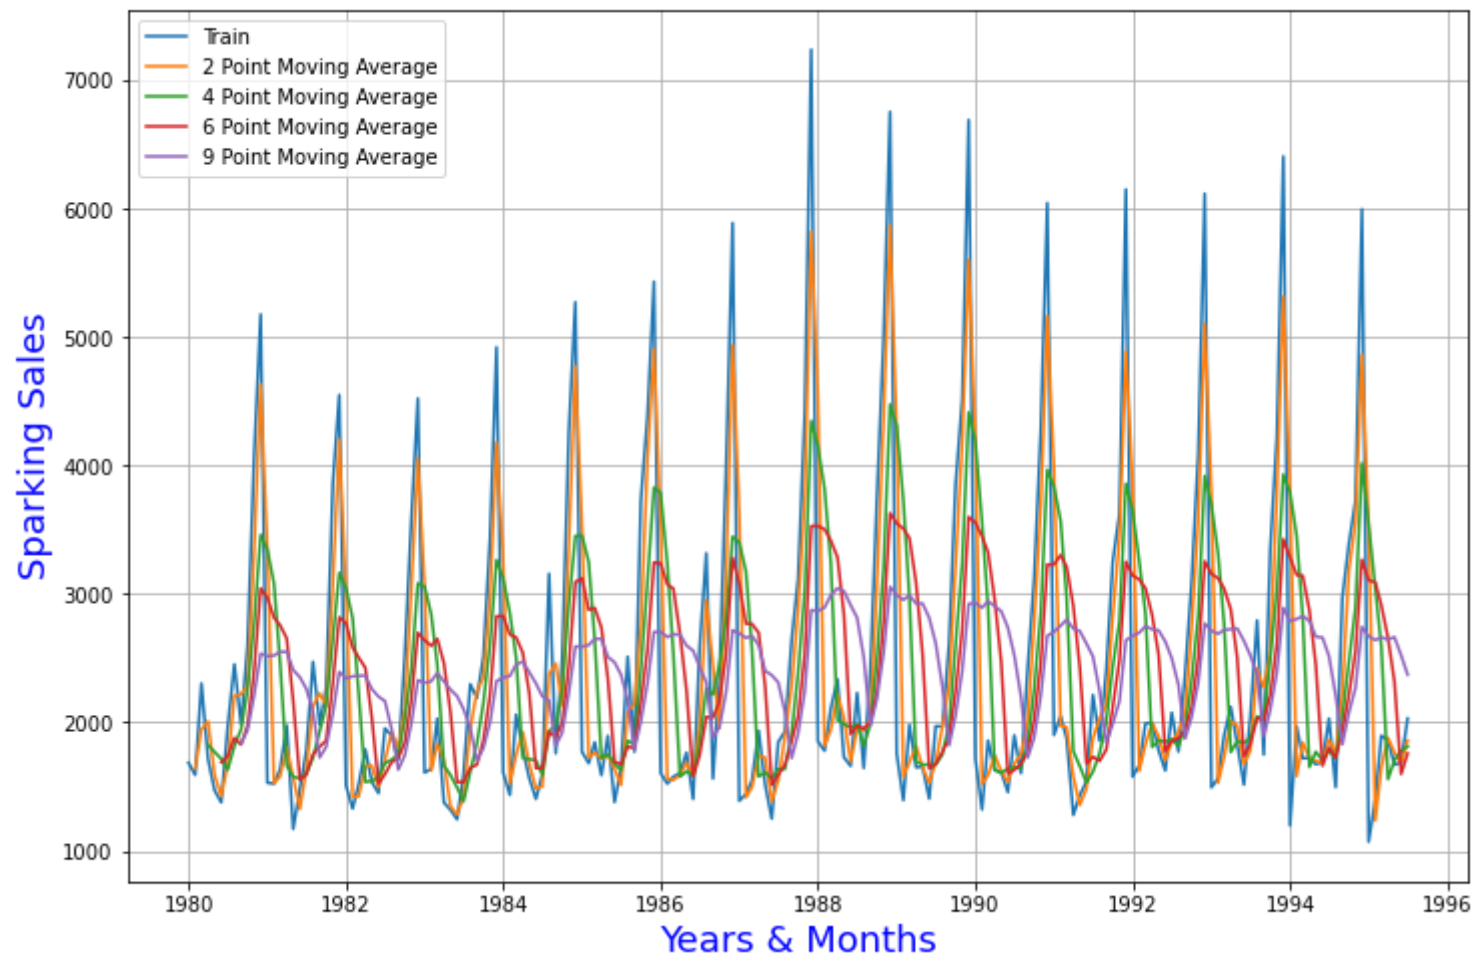
| | Sparkling | Trailing_2 | Trailing_4 | Trailing_6 | Trailing_9 |
|------------|-----------|------------|------------|-------------|-------------|
| YearMonth | | | | | |
| 1980-01-01 | 1686 | NaN | NaN | NaN | NaN |
| 1980-02-01 | 1591 | 1638.5 | NaN | NaN | NaN |
| 1980-03-01 | 2304 | 1947.5 | NaN | NaN | NaN |
| 1980-04-01 | 1712 | 2008.0 | 1823.25 | NaN | NaN |
| 1980-05-01 | 1471 | 1591.5 | 1769.50 | NaN | NaN |
| 1980-06-01 | 1377 | 1424.0 | 1716.00 | 1690.166667 | NaN |
| 1980-07-01 | 1966 | 1671.5 | 1631.50 | 1736.833333 | NaN |
| 1980-08-01 | 2453 | 2209.5 | 1816.75 | 1880.500000 | NaN |
| 1980-09-01 | 1984 | 2218.5 | 1945.00 | 1827.166667 | 1838.222222 |
| 1980-10-01 | 2596 | 2290.0 | 2249.75 | 1974.500000 | 1939.333333 |

In [73]: *# Plotting on the entire data:*

```
plt.plot(MA['Sparkling'], label='Train')
plt.plot(MA['Trailing_2'], label='2 Point Moving Average')
plt.plot(MA['Trailing_4'], label='4 Point Moving Average')
plt.plot(MA['Trailing_6'], label='6 Point Moving Average')
plt.plot(MA['Trailing_9'], label='9 Point Moving Average')

plt.legend(loc = 'best')
plt.grid();
plt.title('Various MA Forecasts on entire data',color='blue',fontsize=20);
plt.xlabel('Years & Months',color='blue',fontsize=18);
plt.ylabel('Sparkling Sales',color='blue',fontsize=18);
```

Various MA Forecasts on entire data



```
In [74]: #Creating train and test set for MA method:
```

```
trailing_MA_train = MA[MA.index<'1991']  
trailing_MA_test= MA[MA.index>='1991']
```

```
In [75]: trailing_MA_train.tail()
```

Out[75]:


| | Sparkling | Trailing_2 | Trailing_4 | Trailing_6 | Trailing_9 |
|------------|-----------|------------|------------|-------------|-------------|
| YearMonth | | | | | |
| 1990-08-01 | 1605 | 1752.0 | 1644.00 | 1677.166667 | 2199.777778 |
| 1990-09-01 | 2424 | 2014.5 | 1846.25 | 1771.333333 | 1725.333333 |
| 1990-10-01 | 3116 | 2770.0 | 2261.00 | 2019.333333 | 1880.444444 |
| 1990-11-01 | 4286 | 3701.0 | 2857.75 | 2464.500000 | 2209.888889 |
| 1990-12-01 | 6047 | 5166.5 | 3968.25 | 3229.500000 | 2675.222222 |

In [76]: *# Plotting on Test data:*

```
plt.figure(figsize=(16,8))
plt.plot(trailing_MA_train['Sparkling'], label='Train')
plt.plot(trailing_MA_test['Sparkling'], label='Test')

plt.plot(trailing_MA_test['Trailing_2'], label='2 Point Trailing Moving Average on Test Set')
plt.plot(trailing_MA_test['Trailing_4'], label='4 Point Trailing Moving Average on Test Set')
plt.plot(trailing_MA_test['Trailing_6'], label = '6 Point Trailing Moving Average on Test Set')
plt.plot(trailing_MA_test['Trailing_9'], label = '9 Point Trailing Moving Average on Test Set')

plt.legend(loc = 'best')
plt.grid();
plt.title('Various MA Forecasts on Test data',color='blue',fontsize=20);
plt.xlabel('Years & Months',color='blue',fontsize=18);
plt.ylabel('Sparkling Sales',color='blue',fontsize=18);
```



In [77]: *# Evaluating using RSME:*

```
from sklearn import metrics
```

```
# 2 point Trailing RSME
```

```
rmse_test_2 = metrics.mean_squared_error(test['Sparkling'],trailing_MA_test['Trailing_2'],squared=False)  
print("For 2 point MA Model, RMSE is %3.3f" %(rmse_test_2))
```

```
# 4 point Trailing RSME
```

```
rmse_test_4 = metrics.mean_squared_error(test['Sparkling'],trailing_MA_test['Trailing_4'],squared=False)  
print("For 4 point MA Model, RMSE is %3.3f" %(rmse_test_4))
```

```
# 6 point Trailing RSME
```

```
rmse_test_6 = metrics.mean_squared_error(test['Sparkling'],trailing_MA_test['Trailing_6'],squared=False)  
print("For 6 point MA Model, RMSE is %3.3f" %(rmse_test_6))
```

```
# 9 point Trailing RSME
```

```
rmse_test_9 = metrics.mean_squared_error(test['Sparkling'],trailing_MA_test['Trailing_9'],squared=False)  
print("For 9 point MA Model, RMSE is %3.3f" %(rmse_test_9))
```

```
For 2 point MA Model, RMSE is 813.401  
For 4 point MA Model, RMSE is 1156.590  
For 6 point MA Model, RMSE is 1283.927  
For 9 point MA Model, RMSE is 1346.278
```

In [78]:

```

results_smoothing_6 = pd.DataFrame({'Test RMSE': [rmse_test_2,rmse_test_4
                                                ,rmse_test_6,rmse_test_9]}
                                   ,index=['2-point MA','4-point MA','6-point MA','9-point MA'])

results = pd.concat([results, results_smoothing_6])
results

```

Out[78]:

| | Test RMSE |
|---|-------------|
| TES: Alpha=0.15, Beta=1.34, Gamma=0.37 | 392.932696 |
| TES_tweaked: Alpha=0.15, Beta=1.28e-21, Gamma=0.37 | 383.138464 |
| DES: Alpha=0.65,Beta=0.0 | 3851.279016 |
| LR RSME | 1389.135175 |
| Naive RSME | 3864.279352 |
| SA RSME | 1275.081804 |
| 2-point MA | 813.400684 |
| 4-point MA | 1156.589694 |
| 6-point MA | 1283.927428 |
| 9-point MA | 1346.278315 |

In [79]: *# Plotting the comparison of all model predictions:*


```
plt.figure(figsize=(16,8))
plt.plot(train['Sparkling'], label='Train')
plt.plot(test['Sparkling'], label='Test')

plt.plot(TES_predict, label='Alpha=0.15, Beta=1.34, Gamma=0.37: TES predictions on Test Set')
plt.plot(TES_predict_tweaked, label='Alpha=0.15,Beta=1.28e-21,Gamma=0.37: Tweaked TES predictions on Test Set')
plt.plot(DESPredict, label='Alpha=0.65,Beta=0.0: DES predictions on Test Set')
plt.plot(LinearRegression_test['LR_on_time'], label='Regression On Time_Test Data')
plt.plot(Naive_test['naive'], label='Naive Forecast on Test Data')
plt.plot(SA_test['mean_forecast'], label='Simple Average on Test Data')

plt.plot(trailing_MA_test['Trailing_2'], label='2-Point Trailing MA on Training data')

plt.legend(loc='best')
plt.title("Model Comparison Plots")
plt.grid();
plt.title('Model Comparison Plots',color='blue',fontsize=20);
plt.xlabel('Years & Months',color='blue',fontsize=18);
plt.ylabel('Sparkling Sales',color='blue',fontsize=18);
```

Model Comparison Plots



```
In [80]: # Checking for stationarity of data using ADF test:
```

```
In [81]: from statsmodels.tsa.stattools import adfuller
```

```
In [82]: stat_test = adfuller(df, regression='ct')
print('Test statistic is %3.3f' % stat_test[0])
print('Test p-value is' , stat_test[1])
print('Number of lags used' , stat_test[2])
```

```
Test statistic is -1.798
Test p-value is 0.7055958459932584
Number of lags used 12
```

```
In [83]: # P-value > alpha, so data is non-stationary. Using level-1 differencing to make data stationary:
```

```
In [84]: stat_test = adfuller(df.diff().dropna(), regression='ct')
print('Test statistic is %3.3f' % stat_test[0])
print('Test p-value is' , stat_test[1])
print('Number of lags used' , stat_test[2])
```

Test statistic is -44.912


Test p-value is 0.0

Number of lags used 10

```
In [85]: # Now data is stationary. Plotting the differenced data:
```

```
In [86]: df.diff().dropna().plot(grid=True);  
plt.title('Sparkling Sales Distribution (Post Differencing)',color='blue',fontsize=16)  
plt.xlabel('Year/Month',color='blue',fontsize=14)  
plt.ylabel('Sparkling Sales',color='blue',fontsize=14)
```


```
Out[86]: Text(0, 0.5, 'Sparkling Sales')
```





```
In [87]: # Plotting the autocorrelation and partial autocorrelation plots on data:
```

```
In [88]: from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

```
In [89]: plot_acf(df,alpha=0.05);
```



```
In [90]: plot_pacf(df, zero=False, alpha=0.05);
```




```
In [91]: #Splitting data to build the models:  
train = df[df.index<'1991']  
test = df[df.index>='1991']
```

In [92]: *#Plotting the train data:*

```
train.plot(grid=True);  
plt.title('Train data',color='blue',fontsize=16)  
plt.xlabel('Year/Month',color='blue',fontsize=14)  
plt.ylabel('Sparkling Sales',color='blue',fontsize=14)
```

Out[92]: Text(0, 0.5, 'Sparkling Sales')



In [93]: *#Checking for stationarity of the train data:*

```
stat_test = adfuller(train,regression='ct')
print('Test statistic is %3.3f' %stat_test[0])
print('Test p-value is' ,stat_test[1])
print('Number of lags used' ,stat_test[2])
```

Test statistic is -2.062
Test p-value is 0.56741103885937
Number of lags used 12

In [94]: *# P-value > alpha, so data is non-stationary. Using level-1 differencing to make data stationary:*

In [95]:


```
stat_test = adfuller(train.diff().dropna(),regression='ct')
print('Test statistic is %3.3f' %stat_test[0])
print('Test p-value is' ,stat_test[1])
print('Number of lags used' ,stat_test[2])
```

Test statistic is -7.968
Test p-value is 8.479210655513744e-11
Number of lags used 11

In [96]: *#Plotting the differenced train data:*

```
train.diff().dropna().plot(grid=True);  
plt.title('Train data (Post Differencing)',color='blue',fontsize=16)  
plt.xlabel('Year/Month',color='blue',fontsize=14)  
plt.ylabel('Sparkling Sales',color='blue',fontsize=14)
```

Out[96]: Text(0, 0.5, 'Sparkling Sales')



```
In [97]: # Since there is seasonality in the data set, we will build SARIMA model:  
        ## Building automated version of SARIMA:
```

```
In [98]: import itertools
p = q = range(0, 4)
d= range(1,2)
D = range(0,1)
pdq = list(itertools.product(p, d, q))
PDQ = [(x[0], x[1], x[2], 12) for x in list(itertools.product(p, D, q))]
print('Examples of the parameter combinations for the Model are')
for i in range(1,len(pdq)):
    print('Model: {}'.format(pdq[i], PDQ[i]))
```

Examples of the parameter combinations for the Model are

```
Model: (0, 1, 1)(0, 0, 1, 12)
Model: (0, 1, 2)(0, 0, 2, 12)
Model: (0, 1, 3)(0, 0, 3, 12)
Model: (1, 1, 0)(1, 0, 0, 12)
Model: (1, 1, 1)(1, 0, 1, 12)
Model: (1, 1, 2)(1, 0, 2, 12)
Model: (1, 1, 3)(1, 0, 3, 12)
Model: (2, 1, 0)(2, 0, 0, 12)
Model: (2, 1, 1)(2, 0, 1, 12)
Model: (2, 1, 2)(2, 0, 2, 12)
Model: (2, 1, 3)(2, 0, 3, 12)
Model: (3, 1, 0)(3, 0, 0, 12)
Model: (3, 1, 1)(3, 0, 1, 12)
Model: (3, 1, 2)(3, 0, 2, 12)
Model: (3, 1, 3)(3, 0, 3, 12)
```

```
In [99]: #Creating dataframe for storing AIC values:
SARIMA_AIC = pd.DataFrame(columns=['param', 'seasonal', 'AIC'])
SARIMA_AIC
```

Out[99]:

| param | seasonal | AIC |
|-------|----------|-----|
|-------|----------|-----|


```
In [100]: import statsmodels.api as sm

for param in pdq:
    for param_seasonal in PDQ:
        SARIMA_model = sm.tsa.statespace.SARIMAX(train['Sparkling'].values,
                                                  order=param,
                                                  seasonal_order=param_seasonal,
                                                  enforce_stationarity=False,
                                                  enforce_invertibility=False)

        results_SARIMA = SARIMA_model.fit(maxiter=1000)
        print('SARIMA{}x{} - AIC:{}'.format(param, param_seasonal, results_SARIMA.aic))
        SARIMA_AIC = SARIMA_AIC.append({'param':param, 'seasonal':param_seasonal, 'AIC': results_SARIMA.aic})
```

```
SARIMA(0, 1, 0)x(0, 0, 0, 12) - AIC:2251.3597196862966
SARIMA(0, 1, 0)x(0, 0, 1, 12) - AIC:1956.2614616843318
SARIMA(0, 1, 0)x(0, 0, 2, 12) - AIC:1723.1533640234898
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/base/model.py:567: ConvergenceWarning:

Maximum Likelihood optimization failed to converge. Check mle_retvals

```
SARIMA(0, 1, 0)x(0, 0, 3, 12) - AIC:2763.334195884304
SARIMA(0, 1, 0)x(1, 0, 0, 12) - AIC:1837.4366022456675
SARIMA(0, 1, 0)x(1, 0, 1, 12) - AIC:1806.990530137321
SARIMA(0, 1, 0)x(1, 0, 2, 12) - AIC:1633.2108735940249
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/base/model.py:567: ConvergenceWarning:

Maximum Likelihood optimization failed to converge. Check mle_retvals

```
In [101]: # Arranging models in ascending order to select least-AIC model:  
SARIMA_AIC.sort_values(by=[ 'AIC' ] ).head( )
```

Out[101]:

| | param | seasonal | AIC |
|------------|-----------|---------------|-------------|
| 236 | (3, 1, 2) | (3, 0, 0, 12) | 1387.234721 |
| 220 | (3, 1, 1) | (3, 0, 0, 12) | 1387.788333 |
| 237 | (3, 1, 2) | (3, 0, 1, 12) | 1388.602615 |
| 221 | (3, 1, 1) | (3, 0, 1, 12) | 1388.681480 |
| 252 | (3, 1, 3) | (3, 0, 0, 12) | 1389.142191 |

In [102]: *#Using the least-AIC model for SARIMAX computation:*

```
import statsmodels.api as sm

auto_SARIMA = sm.tsa.statespace.SARIMAX(train['Sparkling'],
                                         order=(3, 1, 2),
                                         seasonal_order=(3, 0, 0, 12),
                                         enforce_stationarity=False,
                                         enforce_invertibility=False)
results_auto_SARIMA = auto_SARIMA.fit(maxiter=1000)
print(results_auto_SARIMA.summary())
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

SARIMAX Results

```
=====
Dep. Variable:          Sparkling      No. Observations:          132
Model:                SARIMAX(3, 1, 2)x(3, 0, [], 12)  Log Likelihood          -684.617
Date:                  Sat, 09 Oct 2021      AIC                  1387.235
Time:                  07:01:11             BIC                  1409.931
Sample:                01-01-1980           HQIC                 1396.395
                    - 12-01-1990
Covariance Type:                opg
=====
```

| | coef | std err | z | P> z | [0.025 | 0.975] |
|----------|---------|---------|--------|-------|--------|--------|
| ar.L1 | -0.5372 | 0.339 | -1.586 | 0.113 | -1.201 | 0.126 |
| ar.L2 | 0.0256 | 0.187 | 0.137 | 0.891 | -0.340 | 0.392 |
| ar.L3 | 0.0785 | 0.130 | 0.604 | 0.546 | -0.176 | 0.333 |
| ma.L1 | -0.1878 | 0.326 | -0.575 | 0.565 | -0.828 | 0.452 |
| ma.L2 | -0.6875 | 0.272 | -2.531 | 0.011 | -1.220 | -0.155 |
| ar.S.L12 | 0.5713 | 0.103 | 5.542 | 0.000 | 0.369 | 0.773 |
| ar.S.L24 | 0.2605 | 0.117 | 2.222 | 0.026 | 0.031 | 0.490 |
| ar.S.L36 | 0.2126 | 0.111 | 1.916 | 0.055 | -0.005 | 0.430 |

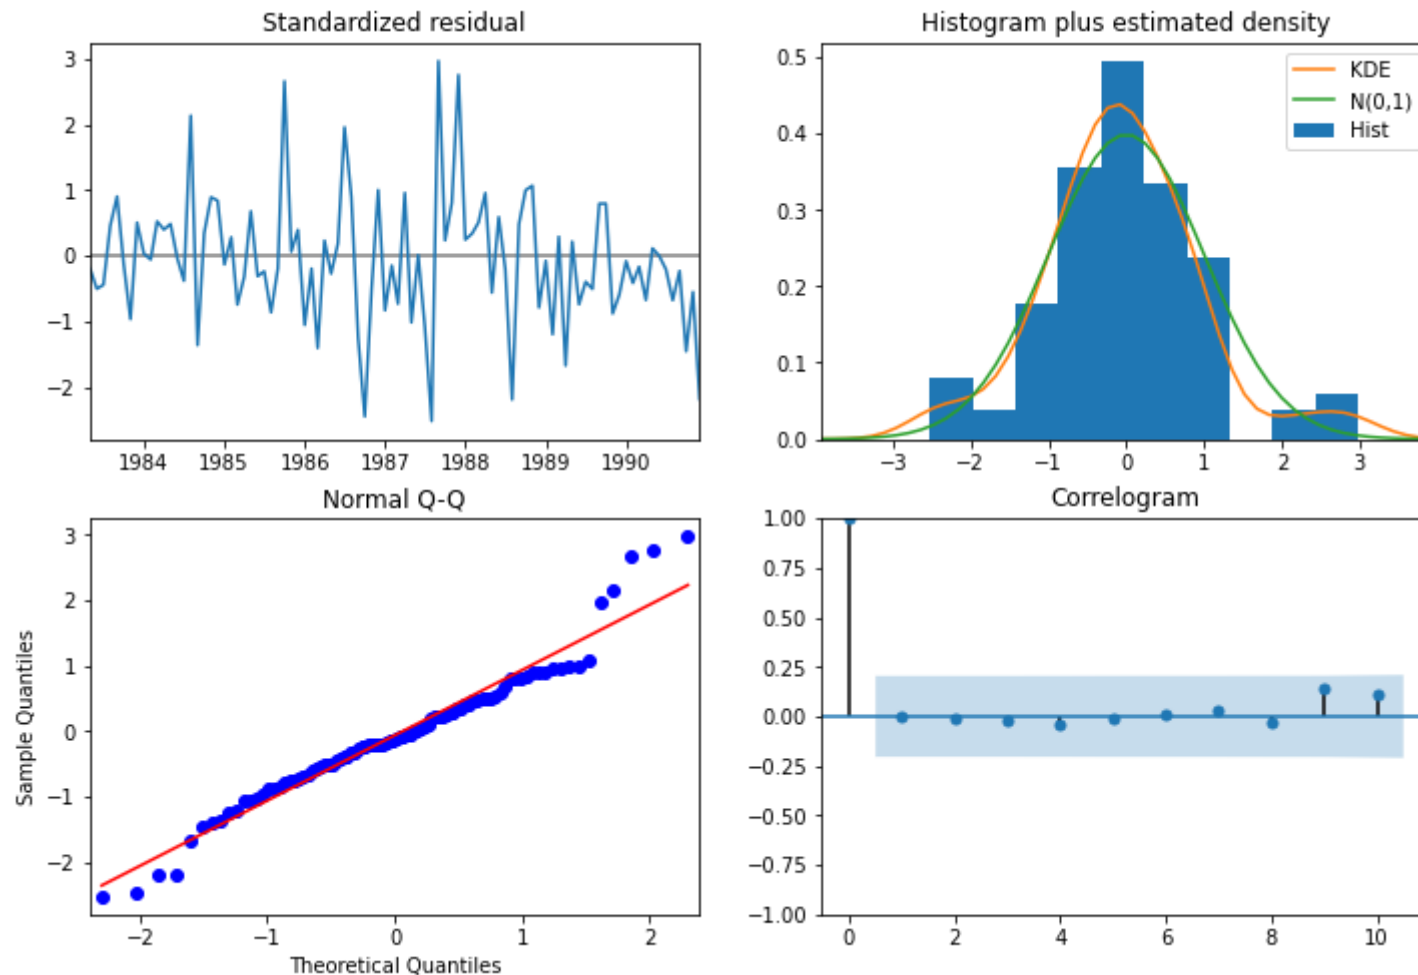
| | | | | | | |
|--------|-----------|----------|-------|-------|----------|----------|
| sigma2 | 1.682e+05 | 2.52e+04 | 6.671 | 0.000 | 1.19e+05 | 2.18e+05 |
|--------|-----------|----------|-------|-------|----------|----------|

| | | | |
|-------------------------|-------|-------------------|------|
| Ljung-Box (Q): | 27.31 | Jarque-Bera (JB): | 8.81 |
| Prob(Q): | 0.94 | Prob(JB): | 0.01 |
| Heteroskedasticity (H): | 1.17 | Skew: | 0.36 |
| Prob(H) (two-sided): | 0.67 | Kurtosis: | 4.33 |

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [103]: results_auto_SARIMA.plot_diagnostics();
```



In [104]: *# Using SARIMA model to predict test set:*

```
predicted_auto_SARIMA = results_auto_SARIMA.get_forecast(steps=len(test))
```

In [105]: *# Defining Mean Absolute Percentage Error (MAPE):*

```
def mean_absolute_percentage_error(y_true, y_pred):
    return np.mean((np.abs(y_true-y_pred))/(y_true))*100
```

```
from sklearn.metrics import mean_squared_error
```

In [106]: *# Evaluating the predictions:*

```
rmse = mean_squared_error(test['Sparkling'],predicted_auto_SARIMA.predicted_mean,squared=False)
mape = mean_absolute_percentage_error(test['Sparkling'],predicted_auto_SARIMA.predicted_mean)
print('RMSE:',rmse,'\nMAPE:',mape)
```

```
RMSE: 543.0401637640175
```

```
MAPE: 23.232450291356987
```

In [107]: *# Storing results for comparison:*


```
results_models = pd.DataFrame({'Test RMSE': rmse,'MAPE':mape}
                               ,index=['SARIMA(3,1,2)(3,0,0,12)'])
```


```
results_models
```

Out[107]:

| | Test RMSE | MAPE |
|--------------------------------|------------|----------|
| SARIMA(3,1,2)(3,0,0,12) | 543.040164 | 23.23245 |

```
In [108]: # Building a manual SARIMA model by selecting values of p, q from correlation plots:  
  
plot_acf(train.diff(),title='Training Data Autocorrelation',missing='drop')  
plot_pacf(train.diff().dropna(),title='Training Data Partial Autocorrelation',zero=False)  
plt.show()
```





```
In [109]: # As per the ACF and PACF plots, we will take the values as p=0, q=0, d=1 and seasonal componenets P=0,
```

```
In [110]: import statsmodels.api as sm
```

```
manual_SARIMA = sm.tsa.statespace.SARIMAX(train['Sparkling'],
                                           order=(0,1,0),
                                           seasonal_order=(0, 0, 0, 12),
                                           enforce_stationarity=False,
                                           enforce_invertibility=False)
results_manual_SARIMA = manual_SARIMA.fit(maxiter=1000)
print(results_manual_SARIMA.summary())
```

SARIMAX Results

```
=====
Dep. Variable:          Sparkling    No. Observations:          132
Model:                SARIMAX(0, 1, 0)    Log Likelihood          -1124.680
Date:                Sat, 09 Oct 2021    AIC                    2251.360
Time:                07:01:13    BIC                    2254.227
Sample:                01-01-1980    HQIC                   2252.525
                    - 12-01-1990
Covariance Type:                opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
sigma2      1.899e+06   1.31e+05    14.543     0.000    1.64e+06    2.16e+06
=====
Ljung-Box (Q):                345.63    Jarque-Bera (JB):                194.29
Prob(Q):                      0.00    Prob(JB):                      0.00
Heteroskedasticity (H):        2.46    Skew:                          -1.92
Prob(H) (two-sided):           0.00    Kurtosis:                      7.60
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

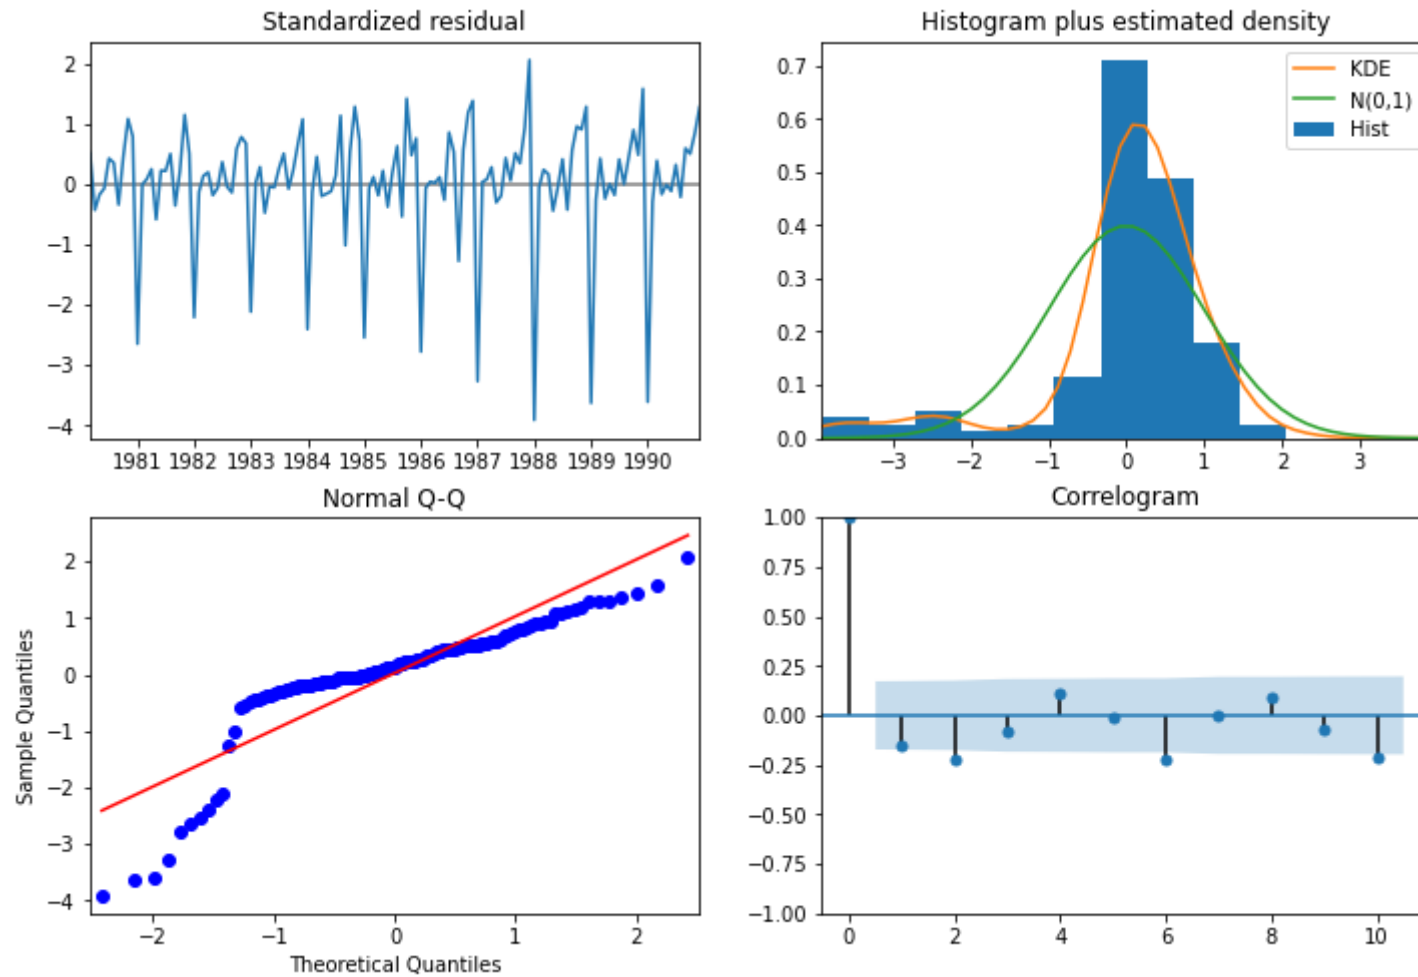
/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

```
In [111]: results_manual_SARIMA.plot_diagnostics()  
plt.show()
```



```
In [112]: predicted_manual_SARIMA = results_manual_SARIMA.get_forecast(steps=len(test))
```

```
In [113]: rmse = mean_squared_error(test['Sparkling'],predicted_manual_SARIMA.predicted_mean,squared=False)
mape = mean_absolute_percentage_error(test['Sparkling'],predicted_manual_SARIMA.predicted_mean)
print('RMSE:',rmse,'\nMAPE:',mape)
```

```
RMSE: 3864.2793518443914
MAPE: 201.32764950352743
```

```
In [114]: results_1 = pd.DataFrame({'Test RMSE': [rmse], 'MAPE':mape}
                                   ,index=['SARIMA(0,1,0)(0,0,0,12)'])
results_models = pd.concat([results_models,results_1])
results_models
```

Out[114]:

| | Test RMSE | MAPE |
|--------------------------------|-------------|-----------|
| SARIMA(3,1,2)(3,0,0,12) | 543.040164 | 23.23245 |
| SARIMA(0,1,0)(0,0,0,12) | 3864.279352 | 201.32765 |

In []:

```
In [115]: # Trying different parameters:
```

```
In [116]: manual_SARIMA = sm.tsa.statespace.SARIMAX(train['Sparkling'],
                                                    order=(1,1,1),
                                                    seasonal_order=(0, 0, 0, 12),
                                                    enforce_stationarity=False,
                                                    enforce_invertibility=False)
results_manual_SARIMA = manual_SARIMA.fit(maxiter=1000)
print(results_manual_SARIMA.summary())
```

SARIMAX Results

```
=====
Dep. Variable:          Sparkling    No. Observations:          132
Model:                SARIMAX(1, 1, 1)    Log Likelihood          -1099.467
Date:                 Sat, 09 Oct 2021    AIC                    2204.934
Time:                 07:01:14           BIC                    2213.513
Sample:               01-01-1980         HQIC                   2208.420
                  - 12-01-1990
Covariance Type:      opg
=====
```

| | coef | std err | z | P> z | [0.025 | 0.975] |
|--------|-----------|----------|---------|-------|----------|----------|
| ar.L1 | 0.4323 | 0.106 | 4.074 | 0.000 | 0.224 | 0.640 |
| ma.L1 | -0.9865 | 0.080 | -12.294 | 0.000 | -1.144 | -0.829 |
| sigma2 | 1.756e+06 | 2.14e+05 | 8.215 | 0.000 | 1.34e+06 | 2.17e+06 |

```
=====
Ljung-Box (Q):          343.21    Jarque-Bera (JB):          11.75
Prob(Q):                0.00     Prob(JB):                0.00
Heteroskedasticity (H):  2.69     Skew:                    0.55
Prob(H) (two-sided):    0.00     Kurtosis:                4.00
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

```
In [117]: predicted_manual_SARIMA = results_manual_SARIMA.get_forecast(steps=len(test))

rmse = mean_squared_error(test['Sparkling'],predicted_manual_SARIMA.predicted_mean,squared=False)
mape = mean_absolute_percentage_error(test['Sparkling'],predicted_manual_SARIMA.predicted_mean)
print('RMSE:',rmse,'\nMAPE:',mape)
```

RMSE: 1325.3488296929725

MAPE: 46.401952104149196

In []:

```
In [118]: results_2 = pd.DataFrame({'Test RMSE': [rmse], 'MAPE':mape}
                                   ,index=['SARIMA(1,1,1)(0,0,0,12)'])
results_models = pd.concat([results_models,results_2])
results_models
```

Out[118]:

| | Test RMSE | MAPE |
|--------------------------------|-------------|------------|
| SARIMA(3,1,2)(3,0,0,12) | 543.040164 | 23.232450 |
| SARIMA(0,1,0)(0,0,0,12) | 3864.279352 | 201.327650 |
| SARIMA(1,1,1)(0,0,0,12) | 1325.348830 | 46.401952 |

```
In [119]: manual_SARIMA = sm.tsa.statespace.SARIMAX(train['Sparkling'],
                                                    order=(3,1,2),
                                                    seasonal_order=(0, 0, 1, 12),
                                                    enforce_stationarity=False,
                                                    enforce_invertibility=False)
results_manual_SARIMA = manual_SARIMA.fit(maxiter=1000)
print(results_manual_SARIMA.summary())
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

SARIMAX Results

```
=====
Dep. Variable:          Sparkling      No. Observations:          132
Model:                SARIMAX(3, 1, 2)x(0, 0, [1], 12)      Log Likelihood          -935.568
Date:                  Sat, 09 Oct 2021      AIC          1885.137
Time:                  07:01:16      BIC          1904.412
Sample:                01-01-1980      HQIC          1892.961
                    - 12-01-1990
Covariance Type:                opg
=====
```

| | coef | std err | z | P> z | [0.025 | 0.975] |
|----------|-----------|----------|----------|-------|----------|----------|
| ar.L1 | 1.1988 | 0.116 | 10.309 | 0.000 | 0.971 | 1.427 |
| ar.L2 | -0.4623 | 0.178 | -2.604 | 0.009 | -0.810 | -0.114 |
| ar.L3 | -0.0446 | 0.117 | -0.383 | 0.702 | -0.273 | 0.184 |
| ma.L1 | -1.9965 | 0.148 | -13.503 | 0.000 | -2.286 | -1.707 |
| ma.L2 | 0.9888 | 0.150 | 6.604 | 0.000 | 0.695 | 1.282 |
| ma.S.L12 | 0.7484 | 0.088 | 8.540 | 0.000 | 0.577 | 0.920 |
| sigma2 | 4.683e+05 | 6.55e-07 | 7.15e+11 | 0.000 | 4.68e+05 | 4.68e+05 |

```
=====
Ljung-Box (Q):          151.18      Jarque-Bera (JB):          9.80
Prob(Q):                0.00      Prob(JB):                0.01
Heteroskedasticity (H):  2.91      Skew:                    0.59
Prob(H) (two-sided):    0.00      Kurtosis:                3.80
=====
```

=====

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

[2] Covariance matrix is singular or near-singular, with condition number 2.73e+27. Standard errors may be unstable.

```
In [120]: predicted_manual_SARIMA = results_manual_SARIMA.get_forecast(steps=len(test))

rmse = mean_squared_error(test['Sparkling'], predicted_manual_SARIMA.predicted_mean, squared=False)
mape = mean_absolute_percentage_error(test['Sparkling'], predicted_manual_SARIMA.predicted_mean)
print('RMSE:', rmse, '\nMAPE:', mape)
```

RMSE: 1206.6591167833199

MAPE: 42.20399327401117

```
In [121]: results_3 = pd.DataFrame({'Test RMSE': [rmse], 'MAPE': mape}
                                   , index=['SARIMA(3,1,2)(0,0,1,12)'])
results_models = pd.concat([results_models, results_3])
results_models
```

Out[121]:

| | Test RMSE | MAPE |
|--------------------------------|-------------|------------|
| SARIMA(3,1,2)(3,0,0,12) | 543.040164 | 23.232450 |
| SARIMA(0,1,0)(0,0,0,12) | 3864.279352 | 201.327650 |
| SARIMA(1,1,1)(0,0,0,12) | 1325.348830 | 46.401952 |
| SARIMA(3,1,2)(0,0,1,12) | 1206.659117 | 42.203993 |

```
In [122]: manual_SARIMA = sm.tsa.statespace.SARIMAX(train['Sparkling'],
                                                    order=(3,1,2),
                                                    seasonal_order=(0, 0, 1, 6),
                                                    enforce_stationarity=False,
                                                    enforce_invertibility=False)
results_manual_SARIMA = manual_SARIMA.fit(maxiter=1000)
print(results_manual_SARIMA.summary())
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

SARIMAX Results

```
=====
Dep. Variable:          Sparkling    No. Observations:          132
Model:                SARIMAX(3, 1, 2)x(0, 0, [1], 6)    Log Likelihood          -1030.311
Date:                  Sat, 09 Oct 2021    AIC              2074.622
Time:                  07:01:17    BIC              2094.250
Sample:                01-01-1980    HQIC             2082.594
                    - 12-01-1990
Covariance Type:                opg
=====
```

| | coef | std err | z | P> z | [0.025 | 0.975] |
|---------|-----------|----------|----------|-------|----------|----------|
| ar.L1 | -0.4444 | 0.125 | -3.554 | 0.000 | -0.690 | -0.199 |
| ar.L2 | 0.2271 | 0.133 | 1.713 | 0.087 | -0.033 | 0.487 |
| ar.L3 | -0.3279 | 0.174 | -1.886 | 0.059 | -0.669 | 0.013 |
| ma.L1 | -0.0246 | 0.205 | -0.120 | 0.904 | -0.426 | 0.377 |
| ma.L2 | -1.0247 | 0.153 | -6.686 | 0.000 | -1.325 | -0.724 |
| ma.S.L6 | -0.1899 | 0.193 | -0.983 | 0.325 | -0.568 | 0.189 |
| sigma2 | 1.153e+06 | 3.04e-07 | 3.79e+12 | 0.000 | 1.15e+06 | 1.15e+06 |

```
=====
Ljung-Box (Q):          294.38    Jarque-Bera (JB):          7.48
Prob(Q):                0.00    Prob(JB):                0.02
Heteroskedasticity (H):    2.50    Skew:                    0.57
Prob(H) (two-sided):      0.00    Kurtosis:                 3.41
=====
```

=====

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

[2] Covariance matrix is singular or near-singular, with condition number 1.6e+29. Standard errors may be unstable.

```
In [123]: predicted_manual_SARIMA = results_manual_SARIMA.get_forecast(steps=len(test))

rmse = mean_squared_error(test['Sparkling'], predicted_manual_SARIMA.predicted_mean, squared=False)
mape = mean_absolute_percentage_error(test['Sparkling'], predicted_manual_SARIMA.predicted_mean)
print('RMSE:', rmse, '\nMAPE:', mape)
```

RMSE: 1268.7392591632458

MAPE: 43.54397235460573

```
In [124]: results_4 = pd.DataFrame({'Test RMSE': [rmse], 'MAPE': mape}
                                   , index=['SARIMA(3,1,2)(0,0,1,6)'])
results_models = pd.concat([results_models, results_4])
results_models
```

Out[124]:

| | Test RMSE | MAPE |
|--------------------------------|-------------|------------|
| SARIMA(3,1,2)(3,0,0,12) | 543.040164 | 23.232450 |
| SARIMA(0,1,0)(0,0,0,12) | 3864.279352 | 201.327650 |
| SARIMA(1,1,1)(0,0,0,12) | 1325.348830 | 46.401952 |
| SARIMA(3,1,2)(0,0,1,12) | 1206.659117 | 42.203993 |
| SARIMA(3,1,2)(0,0,1,6) | 1268.739259 | 43.543972 |

In [125]: *#Building a comparison table for comparing all the different models built:*

```
frames=[results,results_models]
result=pd.concat(frames)
result
```

Out[125]:

| | Test RMSE | MAPE |
|---|-------------|------------|
| TES: Alpha=0.15, Beta=1.34, Gamma=0.37 | 392.932696 | NaN |
| TES_tweaked: Alpha=0.15, Beta=1.28e-21, Gamma=0.37 | 383.138464 | NaN |
| DES: Alpha=0.65,Beta=0.0 | 3851.279016 | NaN |
| LR RSME | 1389.135175 | NaN |
| Naive RSME | 3864.279352 | NaN |
| SA RSME | 1275.081804 | NaN |
| 2-point MA | 813.400684 | NaN |
| 4-point MA | 1156.589694 | NaN |
| 6-point MA | 1283.927428 | NaN |
| 9-point MA | 1346.278315 | NaN |
| SARIMA(3,1,2)(3,0,0,12) | 543.040164 | 23.232450 |
| SARIMA(0,1,0)(0,0,0,12) | 3864.279352 | 201.327650 |
| SARIMA(1,1,1)(0,0,0,12) | 1325.348830 | 46.401952 |
| SARIMA(3,1,2)(0,0,1,12) | 1206.659117 | 42.203993 |
| SARIMA(3,1,2)(0,0,1,6) | 1268.739259 | 43.543972 |

In [126]: *#Comparing RSME of all models, we can go with TES model as best model for time series forecasting:*

In [127]: *#Building the final model using TES model:*

```
model_final = ExponentialSmoothing(df,trend='multiplicative',seasonal='multiplicative')

# Fitting the model
model_final = model_final.fit()

print('')
print('~~~ Holt Winters model Exponential Smoothing Estimated Parameters ~~~')
print('')
print(model_final.params)
```

~~~ Holt Winters model Exponential Smoothing Estimated Parameters ~~~

```
{'smoothing_level': 0.061258628360488704, 'smoothing_slope': 0.06123982807099429, 'smoothing_seasonal': 0.27190057291568354, 'damping_slope': nan, 'initial_level': 1580.0000060857483, 'initial_slope': 0.9964964857441378, 'initial_seasons': array([1.06255314, 1.00343217, 1.44340795, 1.0857587 , 0.9298955 ,
      0.87297234, 1.24270319, 1.55334083, 1.2574631 , 1.64853144,
      2.58866164, 3.28435007]), 'use_boxcox': False, 'lamda': None, 'remove_bias': False}
```


/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

```
In [136]: df.plot()  
model_final.forecast(steps=12).plot()  
  
plt.legend(['Observed series', 'Forecasted series'])  
plt.title('Forecasted Sparkling Sales (TES Model)', color='blue', fontsize=20)  
plt.xlabel('Year/Month', color='blue', fontsize=18)  
plt.ylabel('Sparkling Sales', color='blue', fontsize=18)  
plt.grid();
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:342: FutureWarning:

The 'freq' argument in Timestamp is deprecated and will be removed in a future version.



In [129]: *# Forecasting using SARIMA(3,1,2)(3,0,0,12):*

```
full_data_model = sm.tsa.statespace.SARIMAX(df['Sparkling'],
                                             order=(3,1,2),
                                             seasonal_order=(3, 0, 0, 12),
                                             enforce_stationarity=False,
                                             enforce_invertibility=False)
results_full_data_model = full_data_model.fit(maxiter=1000)
print(results_full_data_model.summary())
```

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

/opt/anaconda3/anaconda3/lib/python3.8/site-packages/statsmodels/tsa/base/tsa\_model.py:159: ValueWarning:

No frequency information was provided, so inferred frequency MS will be used.

#### SARIMAX Results

```
=====
Dep. Variable:          Sparkling      No. Observations:          187
Model:                SARIMAX(3, 1, 2)x(3, 0, [], 12)      Log Likelihood          -1088.251
Date:                  Sat, 09 Oct 2021      AIC                  2194.503
Time:                  07:01:27      BIC                  2221.417
Sample:                01-01-1980      HQIC                 2205.438
                        - 07-01-1995
Covariance Type:                opg
=====
```

|          | coef    | std err | z      | P> z  | [0.025 | 0.975] |
|----------|---------|---------|--------|-------|--------|--------|
| ar.L1    | -0.5213 | 0.694   | -0.751 | 0.453 | -1.882 | 0.840  |
| ar.L2    | 0.0118  | 0.157   | 0.075  | 0.940 | -0.296 | 0.319  |
| ar.L3    | 0.0183  | 0.104   | 0.176  | 0.860 | -0.185 | 0.221  |
| ma.L1    | -0.3226 | 0.698   | -0.462 | 0.644 | -1.691 | 1.046  |
| ma.L2    | -0.6298 | 0.676   | -0.932 | 0.351 | -1.954 | 0.694  |
| ar.S.L12 | 0.4978  | 0.074   | 6.694  | 0.000 | 0.352  | 0.644  |
| ar.S.L24 | 0.3234  | 0.096   | 3.384  | 0.001 | 0.136  | 0.511  |
| ar.S.L36 | 0.1887  | 0.097   | 1.942  | 0.052 | -0.002 | 0.379  |

```

sigma2      1.557e+05   1.69e+04    9.217      0.000    1.23e+05   1.89e+05
=====
Ljung-Box (Q):                23.13   Jarque-Bera (JB):                24.97
Prob(Q):                      0.98   Prob(JB):                      0.00
Heteroskedasticity (H):       0.91   Skew:                          0.52
Prob(H) (two-sided):          0.74   Kurtosis:                      4.74
=====

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [131]: pred_SARIMA = predicted_full_data.summary_frame(alpha=0.05)
pred_SARIMA.head()
```

Out[131]:

| Sparkling  | mean        | mean_se    | mean_ci_lower | mean_ci_upper |
|------------|-------------|------------|---------------|---------------|
| 1995-08-01 | 1941.248027 | 394.580175 | 1167.885096   | 2714.610958   |
| 1995-09-01 | 2451.224570 | 399.355354 | 1668.502460   | 3233.946680   |
| 1995-10-01 | 3305.033728 | 399.450055 | 2522.126007   | 4087.941449   |
| 1995-11-01 | 3954.622572 | 400.668704 | 3169.326341   | 4739.918802   |
| 1995-12-01 | 6171.799475 | 400.683867 | 5386.473526   | 6957.125424   |


```
In [132]: pred_SARIMA.tail()
```

Out[132]:

| Sparkling  | mean        | mean_se    | mean_ci_lower | mean_ci_upper |
|------------|-------------|------------|---------------|---------------|
| 1996-03-01 | 1816.844449 | 401.415537 | 1030.084455   | 2603.604444   |
| 1996-04-01 | 1843.101124 | 401.598084 | 1055.983343   | 2630.218906   |
| 1996-05-01 | 1676.311269 | 401.801687 | 888.794434    | 2463.828104   |
| 1996-06-01 | 1631.770697 | 401.995276 | 843.874435    | 2419.666959   |
| 1996-07-01 | 2012.405778 | 402.193311 | 1224.121373   | 2800.690183   |

```
In [133]: # Plotting the predictions:
```

```
axis = df['Sparkling'].plot(label='Observed series')  
pred_SARIMA['mean'].plot(ax=axis, label='Forecasted series', alpha=0.7)  
  
plt.title('Forecasted Sparkling Sales (SARIMA model)',color='blue',fontsize=20)  
plt.xlabel('Year/Month',color='blue',fontsize=18)  
plt.ylabel('Sparkling Sales',color='blue',fontsize=18)  
  
plt.legend(loc='best')  
plt.grid();
```








```
In [172]: # Plotting the predictions with confidence interval:

axis = df['Sparkling'].plot(label='Observed series')
pred_SARIMA['mean'].plot(ax=axis, label='Forecasted series', alpha=0.7)

axis.fill_between(pred_SARIMA['mean'].index, pred_SARIMA['mean_ci_lower'], pred_SARIMA['mean_ci_upper'], c

plt.title('Forecasted Sparkling Sales (SARIMA model)', color='blue', fontsize=20)
plt.xlabel('Year/Month', color='blue', fontsize=18)
plt.ylabel('Sparkling Sales', color='blue', fontsize=18)

plt.legend(loc='best')
plt.grid();
```



```
In [134]: ## THE END
```